

CPU design effects that can degrade performance of your programs

Jakub Beránek
jakub.beranek@vsb.cz

whoami

- PhD student @ VSB-TUO, Ostrava, Czech Republic
- Research assistant @ IT4Innovations (HPC center)
- HPC, distributed systems, program optimization

How do we get maximum performance?

- Select the right algorithm

How do we get maximum performance?

- Select the right algorithm
- Use a low-overhead language

How do we get maximum performance?

- Select the right algorithm
- Use a low-overhead language
- Compile properly

How do we get maximum performance?

- Select the right algorithm
- Use a low-overhead language
- Compile properly
- **Tune to the underlying hardware**

Why should we care?

- We write code for the C++ abstract machine

Why should we care?

- We write code for the C++ abstract machine
- Intel CPUs fulfill the contract of this abstract machine

Why should we care?

- We write code for the C++ abstract machine
- Intel CPUs fulfill the contract of this abstract machine
 - But inside they can do whatever they want

Why should we care?

- We write code for the C++ abstract machine
- Intel CPUs fulfill the contract of this abstract machine
 - But inside they can do whatever they want
- Understanding CPU trade-offs can get us more performance

C++ abstract machine example

```
void foo(int* arr, int count)
{
    for (int i = 0; i < count; i++)
    {
        arr[i]++;
    }
}
```

How fast are the individual array increments?

Hardware effects

- Performance effects caused by a specific CPU/memory implementation

Hardware effects

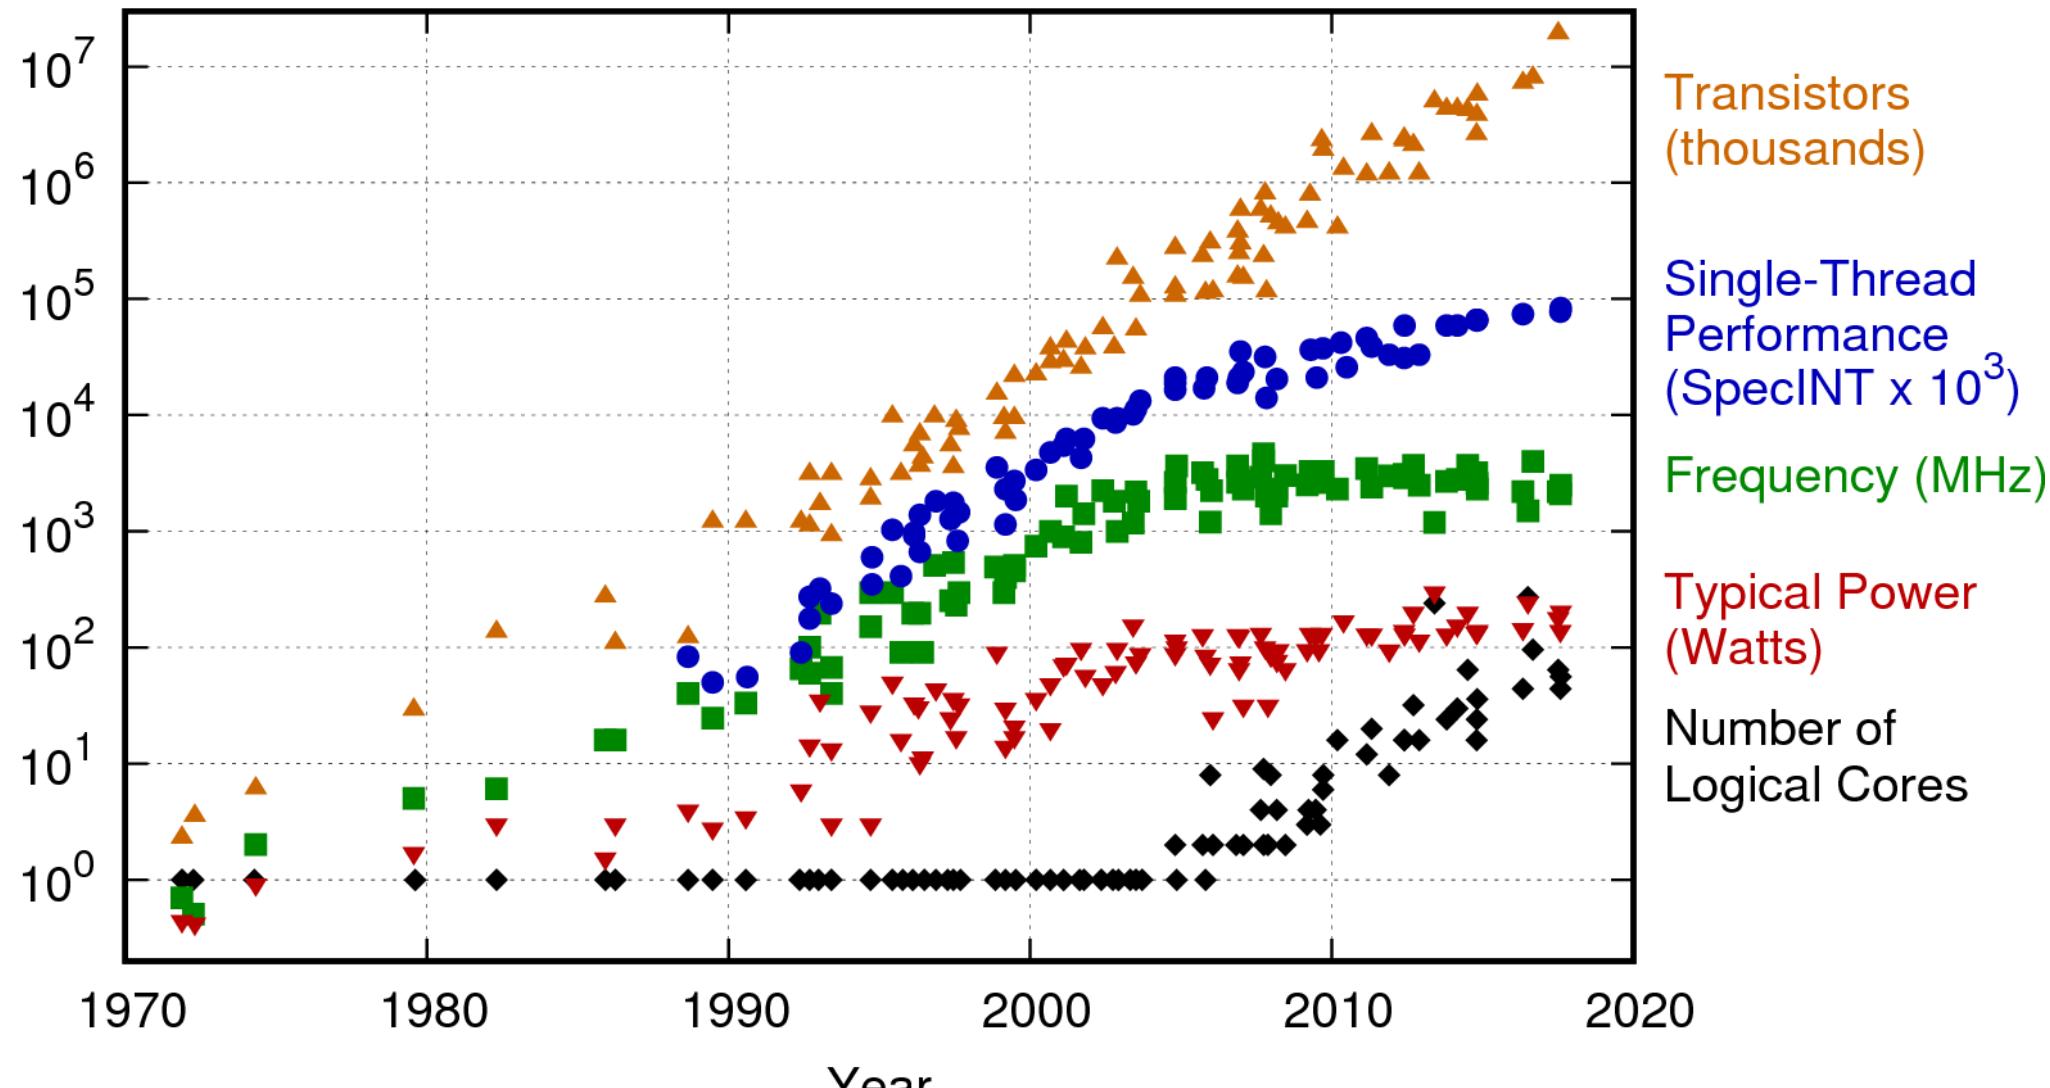
- Performance effects caused by a specific CPU/memory implementation
- Demonstrate some CPU/memory trade-off or assumption

Hardware effects

- Performance effects caused by a specific CPU/memory implementation
- Demonstrate some CPU/memory trade-off or assumption
- Impossible to predict from (C++) code alone

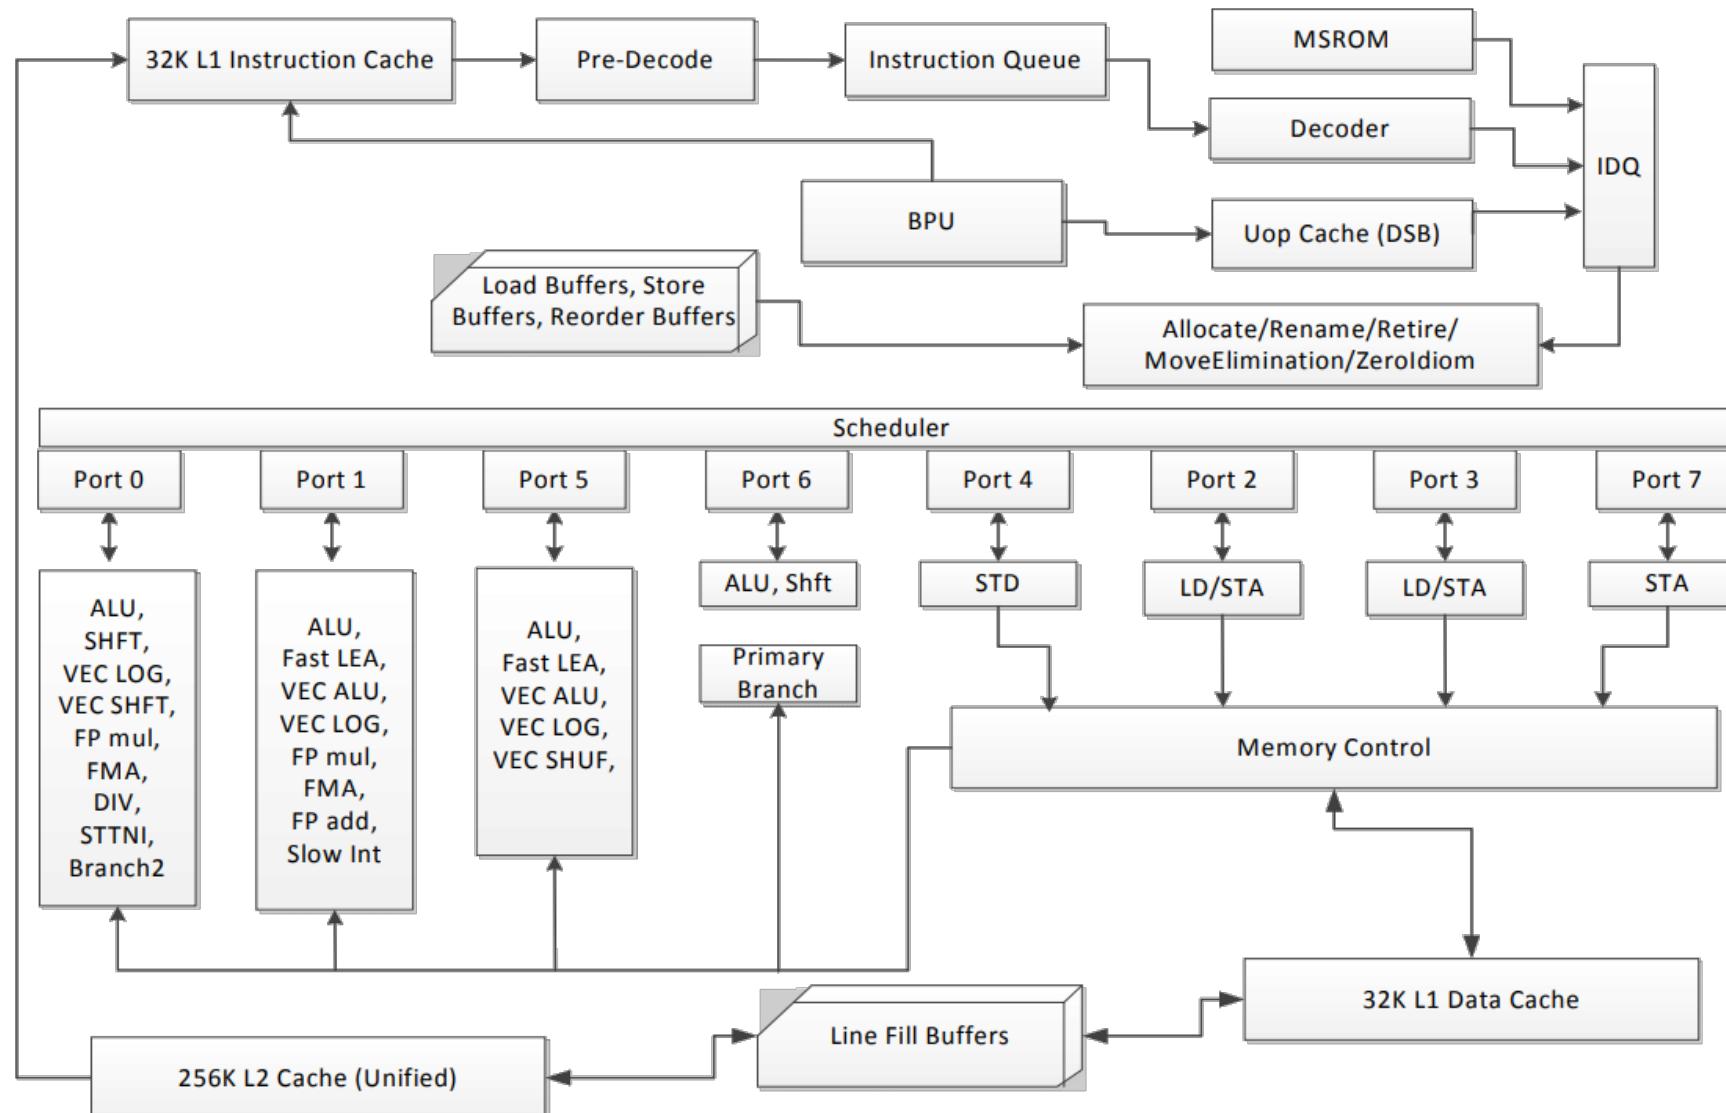
Hardware is getting more and more complex

42 Years of Microprocessor Trend Data



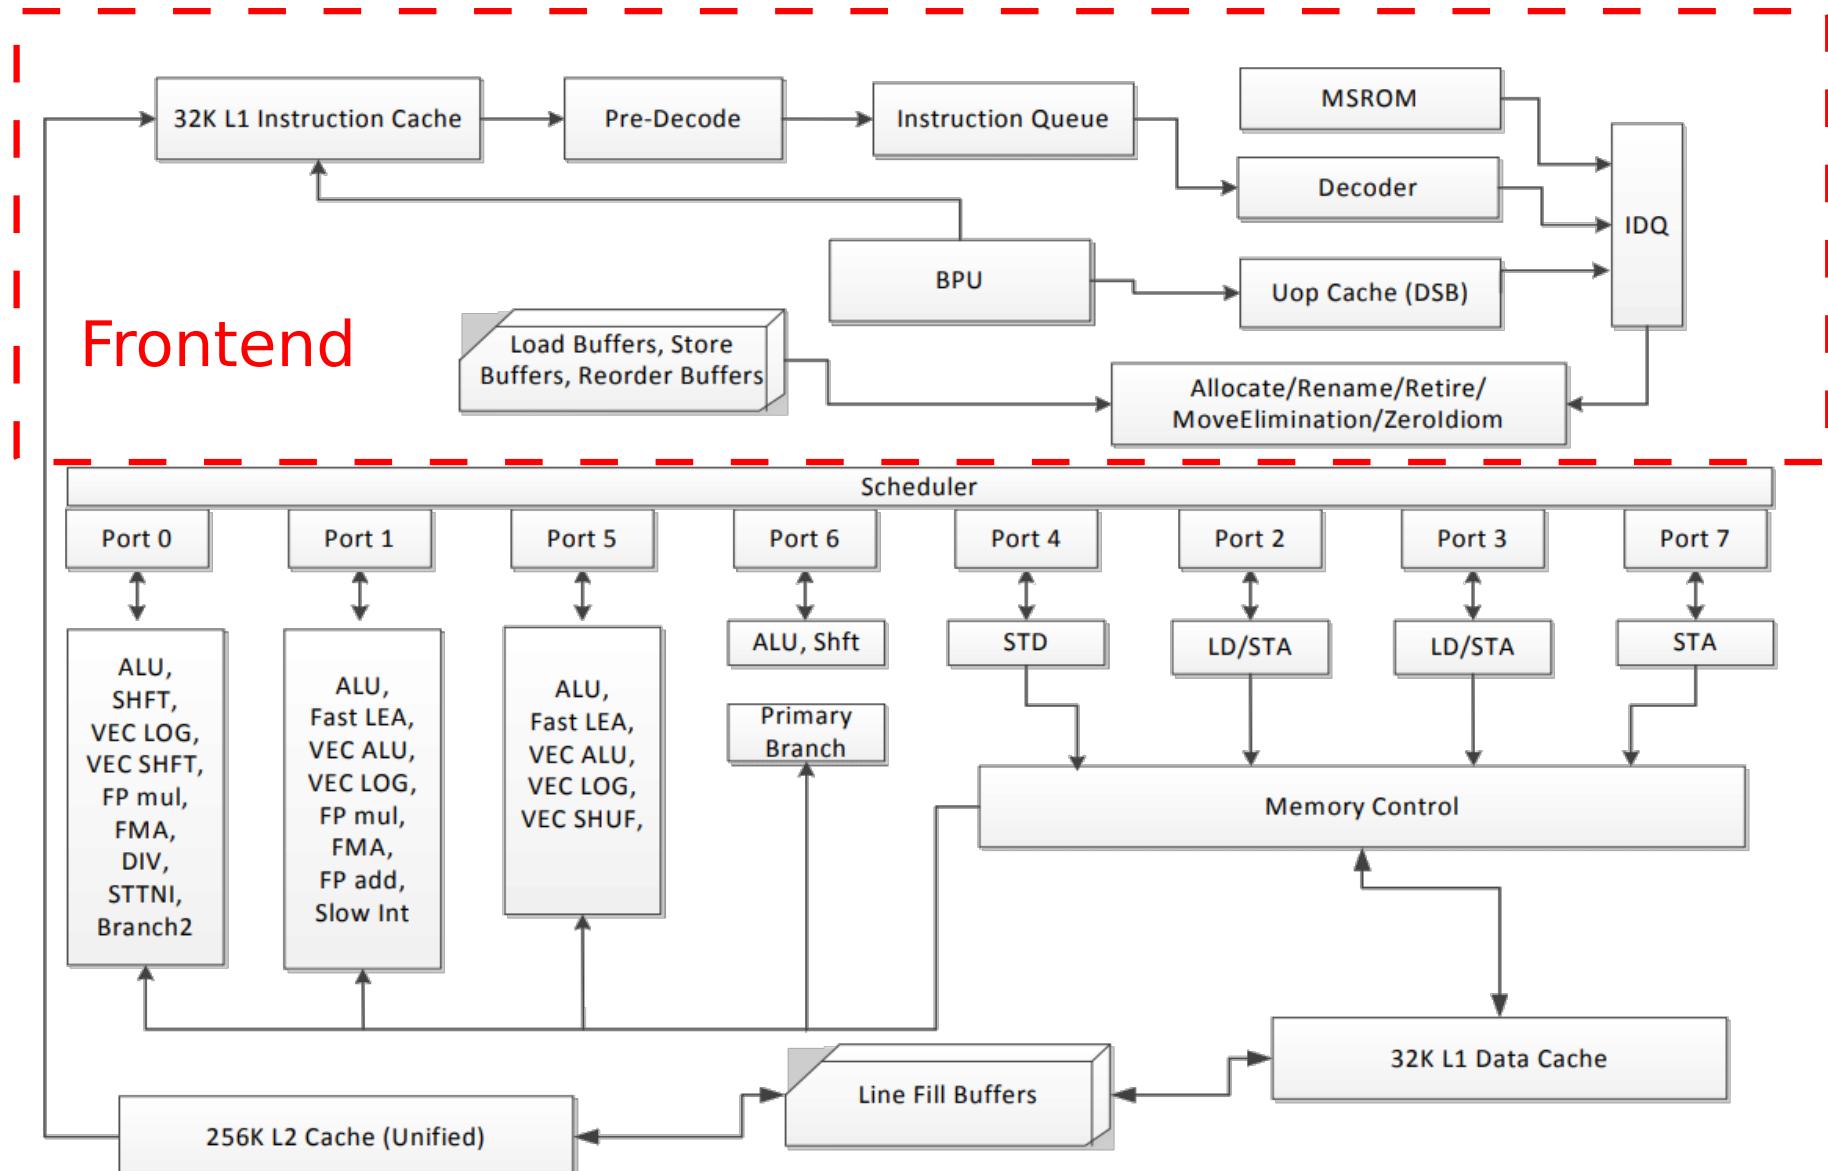
Source: karlrupp.net

Microarchitecture (Haswell)



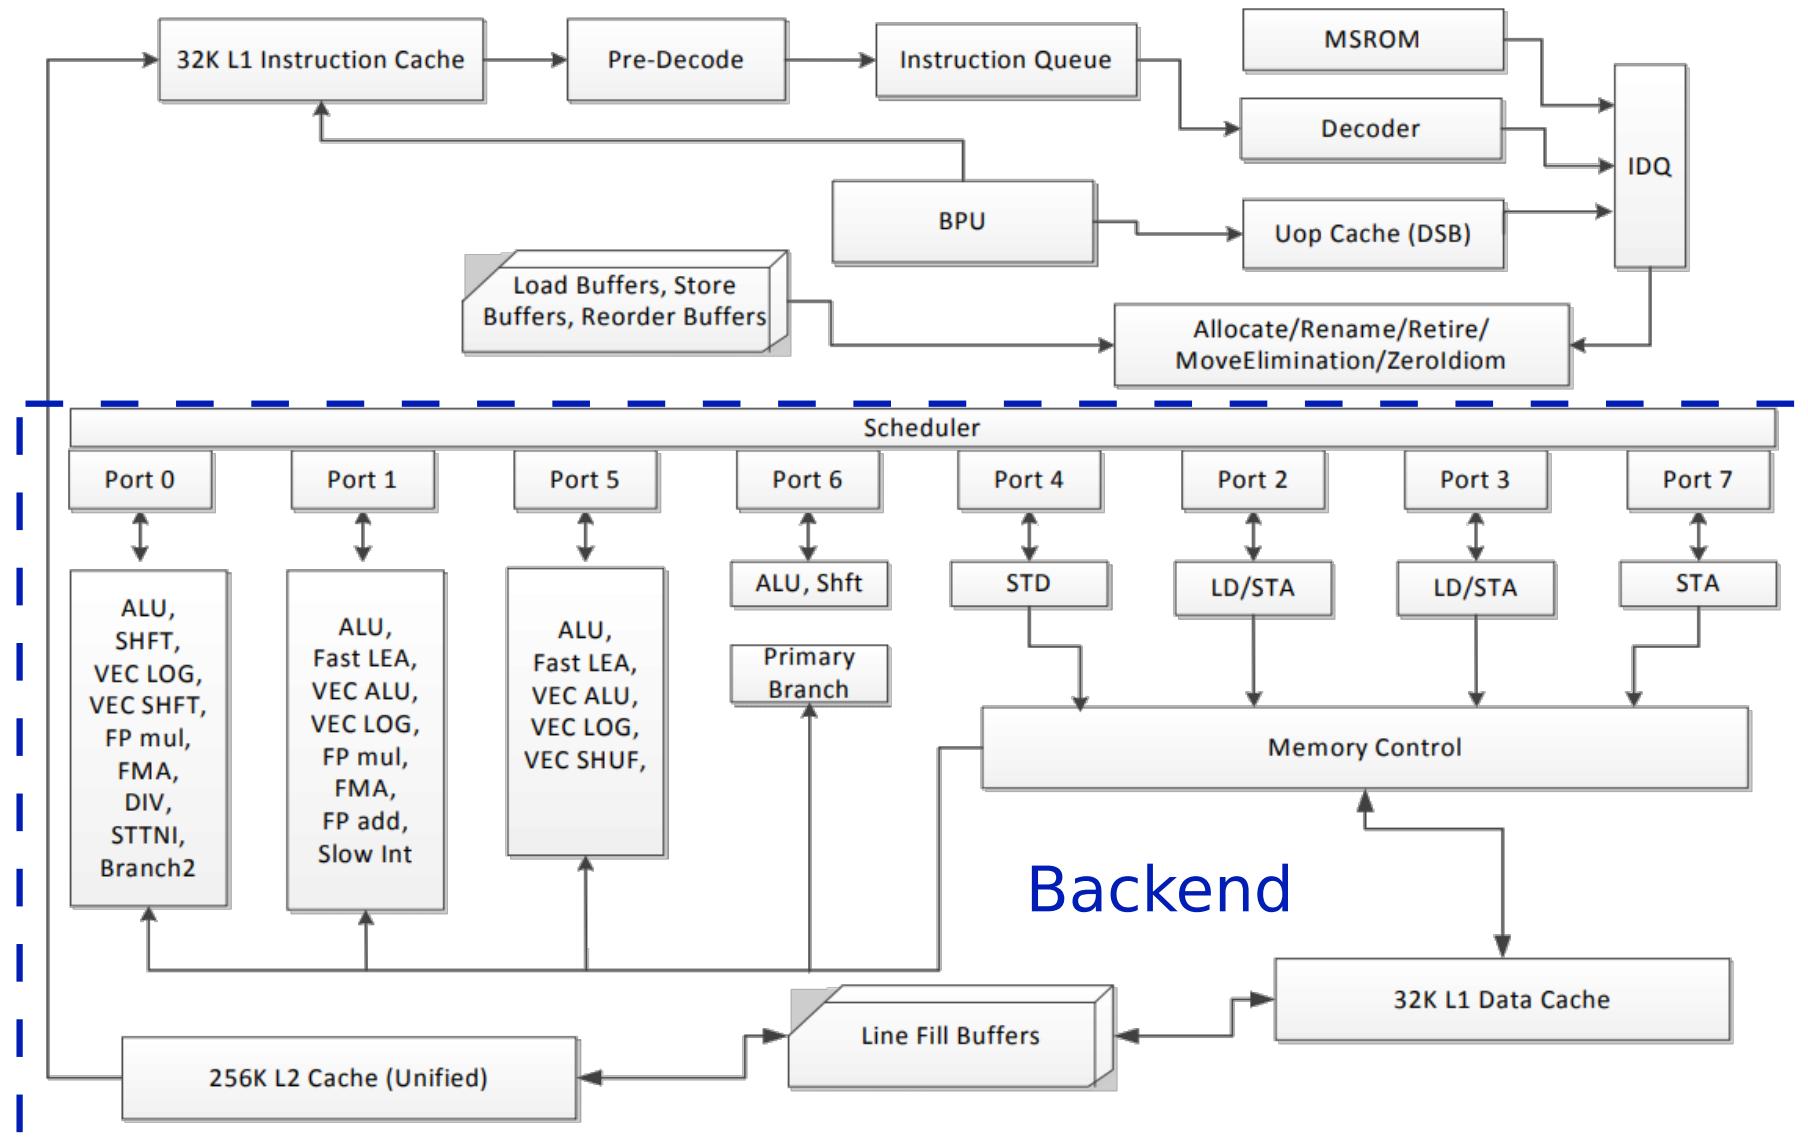
Source: Intel Architectures Optimization Reference Manual

Microarchitecture (Haswell)



Source: Intel Architectures Optimization Reference Manual

Microarchitecture (Haswell)



Source: Intel Architectures Optimization Reference Manual

How bad is it?

- C++ 17 final draft:

<http://www.open-std.org/jtc1/sc22/wg21/docs/papers/2017/n4659.pdf>

<https://software.intel.com/sites/default/files/managed/39/c5/325462-sdm-vol-1-2abcd-3abcd.pdf>

<https://software.intel.com/sites/default/files/managed/9e/bc/64-ia-32-architectures-optimization-manual.pdf>

How bad is it?

- C++ 17 final draft: 1622 pages

<http://www.open-std.org/jtc1/sc22/wg21/docs/papers/2017/n4659.pdf>

<https://software.intel.com/sites/default/files/managed/39/c5/325462-sdm-vol-1-2abcd-3abcd.pdf>

<https://software.intel.com/sites/default/files/managed/9e/bc/64-ia-32-architectures-optimization-manual.pdf>

How bad is it?

- C++ 17 final draft: 1622 pages
- Intel x86 manual:

<http://www.open-std.org/jtc1/sc22/wg21/docs/papers/2017/n4659.pdf>

<https://software.intel.com/sites/default/files/managed/39/c5/325462-sdm-vol-1-2abcd-3abcd.pdf>

<https://software.intel.com/sites/default/files/managed/9e/bc/64-ia-32-architectures-optimization-manual.pdf>

How bad is it?

- C++ 17 final draft: 1622 pages
- Intel x86 manual: **5764** pages!

<http://www.open-std.org/jtc1/sc22/wg21/docs/papers/2017/n4659.pdf>

<https://software.intel.com/sites/default/files/managed/39/c5/325462-sdm-vol-1-2abcd-3abcd.pdf>

<https://software.intel.com/sites/default/files/managed/9e/bc/64-ia-32-architectures-optimization-manual.pdf>

Plan of attack

- Show example C++ programs

Plan of attack

- Show example C++ programs
 - short, (hopefully) comprehensible

Plan of attack

- Show example C++ programs
 - short, (hopefully) comprehensible
 - compiled with -O3

Plan of attack

- Show example C++ programs
 - short, (hopefully) comprehensible
 - compiled with -O3
- Demonstrate weird performance behaviour

Plan of attack

- Show example C++ programs
 - short, (hopefully) comprehensible
 - compiled with -O3
- Demonstrate weird performance behaviour
- Let you guess what might cause it

Plan of attack

- Show example C++ programs
 - short, (hopefully) comprehensible
 - compiled with -O3
- Demonstrate weird performance behaviour
- Let you guess what might cause it
- Explain (a possible) cause

Plan of attack

- Show example C++ programs
 - short, (hopefully) comprehensible
 - compiled with -O3
- Demonstrate weird performance behaviour
- Let you guess what might cause it
- Explain (a possible) cause
- Show how to measure and fix it

Plan of attack

- Show example C++ programs
 - short, (hopefully) comprehensible
 - compiled with -O3
- Demonstrate weird performance behaviour
- Let you guess what might cause it
- Explain (a possible) cause
- Show how to measure and fix it
- Disclaimer #1: Everything will be Intel x86 specific

Plan of attack

- Show example C++ programs
 - short, (hopefully) comprehensible
 - compiled with -O3
- Demonstrate weird performance behaviour
- Let you guess what might cause it
- Explain (a possible) cause
- Show how to measure and fix it
- Disclaimer #1: Everything will be Intel x86 specific
- Disclaimer #2: I'm not an expert on this and I may be wrong :-)



Let's see some examples...



Code (backup)

```
std::vector<float> data = /* 32K random floats in [1, 10] */;
float sum = 0;
// std::sort(data.begin(), data.end());
for (auto x : data)
{
    if (x < 6.0f)
    {
        sum += x;
    }
}
```

Result (backup)

Benchmark	Time	CPU	Iterations
<hr/>			
filter_nosort/32768	133460 ns	132992 ns	5284
filter_sorted/32768	63069 ns	62991 ns	12547



Most upvoted Stack Overflow question

Why is it faster to process a sorted array than an unsorted array?



Here is a piece of C++ code that seems very peculiar. For some strange reason, sorting the data miraculously makes the code almost six times faster.

22447



10294

```
#include <algorithm>
#include <ctime>
#include <iostream>

int main()
{
    // Generate data
    const unsigned arraySize = 32768;
    int data[arraySize];

    for (unsigned c = 0; c < arraySize; ++c)
        data[c] = std::rand() % 256;

    // !!! With this, the next loop runs faster
    std::sort(data, data + arraySize);

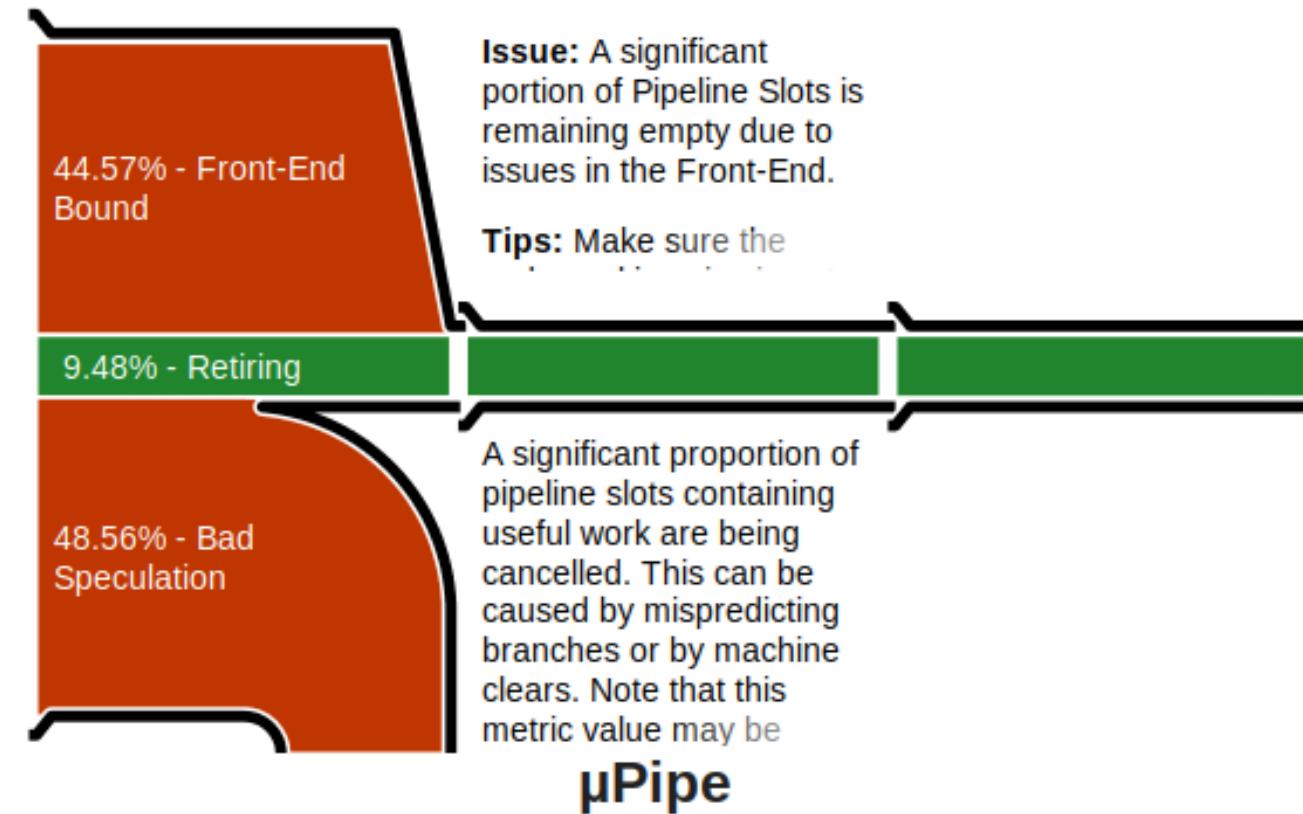
    // Test
    clock_t start = clock();
    long long sum = 0;

    for (unsigned i = 0; i < 100000; ++i)
    {
        // Primary loop
        for (unsigned c = 0; c < arraySize; ++c)
        {
            if (data[c] >= 128)
                sum += data[c];
        }
    }
}
```

What is going on? (Intel Amplifier - VTune)

Elapsed Time [?]: 2.222s

Clockticks:	6,736,800,000
Instructions Retired:	2,942,800,000
CPI Rate [?] :	2.289
MUX Reliability [?] :	0.802
⌚ Retiring [?] :	9.5% of Pipeline Slots
⌚ Front-End Bound [?] :	44.6% of Pipeline Slots
⌚ Front-End Latency [?] :	31.1% of Pipeline Slots
ICache Misses [?] :	0.0% of Clockticks
ITLB Overhead [?] :	0.2% of Clockticks
⌚ Branch Resteers [?] :	13.4% of Clockticks
Mispredicts Resteers [?] :	12.6% of Clockticks
Clears Resteers [?] :	0.0% of Clockticks
Unknown Branches [?] :	0.8% of Clockticks
DSB Switches [?] :	0.0% of Clockticks
Length Changing Prefixes [?] :	0.0% of Clockticks
MS Switches [?] :	4.5% of Clockticks
⌚ Front-End Bandwidth [?] :	13.5% of Pipeline Slots
⌚ Bad Speculation [?] :	48.6% of Pipeline Slots
Branch Mispredict [?] :	48.6% of Pipeline Slots
Machine Clears [?] :	0.0% of Pipeline Slots
⌚ Back-End Bound [?] :	0.0% of Pipeline Slots
Total Thread Count:	1
Paused Time [?] :	0s



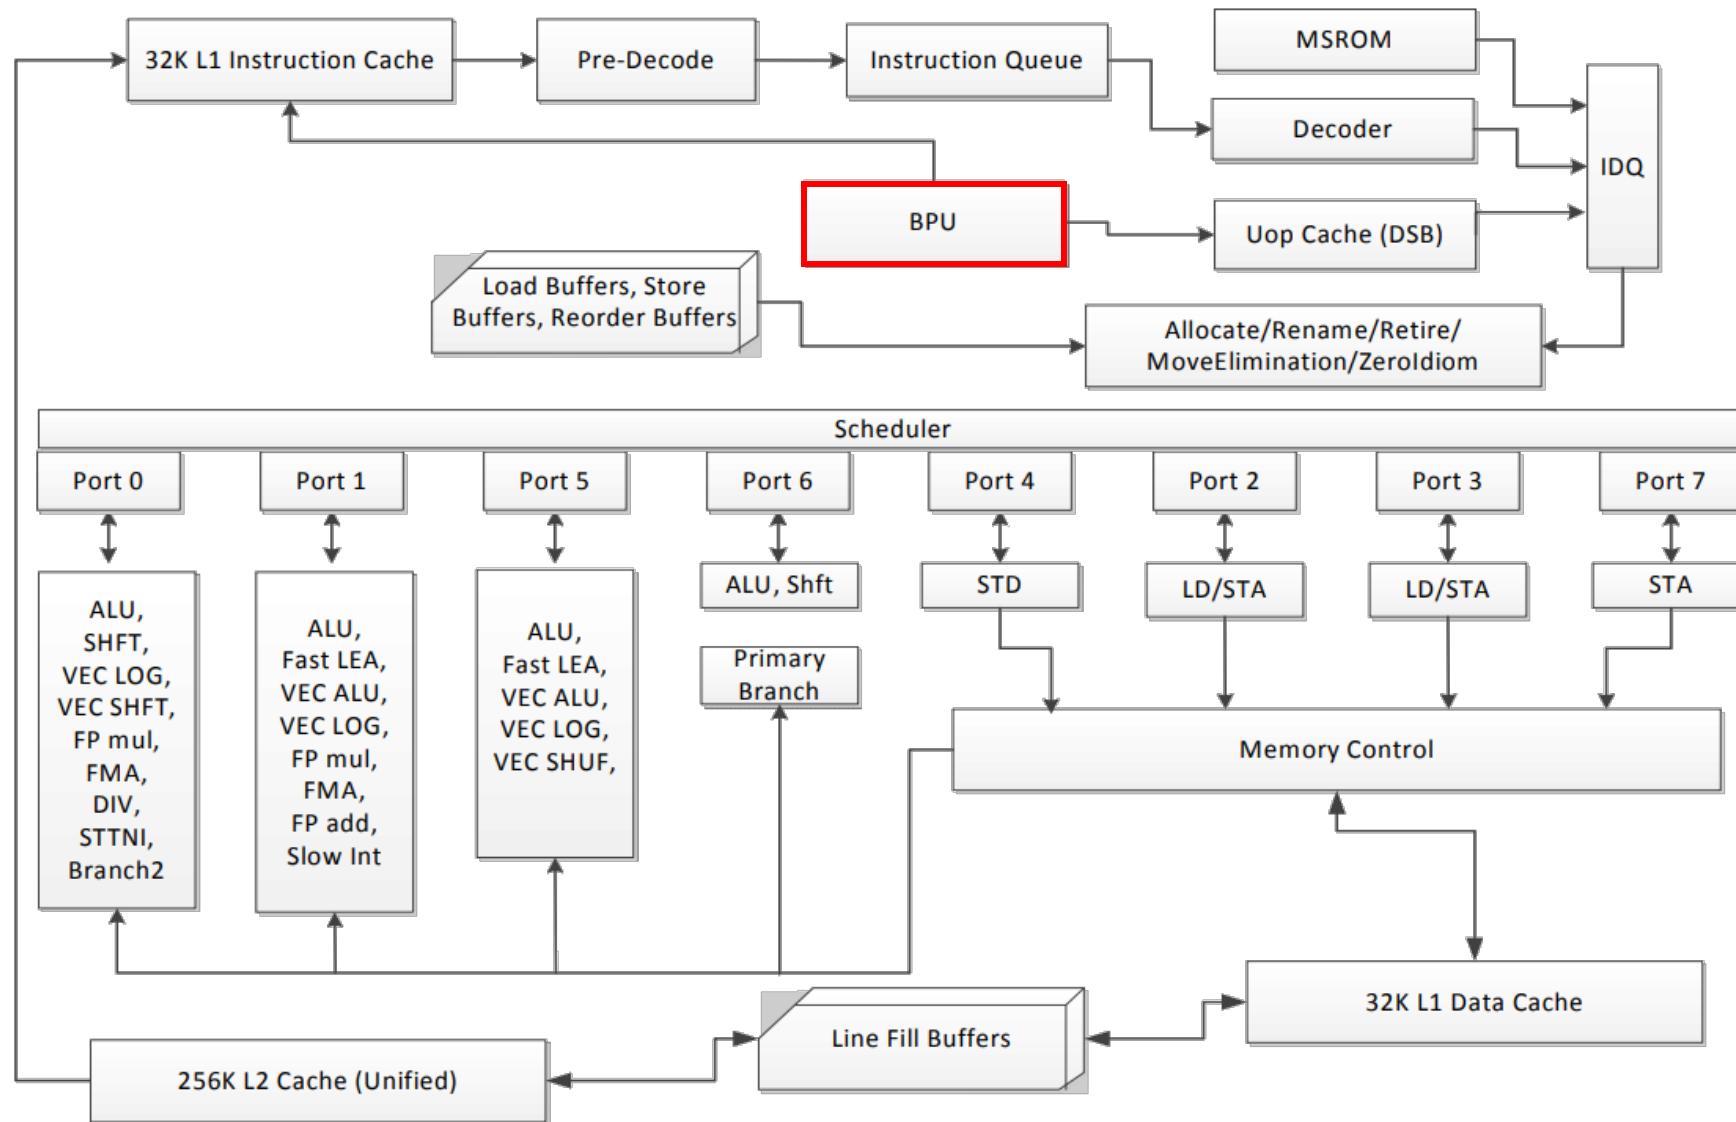
μPipe

This diagram represents inefficiencies in CPU usage. Treat it as a pipe with an output flow equal to the "pipe efficiency" ratio: (Actual Instructions Retired)/(Maximum Possible Instruction Retired). If there are pipeline stalls decreasing the pipe efficiency, the pipe shape gets more narrow.

What is going on? (perf)

```
$ perf stat ./example0a --benchmark_filter=nosort
 853,672 012  task-clock (msec) # 0,997 CPUs utilized
          30  context-switches  # 0,035 K/sec
          0  cpu-migrations  # 0,000 K/sec
         199  page-faults    # 0,233 K/sec
 3 159 530 915  cycles          # 3,701 GHz
 1 475 799 619  instructions   # 0,47 insn per cycle
 419 608 357  branches        # 491,533 M/sec
 102 425 035  branch-misses  # 24,41% of all branches
```

Branch predictor



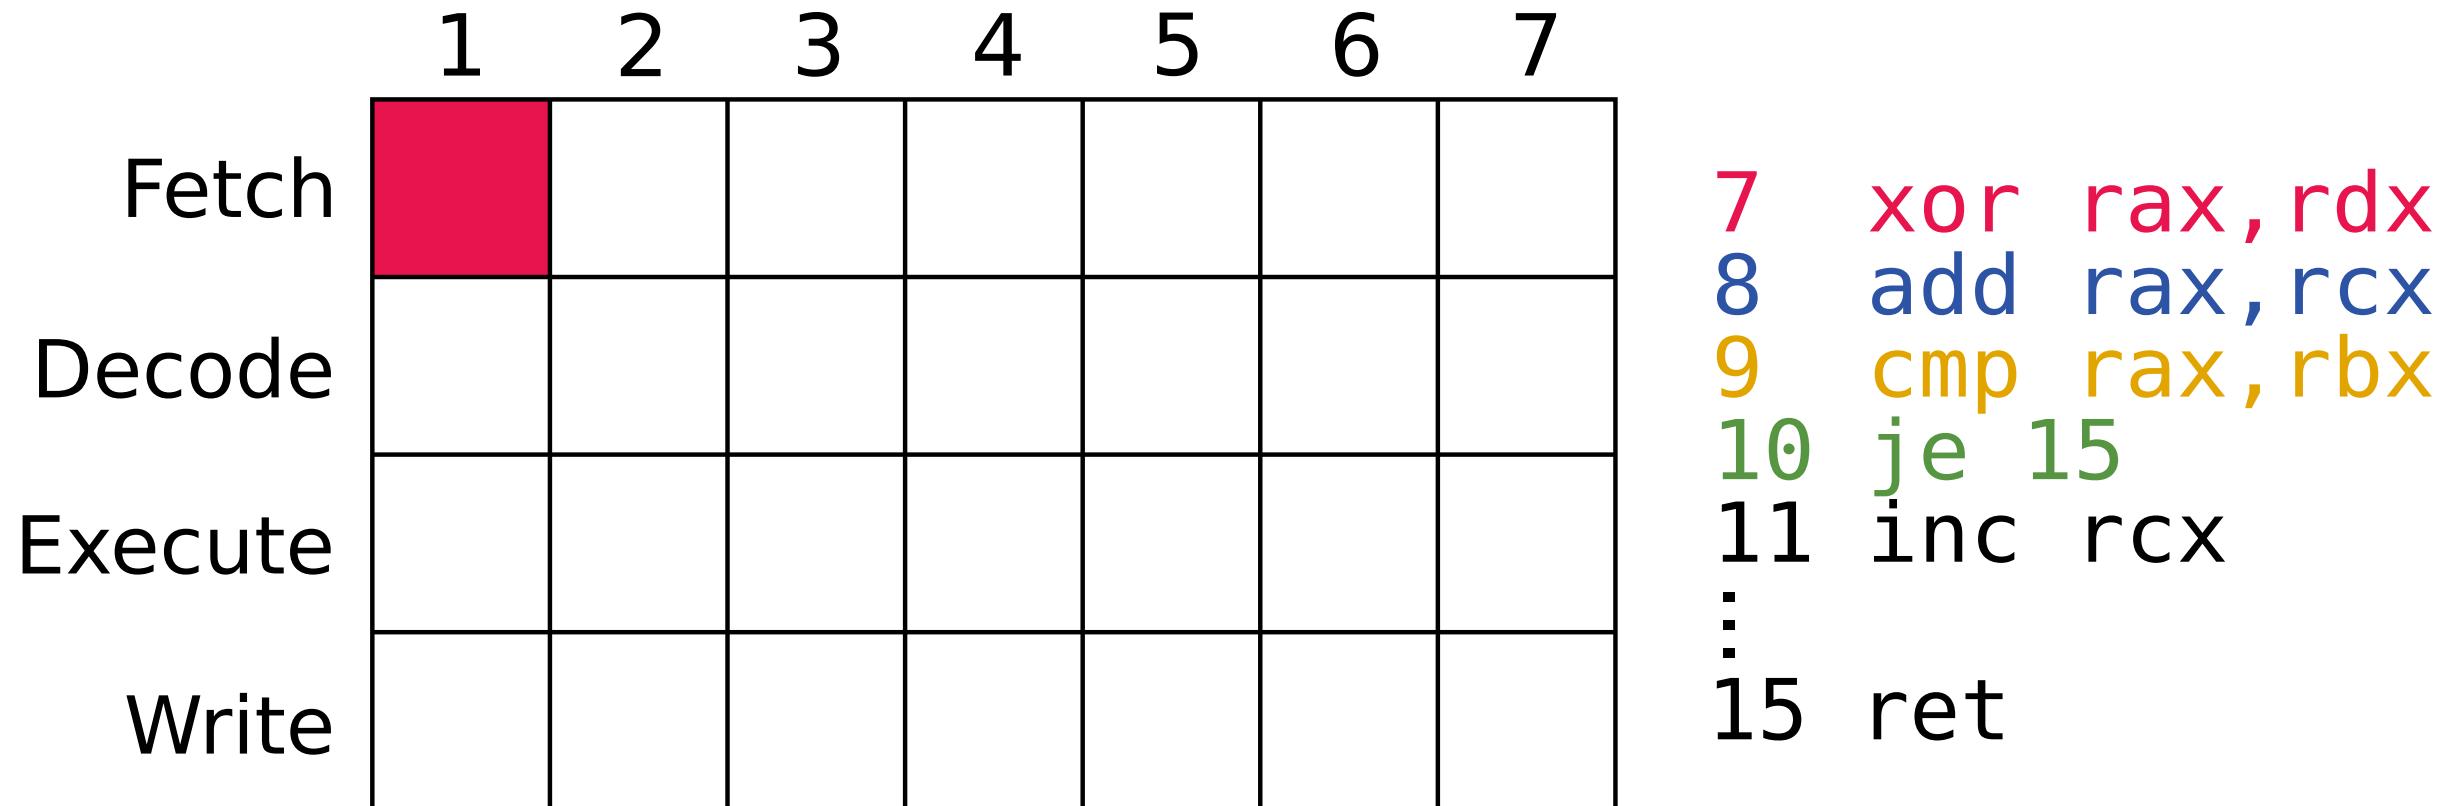
CPU pipeline 101

	1	2	3	4	5	6	7
Fetch							
Decode							
Execute							
Write							

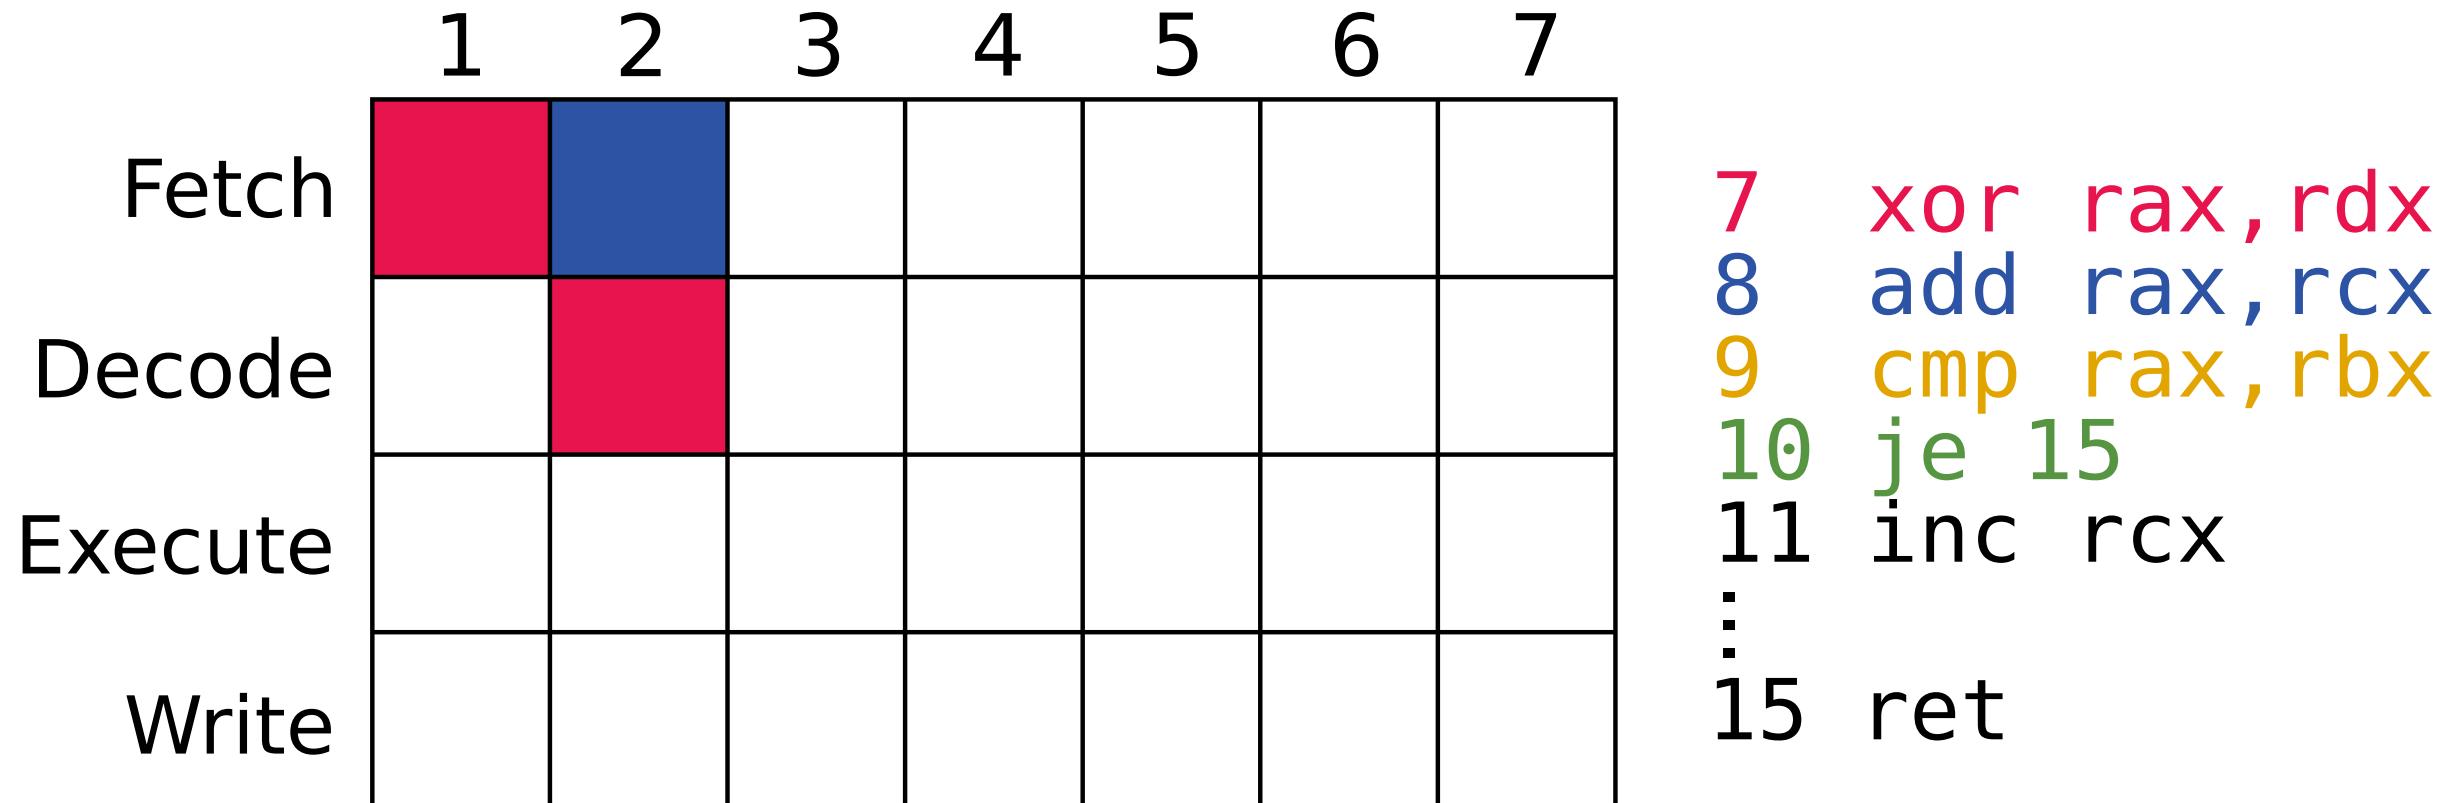
7 xor rax, rdx
8 add rax, rcx
9 cmp rax, rbx
10 je 15
11 inc rcx
...
15 ret



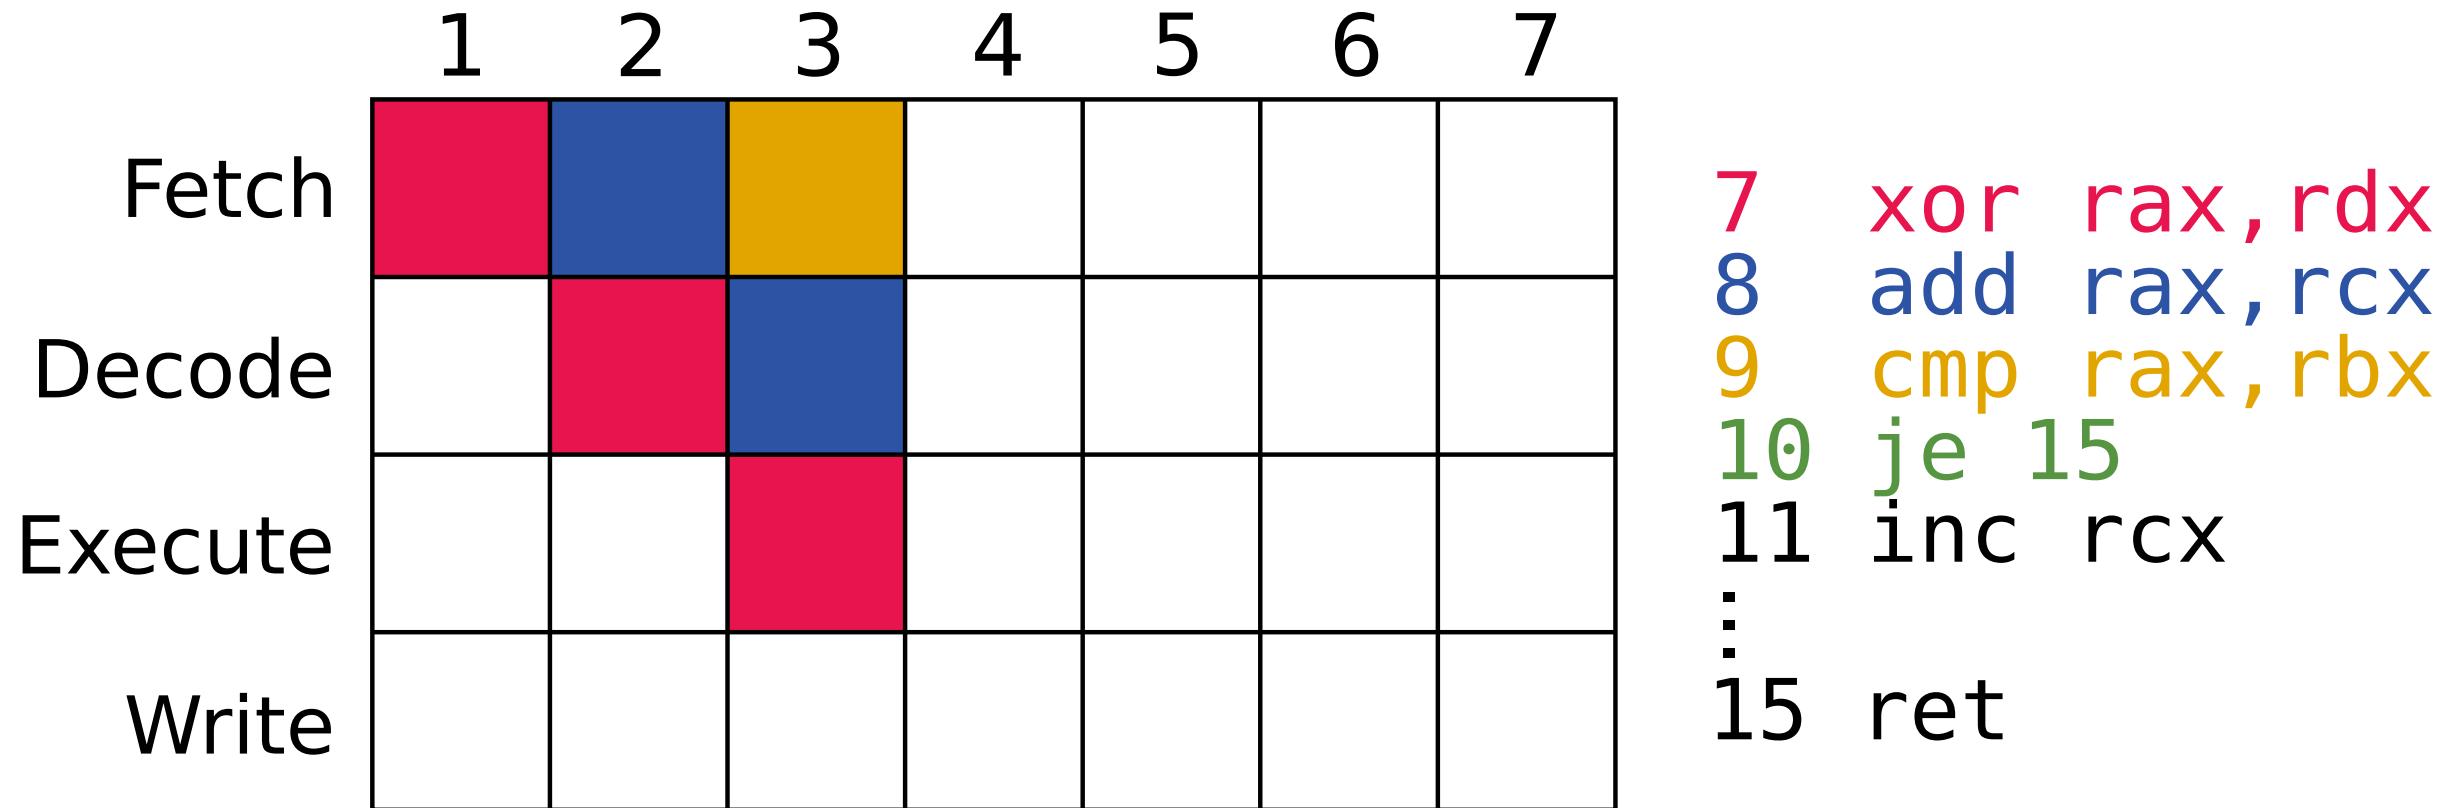
CPU pipeline 101



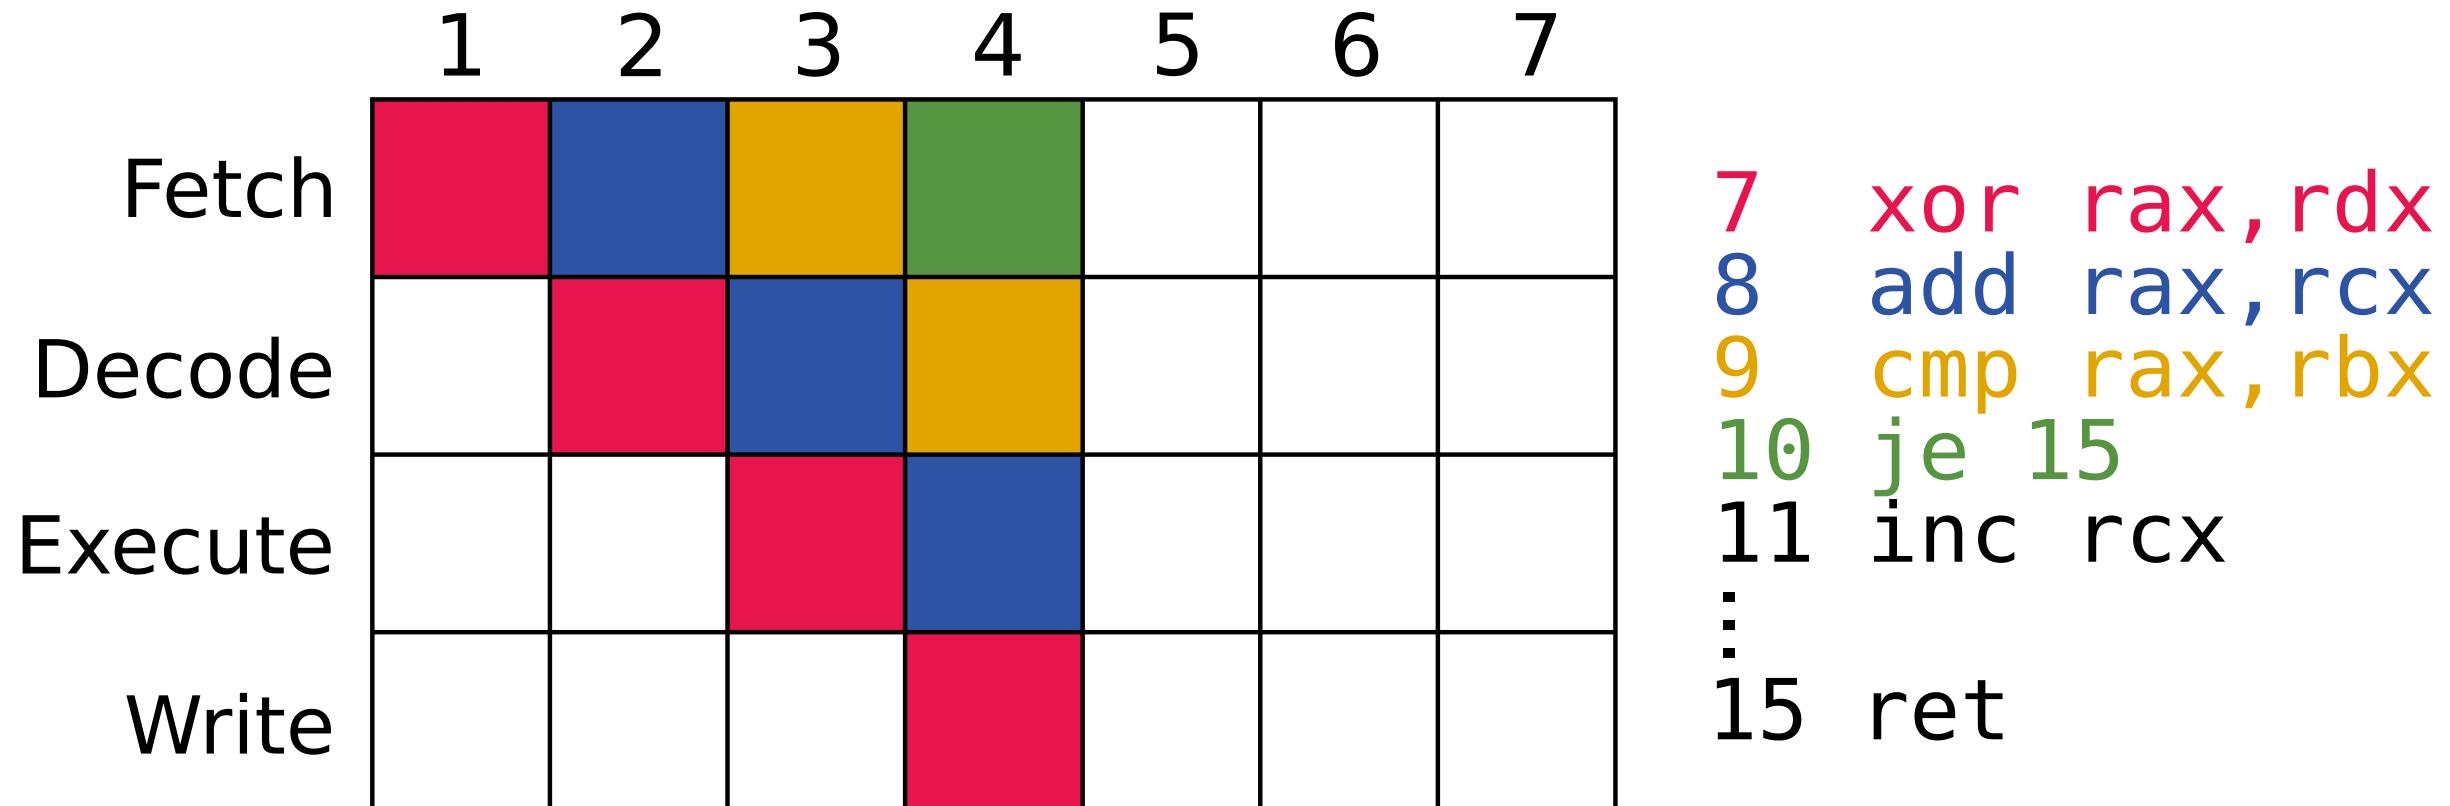
CPU pipeline 101



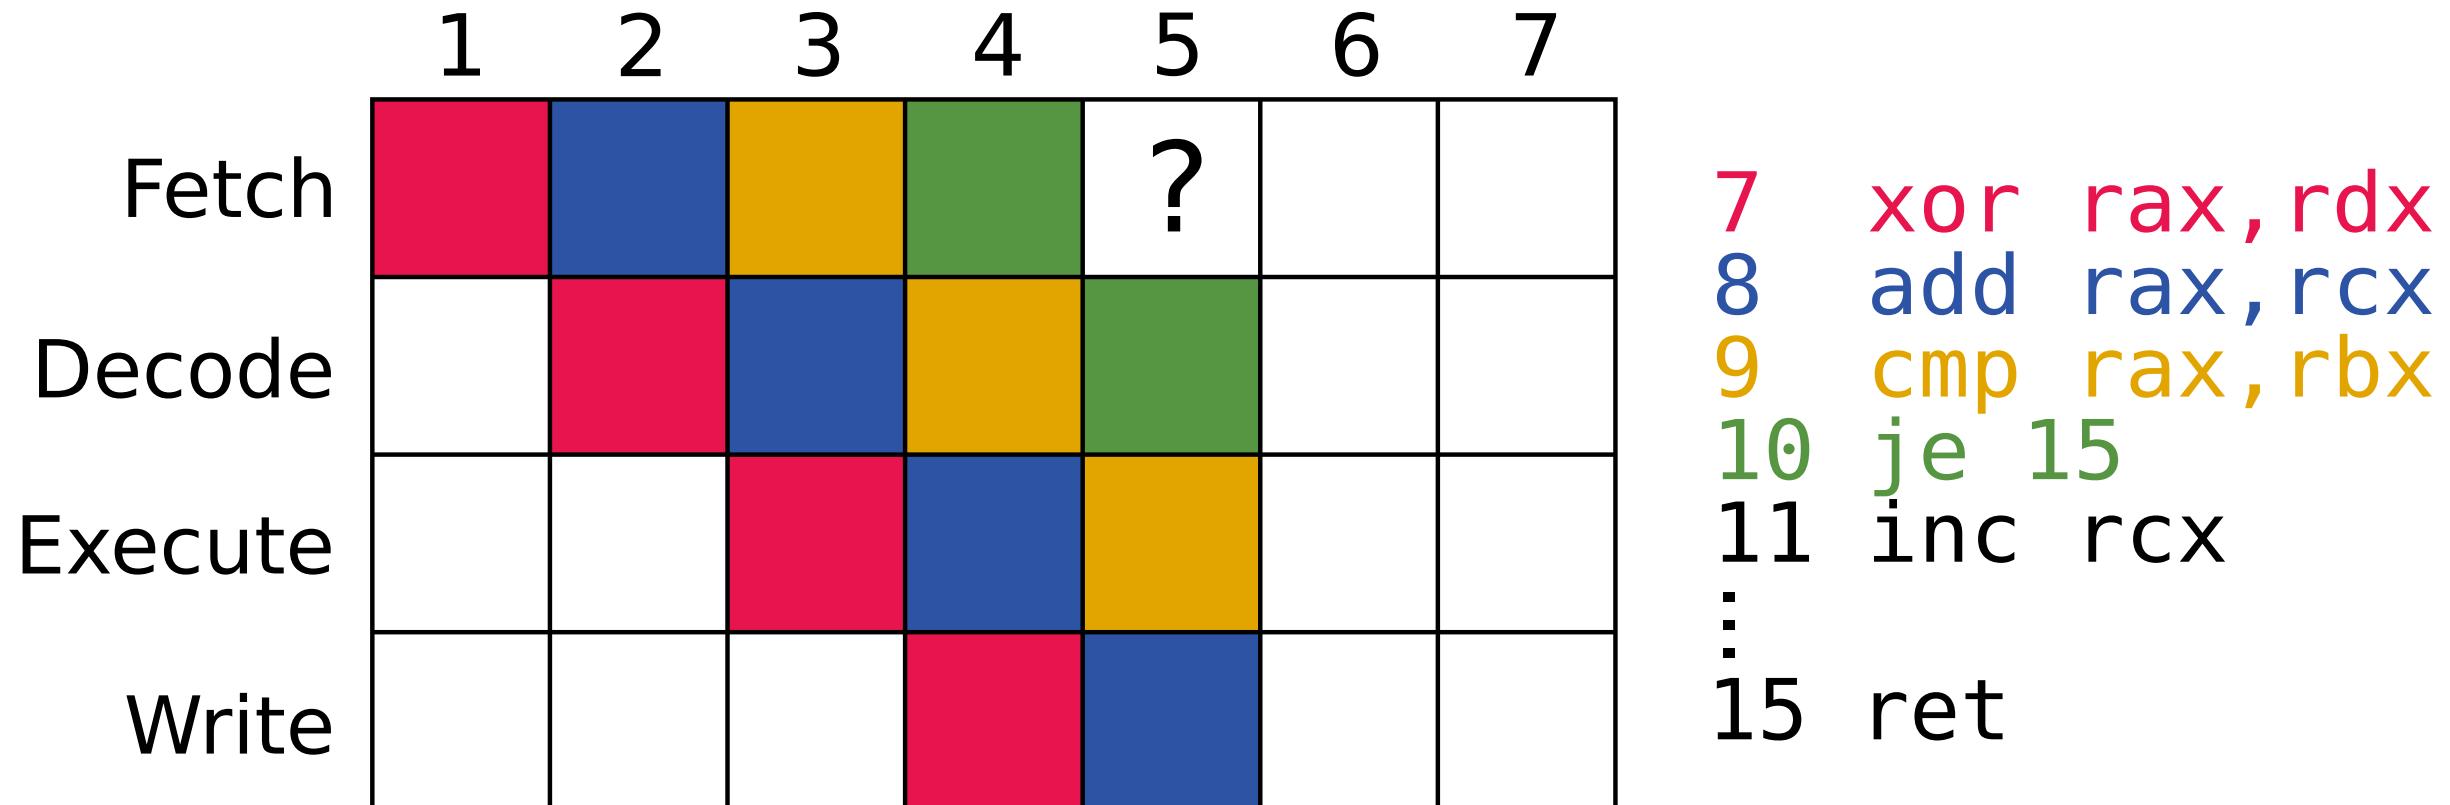
CPU pipeline 101



CPU pipeline 101



CPU pipeline 101



Branch predictor

- CPU tries to predict results of branches



Branch predictor

- CPU tries to predict results of branches
- Misprediction can cost ~15-20 cycles!



Simple branch predictor - unsorted array

```
if (data[i] < 6) {  
    ...  
}
```

6	2	1	7	4	8	3	9
---	---	---	---	---	---	---	---

Prediction: Not taken

Simple branch predictor - unsorted array

```
if (data[i] < 6) {  
    ...  
}
```

6	2	1	7	4	8	3	9
---	---	---	---	---	---	---	---

6 < 6?

Prediction: Not taken

Simple branch predictor - unsorted array

```
if (data[i] < 6) {  
    ...  
}
```

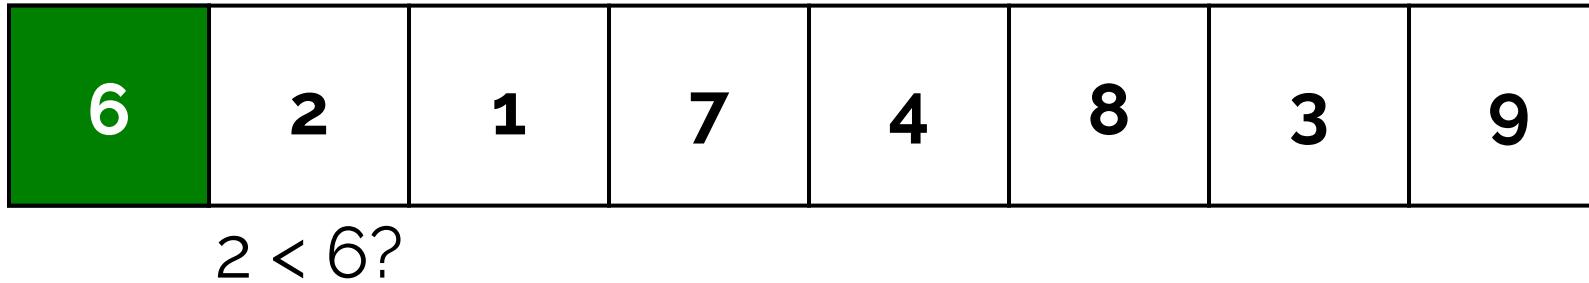
6	2	1	7	4	8	3	9
---	---	---	---	---	---	---	---

6 < 6?

Prediction: Not taken

Simple branch predictor - unsorted array

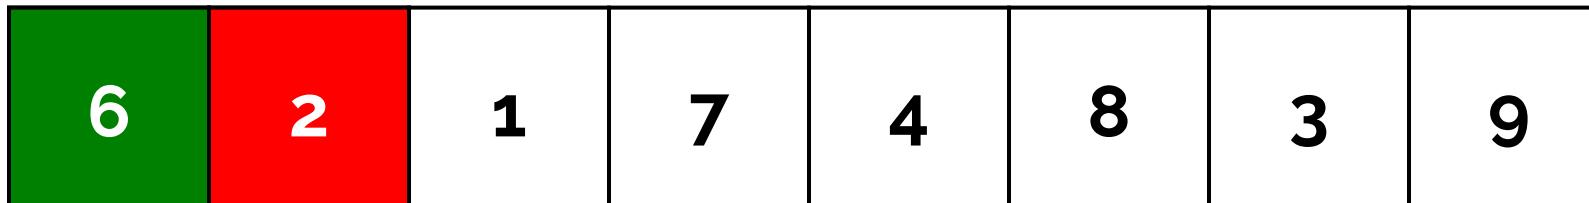
```
if (data[i] < 6) {  
    ...  
}
```



Prediction: Not taken

Simple branch predictor - unsorted array

```
if (data[i] < 6) {  
    ...  
}
```



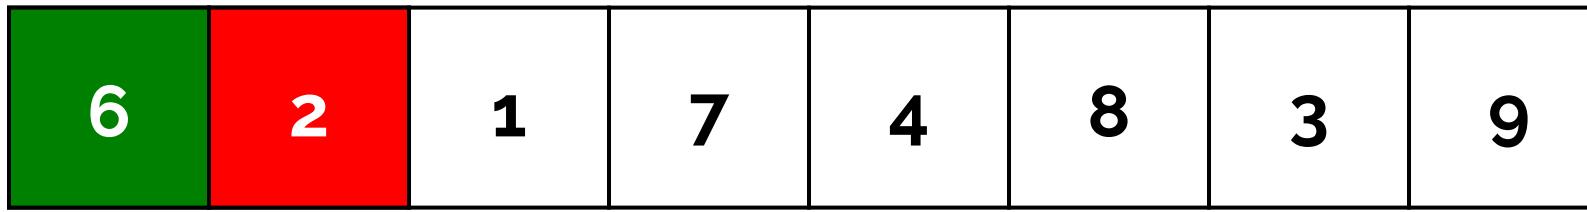
$2 < 6?$

Prediction: Taken



Simple branch predictor - unsorted array

```
if (data[i] < 6) {  
    ...  
}
```



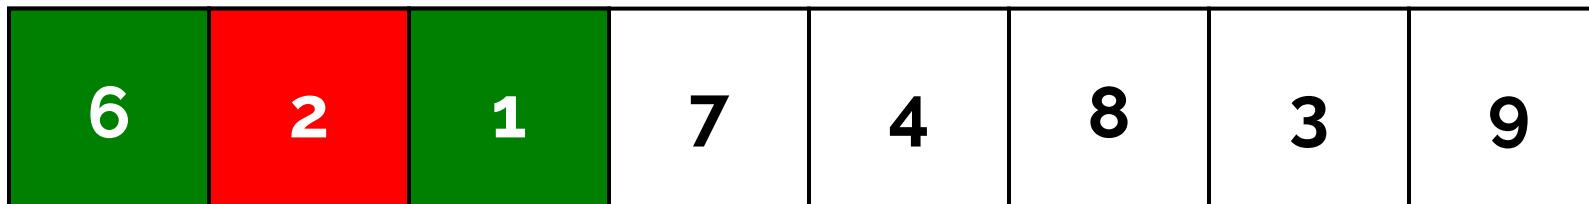
$1 < 6?$

Prediction: Taken



Simple branch predictor - unsorted array

```
if (data[i] < 6) {  
    ...  
}
```

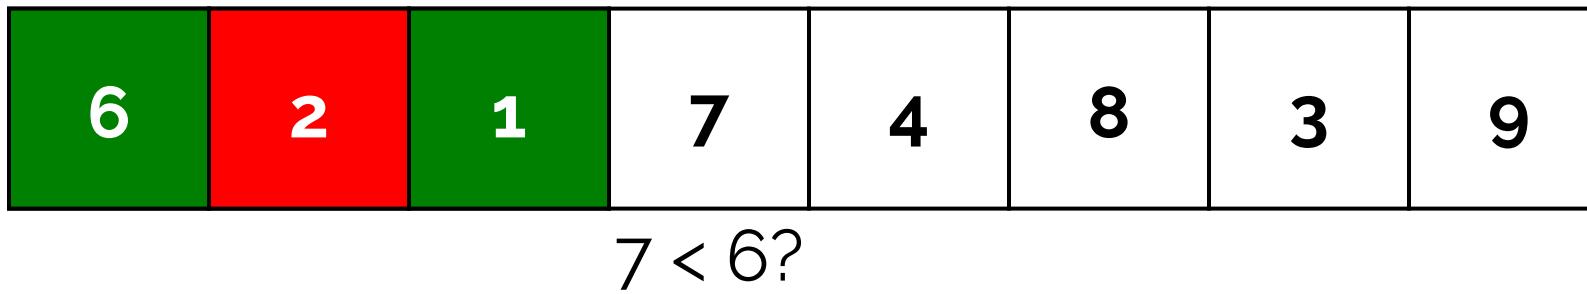


$1 < 6?$

Prediction: Taken

Simple branch predictor - unsorted array

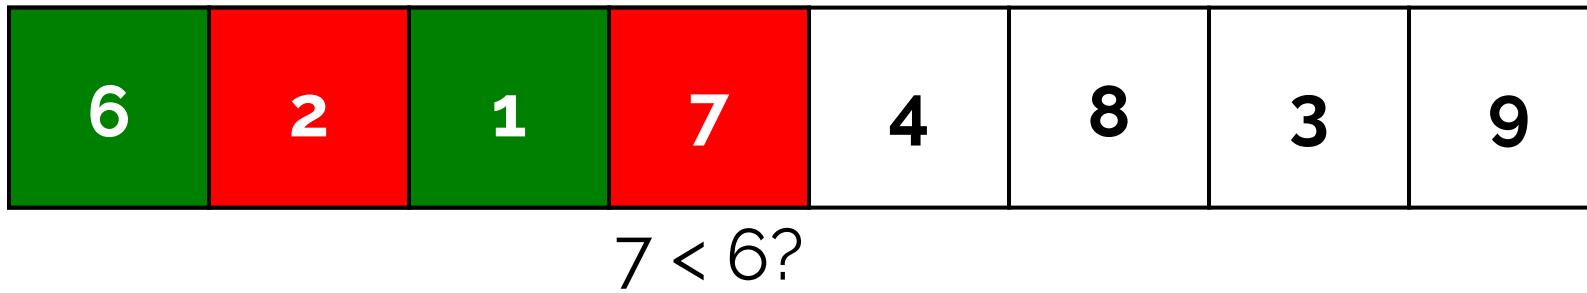
```
if (data[i] < 6) {  
    ...  
}
```



Prediction: Taken

Simple branch predictor - unsorted array

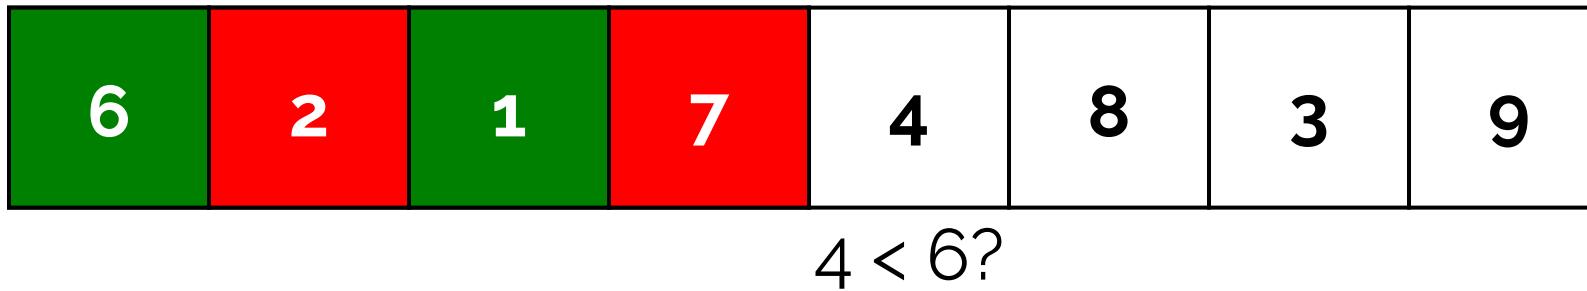
```
if (data[i] < 6) {  
    ...  
}
```



Prediction: Not taken

Simple branch predictor - unsorted array

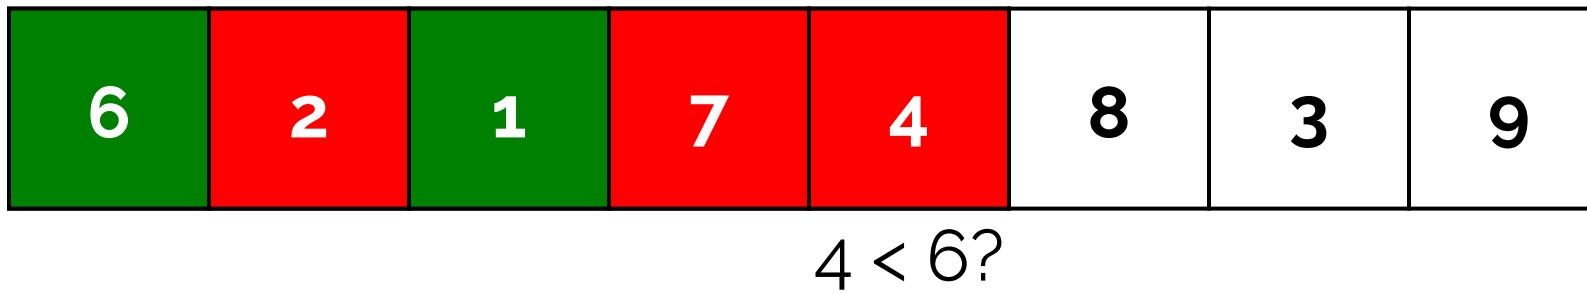
```
if (data[i] < 6) {  
    ...  
}
```



Prediction: Not taken

Simple branch predictor - unsorted array

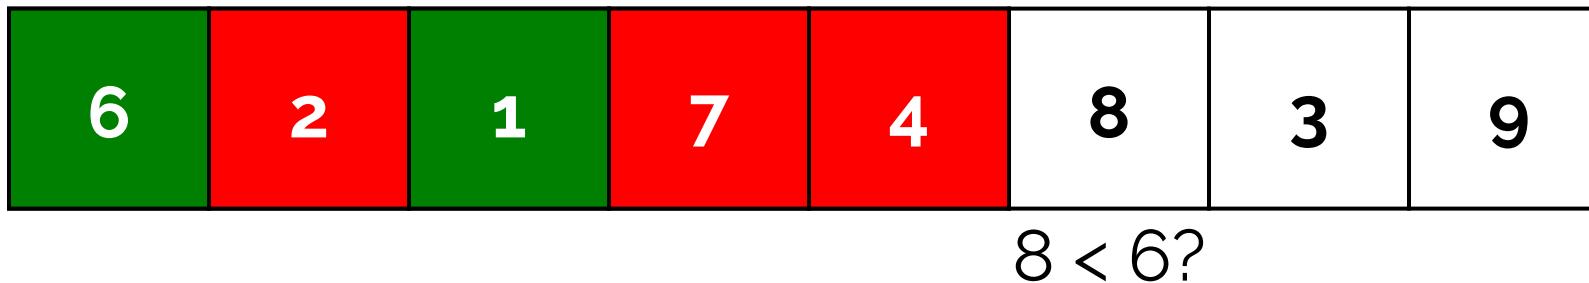
```
if (data[i] < 6) {  
    ...  
}
```



Prediction: Taken

Simple branch predictor - unsorted array

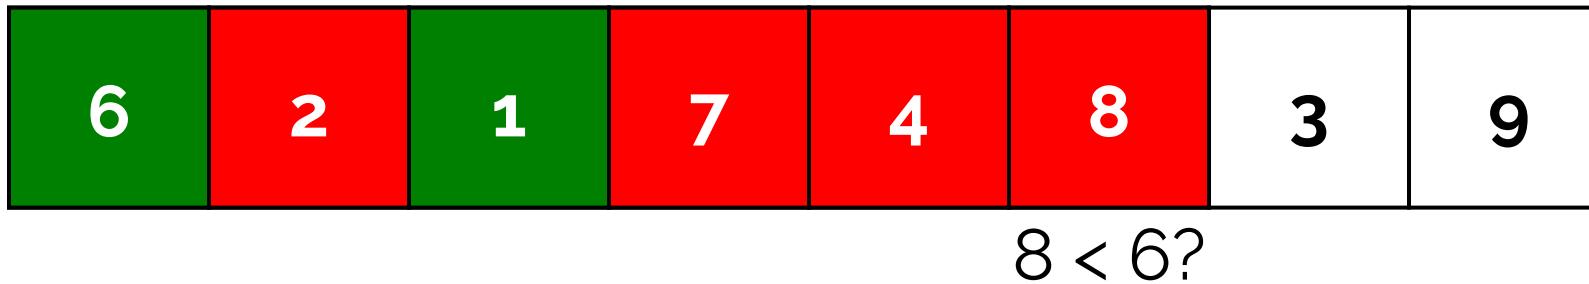
```
if (data[i] < 6) {  
    ...  
}
```



Prediction: Taken

Simple branch predictor - unsorted array

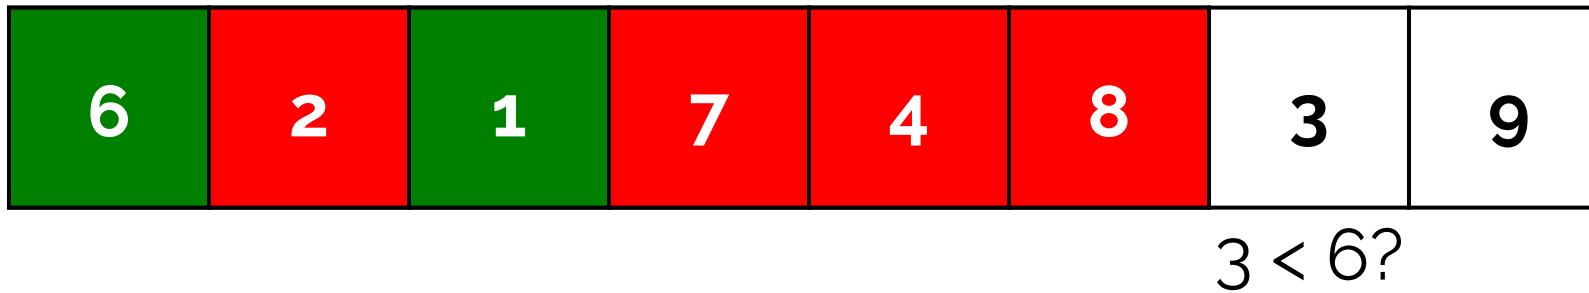
```
if (data[i] < 6) {  
    ...  
}
```



Prediction: Not taken

Simple branch predictor - unsorted array

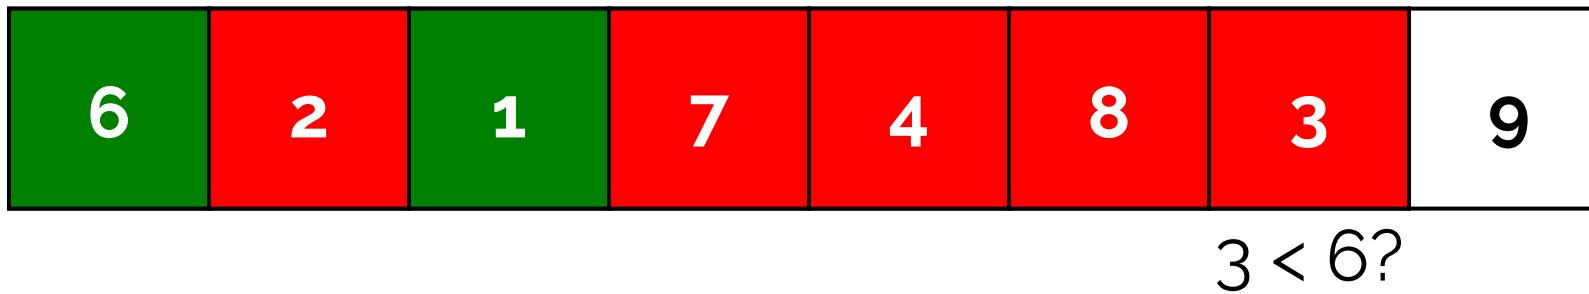
```
if (data[i] < 6) {  
    ...  
}
```



Prediction: Not taken

Simple branch predictor - unsorted array

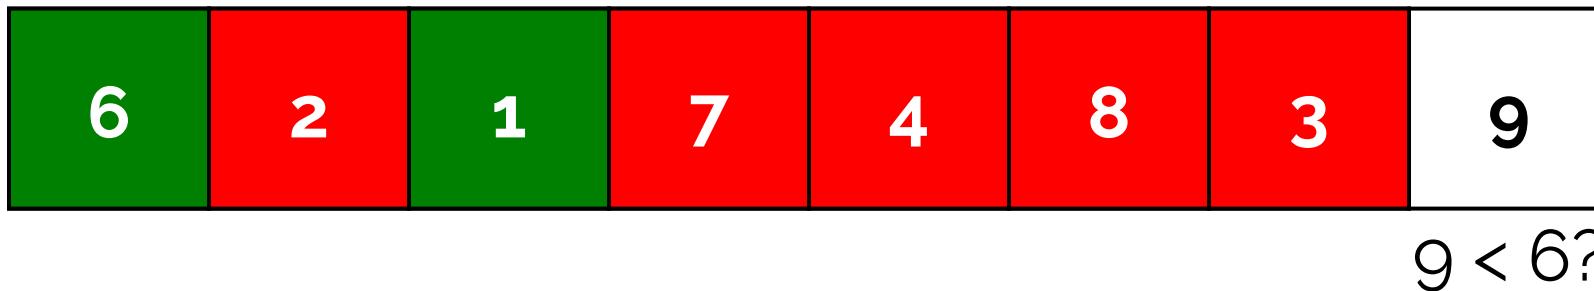
```
if (data[i] < 6) {  
    ...  
}
```



Prediction: Taken

Simple branch predictor - unsorted array

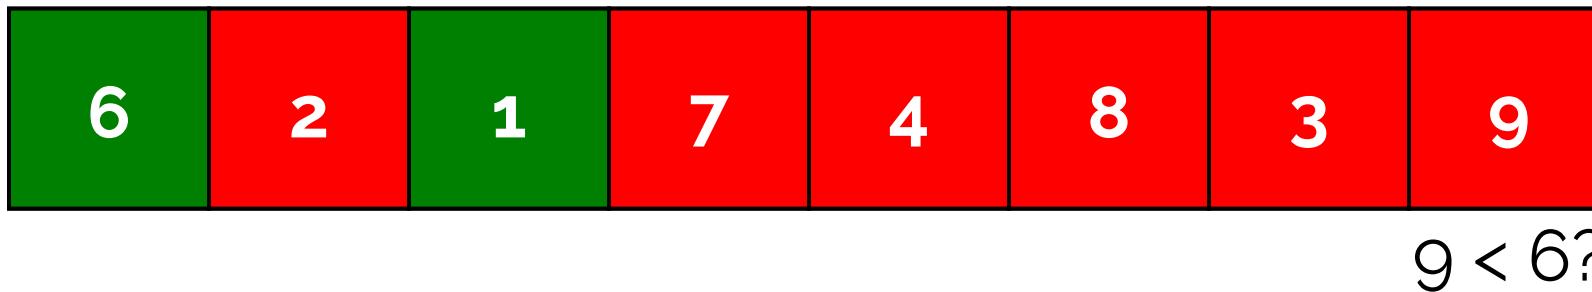
```
if (data[i] < 6) {  
    ...  
}
```



Prediction: Taken

Simple branch predictor - unsorted array

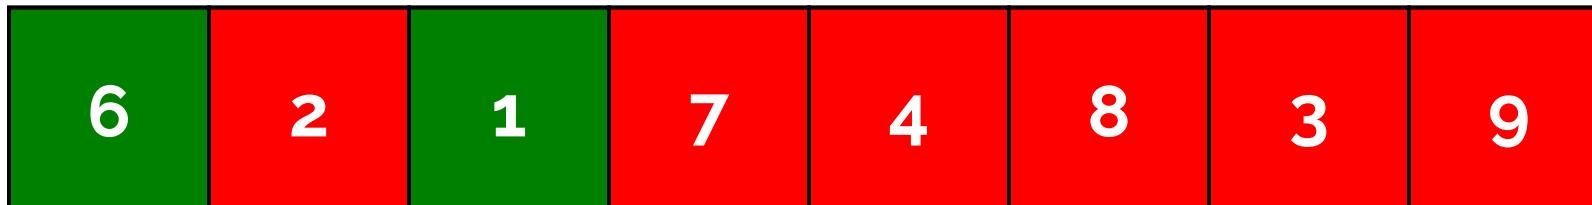
```
if (data[i] < 6) {  
    ...  
}
```



Prediction: Not taken

Simple branch predictor - unsorted array

```
if (data[i] < 6) {  
    ...  
}
```



Prediction: Not taken

2 hits, 6 misses (25% hit rate)



Simple branch predictor - sorted array

```
if (data[i] < 6) {  
    ...  
}
```

1	2	3	4	6	7	8	9
---	---	---	---	---	---	---	---

Prediction: Not taken

Simple branch predictor - sorted array

```
if (data[i] < 6) {  
    ...  
}
```

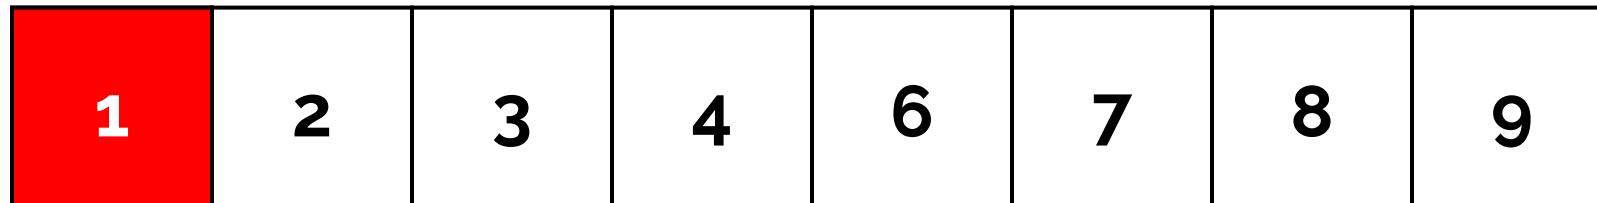
1	2	3	4	6	7	8	9
---	---	---	---	---	---	---	---

1 < 6?

Prediction: Not taken

Simple branch predictor - sorted array

```
if (data[i] < 6) {  
    ...  
}
```

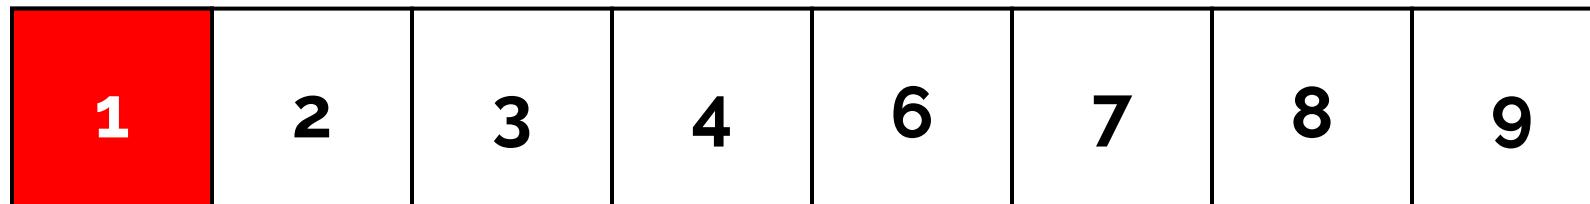


1 < 6?

Prediction: Taken

Simple branch predictor - sorted array

```
if (data[i] < 6) {  
    ...  
}
```



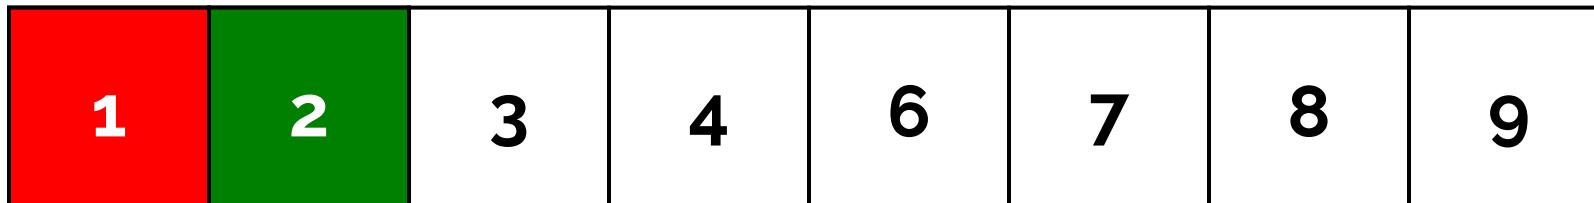
$2 < 6?$

Prediction: Taken



Simple branch predictor - sorted array

```
if (data[i] < 6) {  
    ...  
}
```

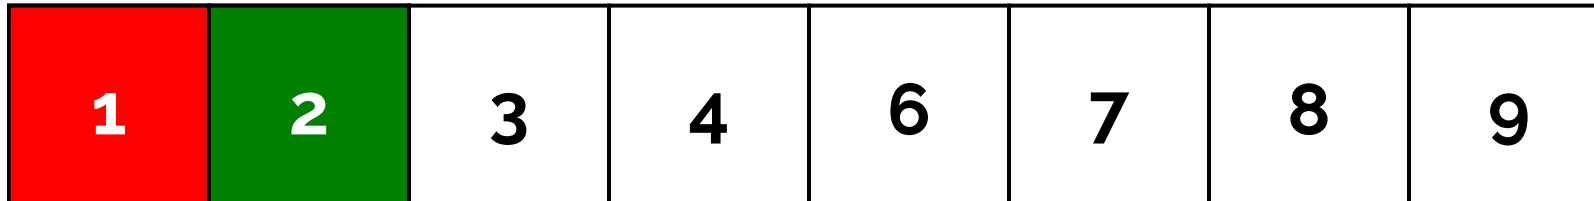


$2 < 6?$

Prediction: Taken

Simple branch predictor - sorted array

```
if (data[i] < 6) {  
    ...  
}
```

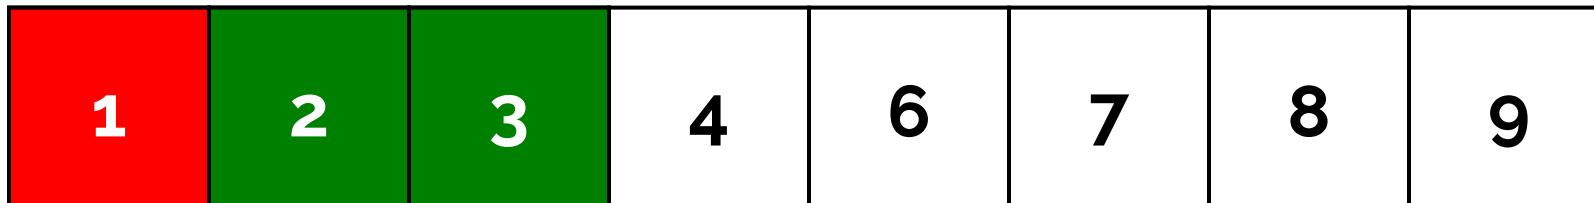


3 < 6?

Prediction: Taken

Simple branch predictor - sorted array

```
if (data[i] < 6) {  
    ...  
}
```



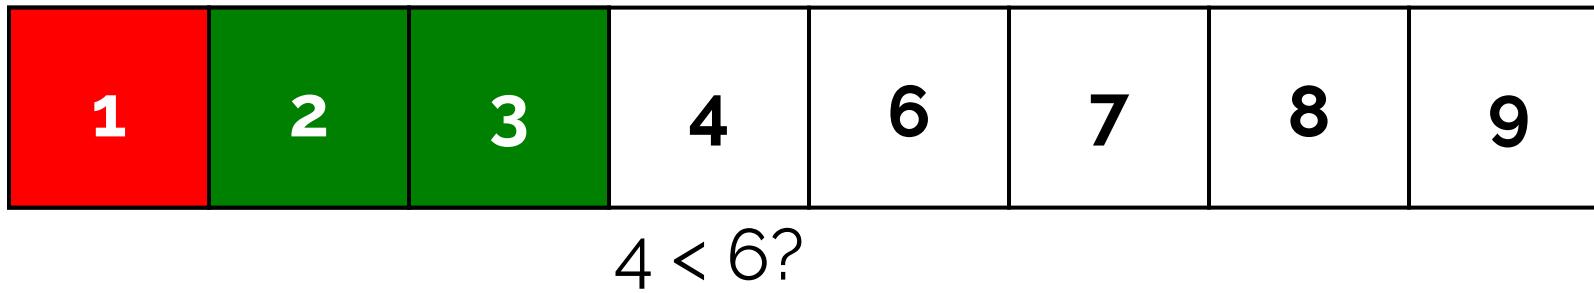
$3 < 6?$

Prediction: Taken



Simple branch predictor - sorted array

```
if (data[i] < 6) {  
    ...  
}
```

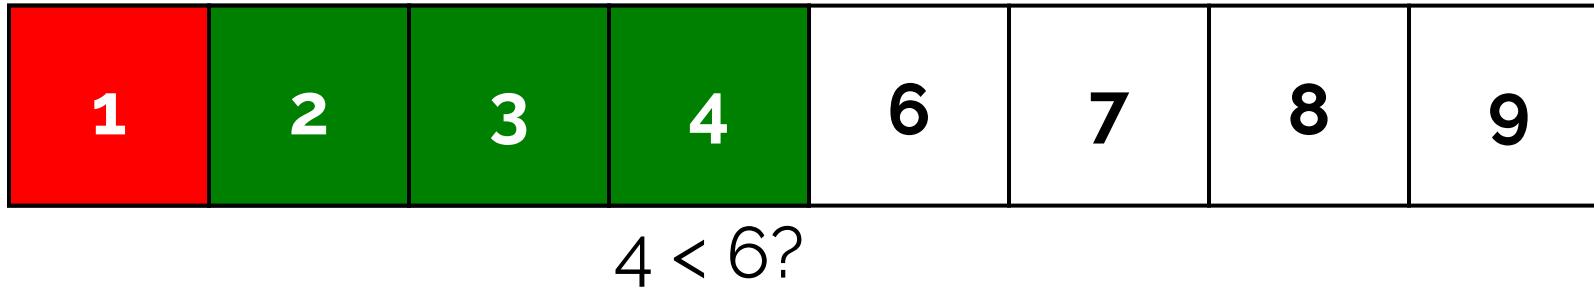


Prediction: Taken



Simple branch predictor - sorted array

```
if (data[i] < 6) {  
    ...  
}
```

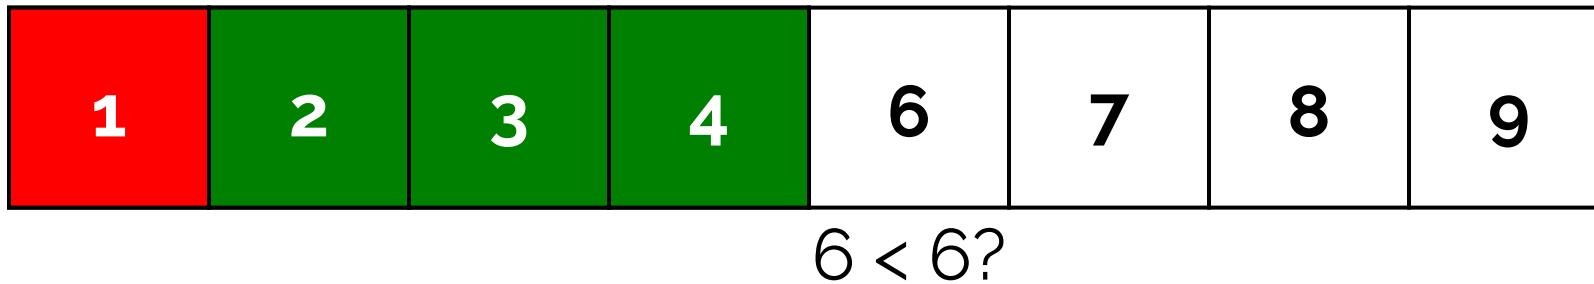


Prediction: Taken



Simple branch predictor - sorted array

```
if (data[i] < 6) {  
    ...  
}
```

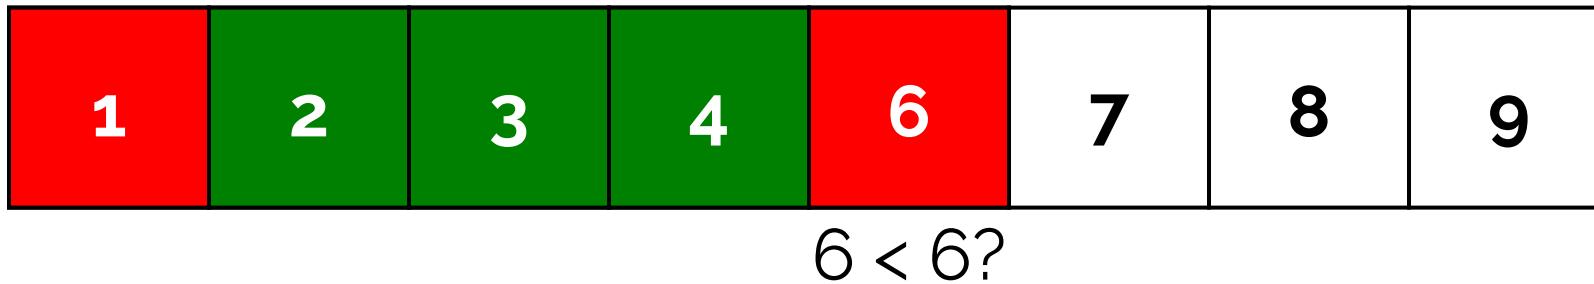


Prediction: Taken



Simple branch predictor - sorted array

```
if (data[i] < 6) {  
    ...  
}
```

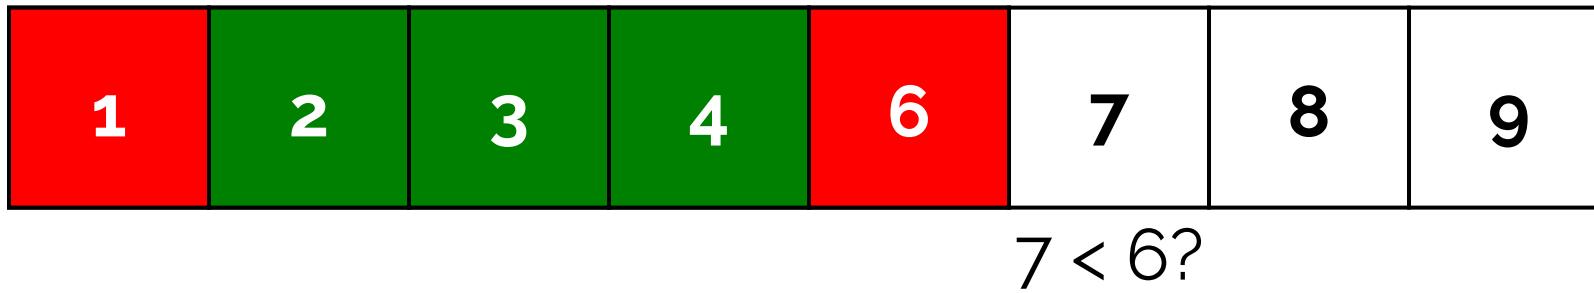


Prediction: Not taken



Simple branch predictor - sorted array

```
if (data[i] < 6) {  
    ...  
}
```

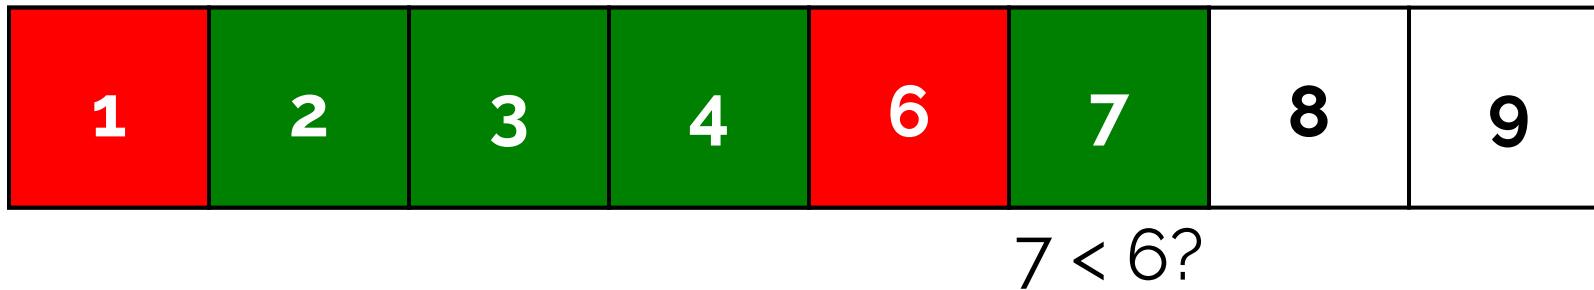


Prediction: Not taken



Simple branch predictor - sorted array

```
if (data[i] < 6) {  
    ...  
}
```

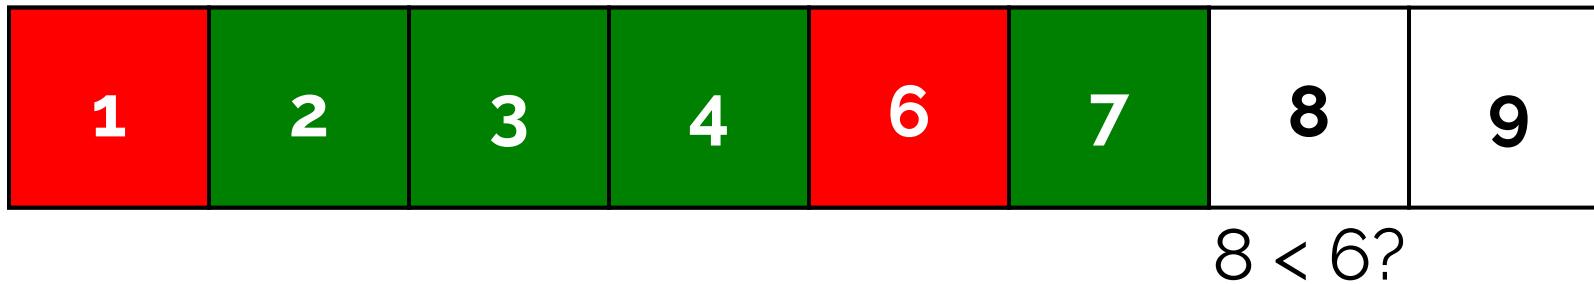


Prediction: Not taken



Simple branch predictor - sorted array

```
if (data[i] < 6) {  
    ...  
}
```

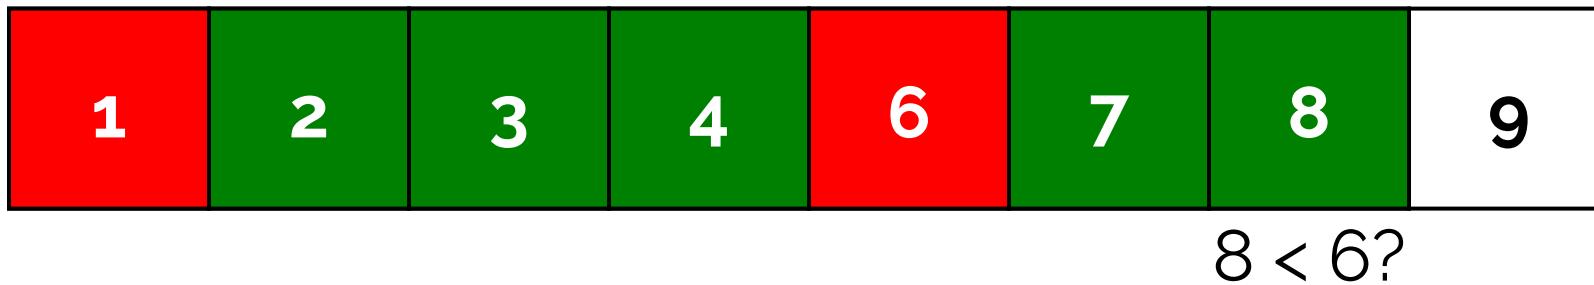


Prediction: Not taken



Simple branch predictor - sorted array

```
if (data[i] < 6) {  
    ...  
}
```

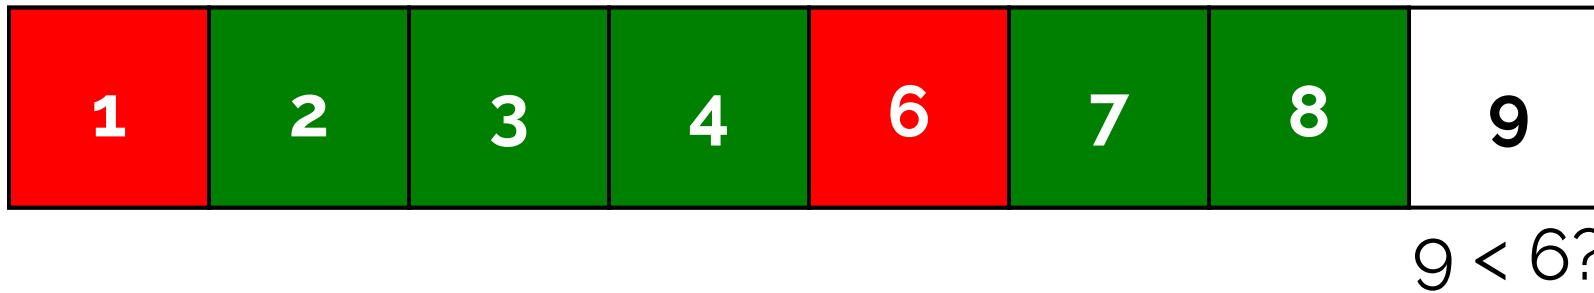


Prediction: Not taken



Simple branch predictor - sorted array

```
if (data[i] < 6) {  
    ...  
}
```

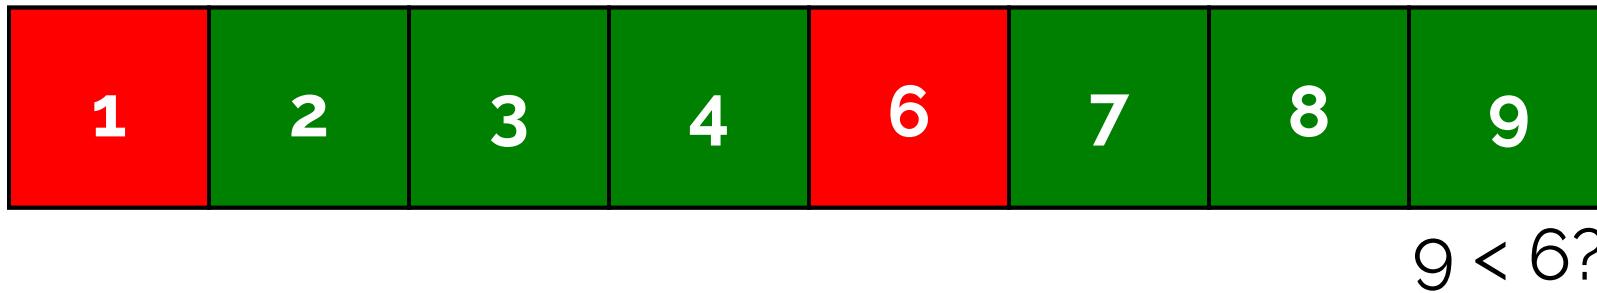


Prediction: Not taken



Simple branch predictor - sorted array

```
if (data[i] < 6) {  
    ...  
}
```

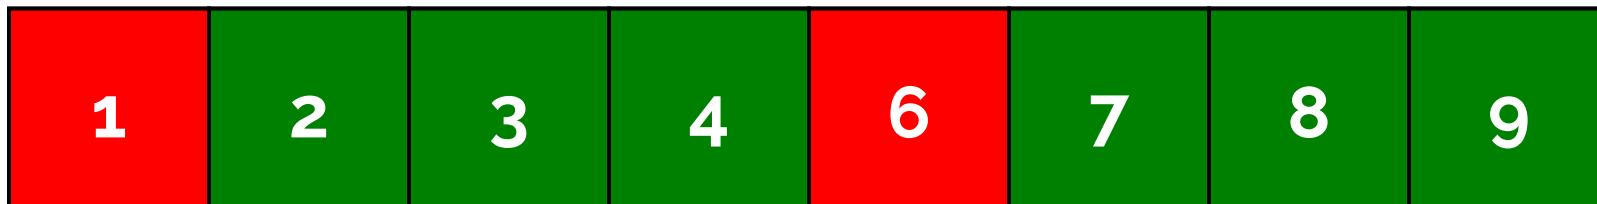


Prediction: Not taken



Simple branch predictor - sorted array

```
if (data[i] < 6) {  
    ...  
}
```



Prediction: Not taken

6 hits, 2 misses (75% hit rate)



How can the compiler help?

```
float sum = 0;
for (auto _ : state)
{
    for (auto x : data)
    {
        if (x < 6)
        {
            sum += x;
        }
    }
}
```

47	mov	rax, rcx
48	cmp	rcx, rdx
49	je	.L9
50	.L12:	
51	movss	xmm0, DWORD PTR [rax]
52	comiss	xmm2, xmm0
53	jbe	.L10
54	addss	xmm1, xmm0
55	.L10:	
56	add	rax, 4
57	cmp	rdx, rax
58	jne	.L12
59	.L9:	
60	sub	rbx, 1
61	jne	.L13
62	jmp	.L6

With float, there are two branches per iteration



How can the compiler help?

```
int sum = 0;
for (auto _ : state)
{
    for (auto x : data)
    {
        if (x < 6)
        {
            sum += x;
        }
    }
}
```

```
41 .L12:
42     mov     rax, r8
43     cmp     r8, rdi
44     je      .L9

45 .L11:
46     mov     edx, DWORD PTR [rax]
47     cmp     edx, 6
48     lea     ecx, [rbx+rdx]
49     cmovl  ebx, ecx
50     add     rax, 4
51     cmp     rdi, rax
52     jne     .L11

53 .L9:
54     sub     rbp, 1
55     jne     .L12
56     mov     rdi, r12
```

With `int`, one branch is removed (using `cmov`)



How to measure?

branch-misses

How many times was a branch mispredicted?



How to measure?

branch-misses

How many times was a branch mispredicted?

```
$ perf stat -e branch-misses ./example0a
with      sort ->      383 902
without    sort -> 101 652 009
```



How to help the branch predictor?

- More predictable data



How to help the branch predictor?

- More predictable data
- Profile-guided optimization



How to help the branch predictor?

- More predictable data
- Profile-guided optimization
- Remove (unpredictable) branches



How to help the branch predictor?

- More predictable data
- Profile-guided optimization
- Remove (unpredictable) branches
- Compiler hints (use with caution)

```
if (__builtin_expect(will_it_blend(), 0)) {  
    // this branch is not likely to be taken  
}
```



Branch target prediction

- Target of a jump is not known at compile time:



Branch target prediction

- Target of a jump is not known at compile time:
 - Function pointer



Branch target prediction

- Target of a jump is not known at compile time:
 - Function pointer
 - Function return address



Branch target prediction

- Target of a jump is not known at compile time:
 - Function pointer
 - Function return address
 - Virtual method



Code (backup)

```
struct A { virtual void handle(size_t* data) const = 0; };
struct B: public A { void handle(size_t* data) const final { *data += 1; } };
struct C: public A { void handle(size_t* data) const final { *data += 2; } };

std::vector<std::unique_ptr<A>> data = /* 4K random B/C instances */;
// std::sort(data.begin(), data.end(), /* sort by instance type */);
size_t sum = 0;
for (auto& x : data)
{
    x->handle(&sum);
}
```



Result (backup)

Benchmark	Time	CPU	Iterations
handle_nosort/4096	23350 ns	23349 ns	30734
handle_sorted/4096	7448 ns	7448 ns	86814



perf (backup)

```
$ perf stat -e branch-misses ./example0b
with sort  ->    337 274
without sort ->  84 183 161
```



Code (backup)

```
// Addresses of N integers, each `offset` bytes apart
std::vector<int*> data = ...;
for (auto ptr: data)
{
    *ptr += 1;
}
// Offsets: 4, 64, 4000, 4096, 4128
```

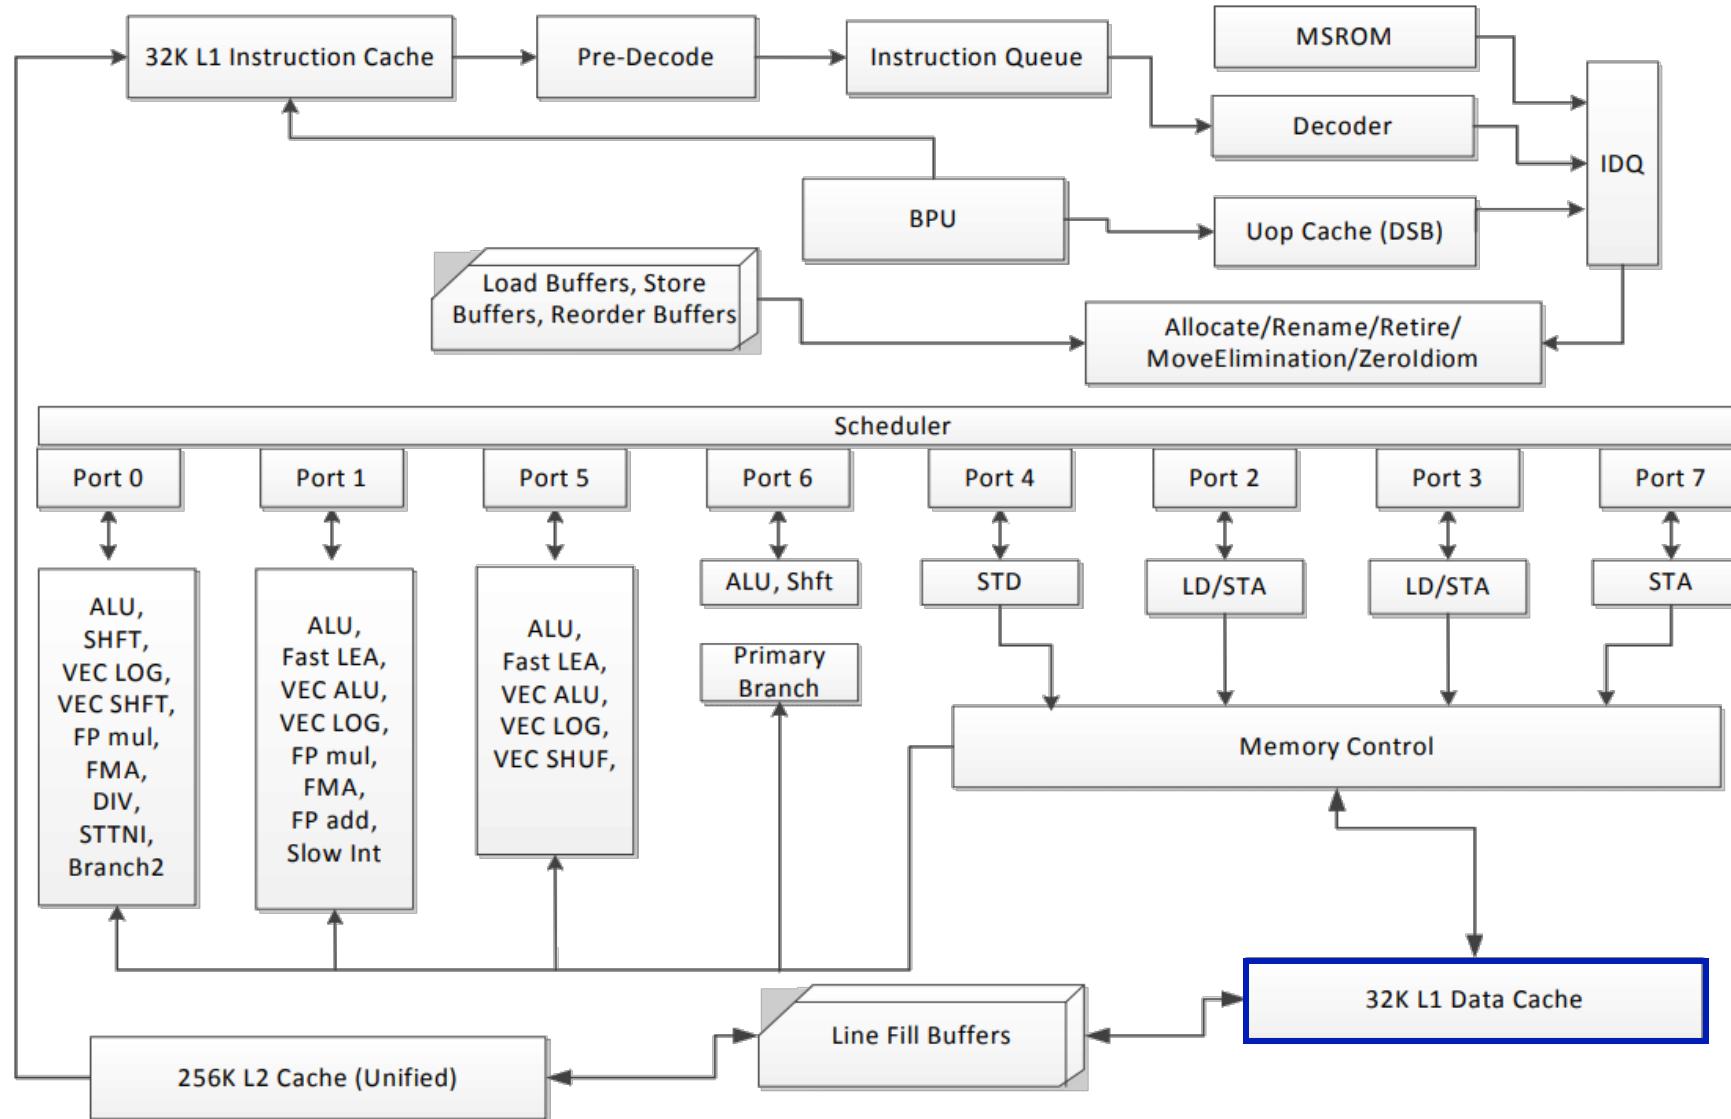


Result (backup)

Benchmark	Time	CPU	Iterations
<hr/>			
offset4/7	2.83 ns	2.83 ns	255750233
offset4/8	3.03 ns	3.02 ns	177109965
offset4/9	3.76 ns	3.75 ns	157739295
offset4/10	4.62 ns	4.61 ns	177899906
offset4/11	4.93 ns	4.92 ns	162959140
offset64/7	3.19 ns	3.18 ns	179723151
offset64/8	3.83 ns	3.65 ns	216288609
offset64/9	3.74 ns	3.74 ns	201008685
offset64/10	4.41 ns	4.40 ns	159949703
offset64/11	4.41 ns	4.41 ns	128933855
offset4000/7	3.69 ns	3.69 ns	187745245
offset4000/8	3.27 ns	3.26 ns	226401022
offset4000/9	4.19 ns	4.18 ns	157866983
offset4000/10	4.49 ns	4.48 ns	173084452
offset4000/11	4.53 ns	4.52 ns	128906229
offset4096/7	9.05 ns	9.05 ns	78087527
offset4096/8	10.4 ns	10.4 ns	67550724
offset4096/9	18.7 ns	18.7 ns	38875870
offset4096/10	25.5 ns	25.5 ns	26893946
offset4096/11	32.7 ns	32.7 ns	21369400
offset4128/7	3.23 ns	3.22 ns	250263727
offset4128/8	3.13 ns	3.13 ns	218371877
offset4128/9	3.75 ns	3.71 ns	157448182
offset4128/10	4.28 ns	4.25 ns	144839049
offset4128/11	5.47 ns	5.44 ns	128547528



Cache memory



How are (L1) caches implemented

- N-way set associative table
 - Hardware hash table



How are (L1) caches implemented

- N-way set associative table
 - Hardware hash table
 - Key = address (8B)

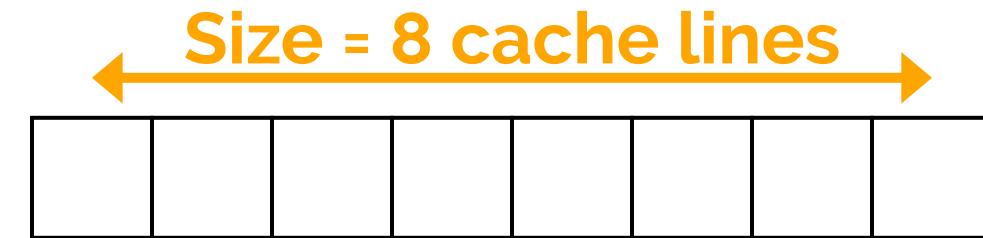


How are (L1) caches implemented

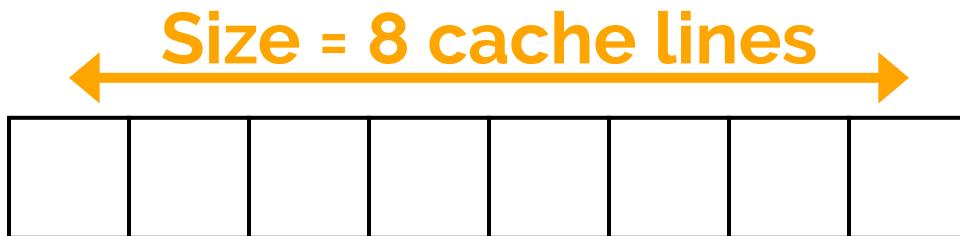
- N-way set associative table
 - Hardware hash table
- Key = address (8B)
- Entry = cache line (64B)



N-way set associative cache



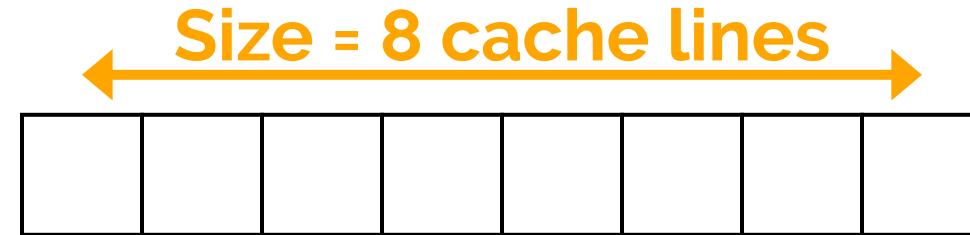
N-way set associative cache



Associativity (N) - # of cache lines per bucket



N-way set associative cache

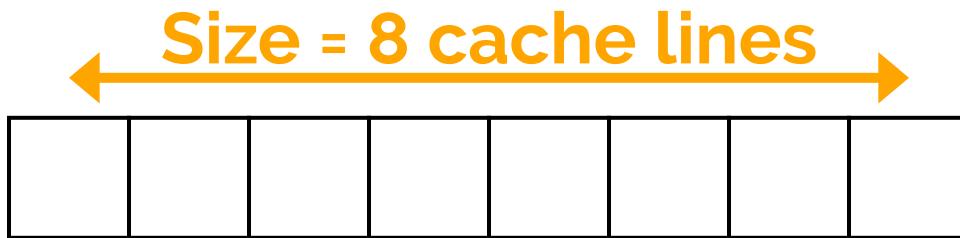


Associativity (N) - # of cache lines per bucket

$$\# \text{ of buckets} = \text{Size} / N$$

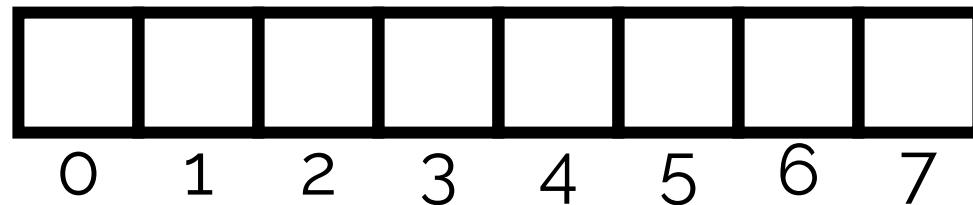


N-way set associative cache

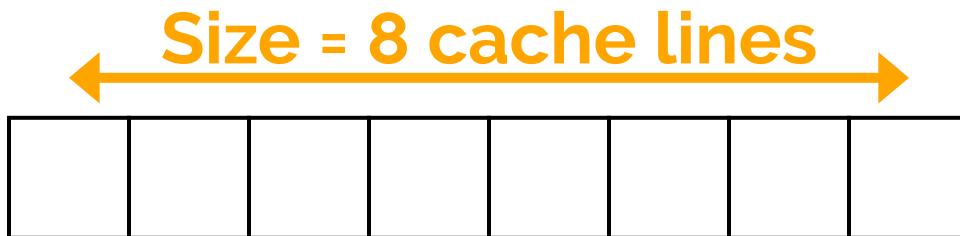


Associativity (N) - # of cache lines per bucket
 $\# \text{ of buckets} = \text{Size} / N$

$N = 1$ (direct mapped)

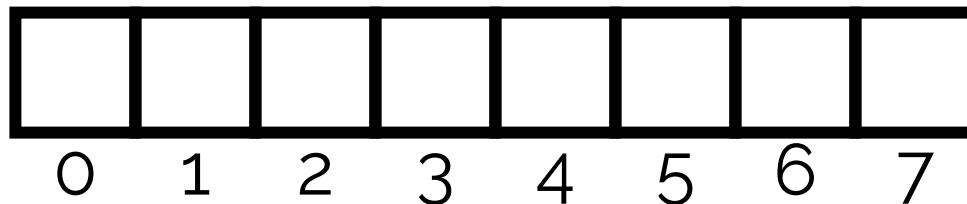


N-way set associative cache

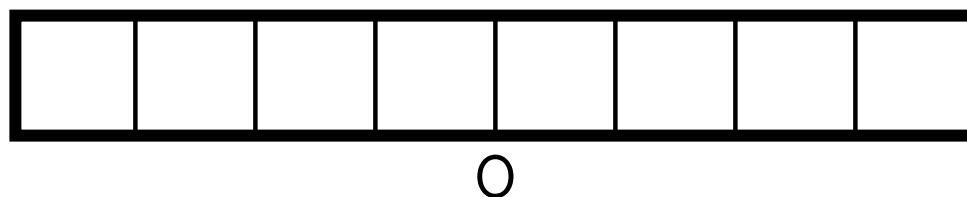


Associativity (N) - # of cache lines per bucket
 $\# \text{ of buckets} = \text{Size} / N$

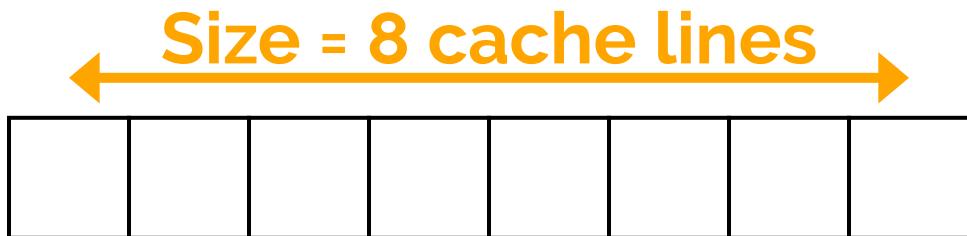
$N = 1$ (direct mapped)



$N = 8$ (fully associative)

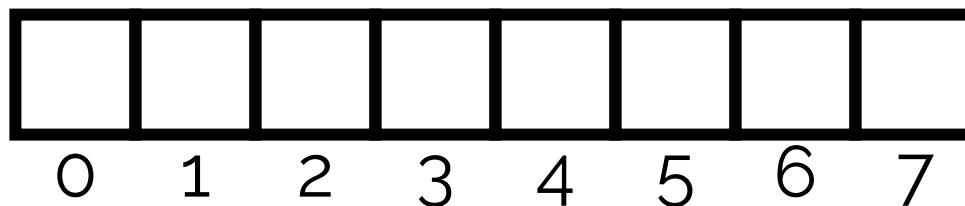


N-way set associative cache

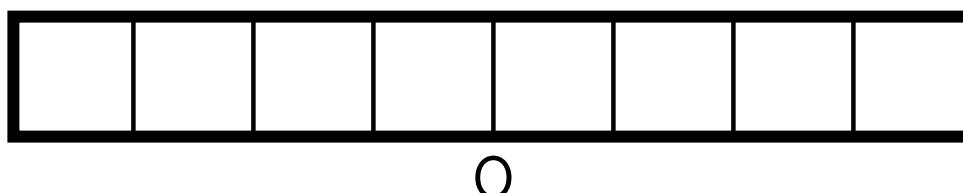


Associativity (N) - # of cache lines per bucket
 $\# \text{ of buckets} = \text{Size} / N$

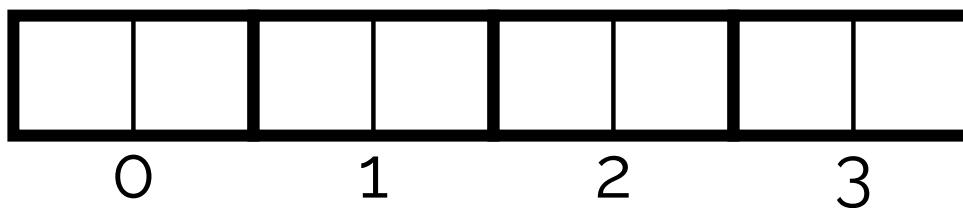
$N = 1$ (direct mapped)



$N = 8$ (fully associative)



$N = 2$



How are addresses hashed?



How are addresses hashed?



- **Offset**

- Selects byte within a cache line
- $\log_2(\text{cache line size})$ bits



How are addresses hashed?



- **Offset**

- Selects byte within a cache line
- $\log_2(\text{cache line size})$ bits

- **Index**

- Selects bucket within the cache
- $\log_2(\text{bucket count})$ bits



How are addresses hashed?

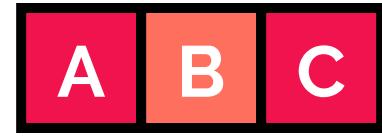


- **Offset**
 - Selects byte within a cache line
 - $\log_2(\text{cache line size})$ bits
- **Index**
 - Selects bucket within the cache
 - $\log_2(\text{bucket count})$ bits
- **Tag**
 - Used for matching



N-way set associative cache

Cache lines:



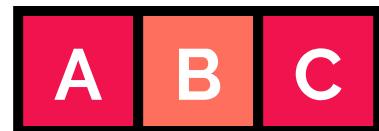
Index bits:

0 1 0



N-way set associative cache

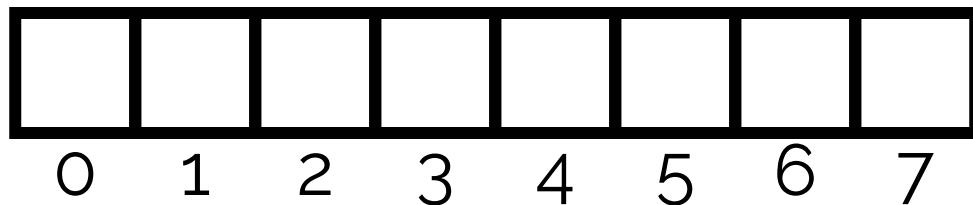
Cache lines:



Index bits:

0 1 0

$N = 1$



N-way set associative cache

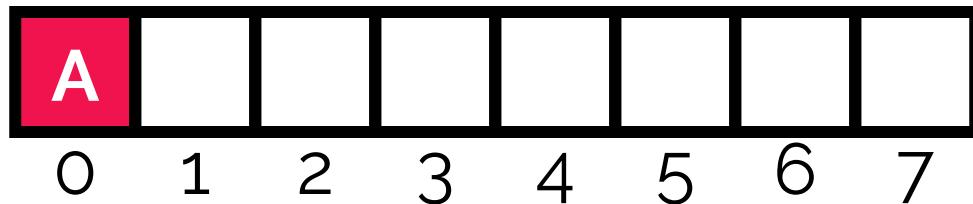
Cache lines:



Index bits:

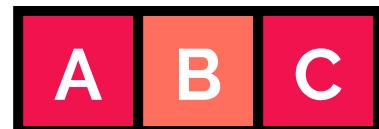
0 1 0

$N = 1$



N-way set associative cache

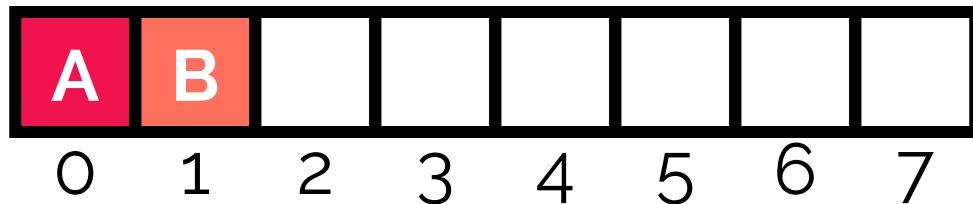
Cache lines:



Index bits:

0 1 0

$N = 1$



N-way set associative cache

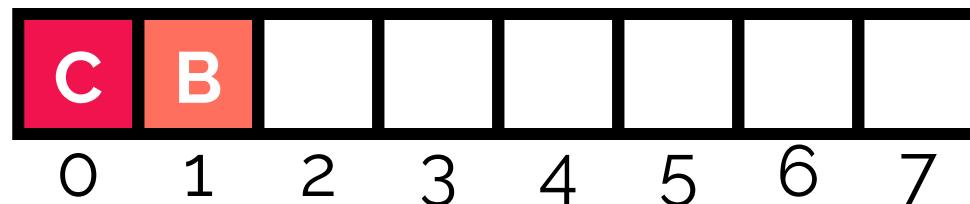
Cache lines:



Index bits:

0 1 0

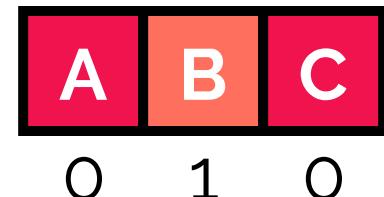
$N = 1$



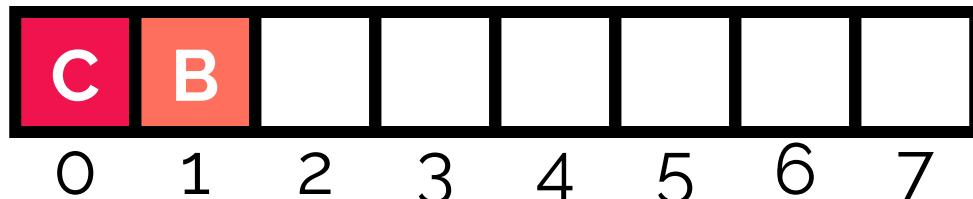
N-way set associative cache

Cache lines:

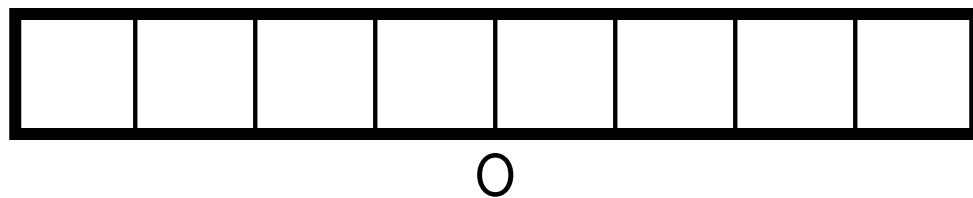
Index bits:



$N = 1$



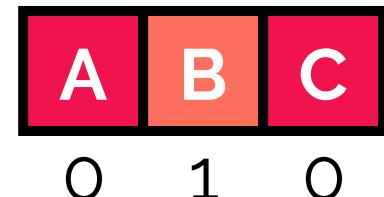
$N = 8$



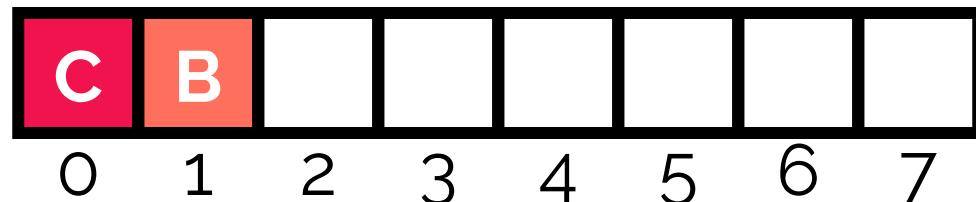
N-way set associative cache

Cache lines:

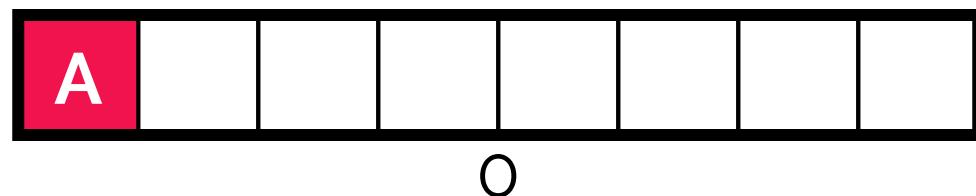
Index bits:



$N = 1$



$N = 8$



N-way set associative cache

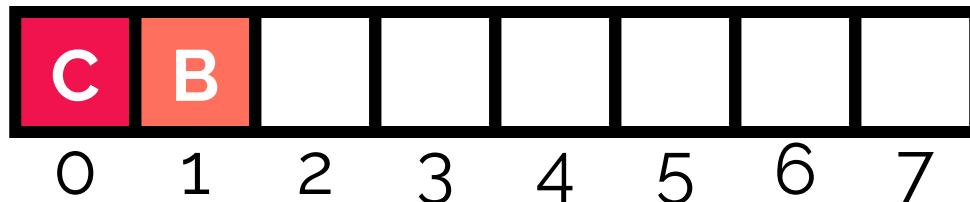
Cache lines:



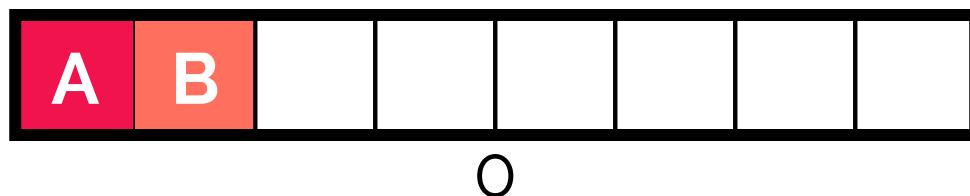
Index bits:

0 1 0

$N = 1$

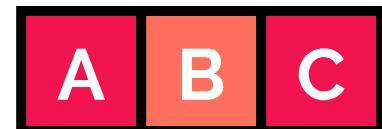


$N = 8$



N-way set associative cache

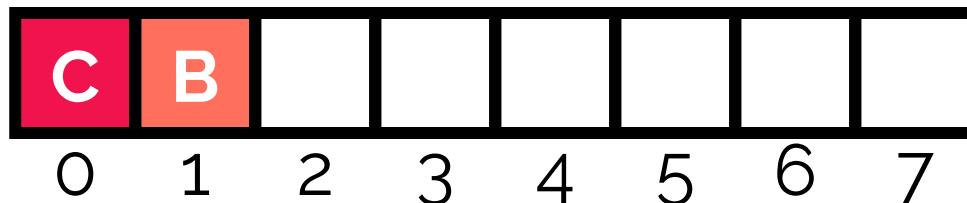
Cache lines:



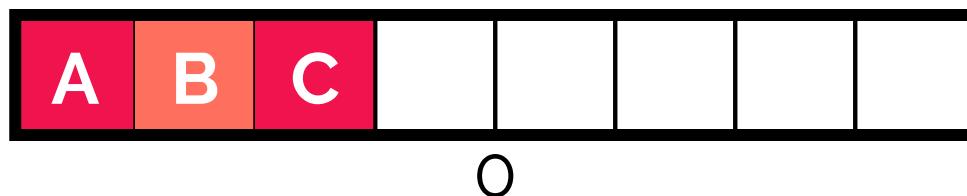
Index bits:

0 1 0

$N = 1$



$N = 8$



N-way set associative cache

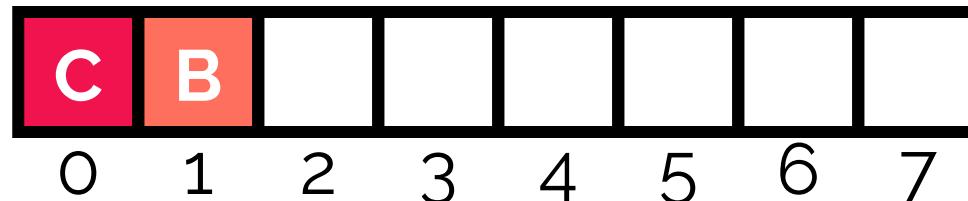
Cache lines:



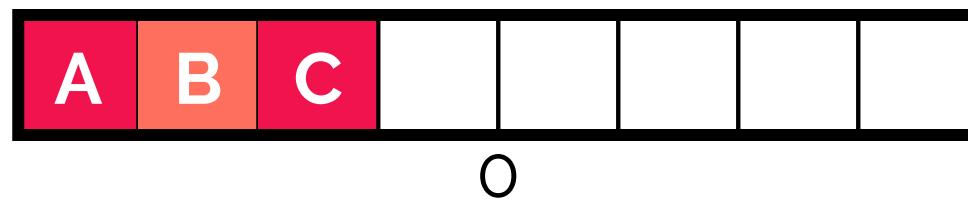
Index bits:

0 1 0

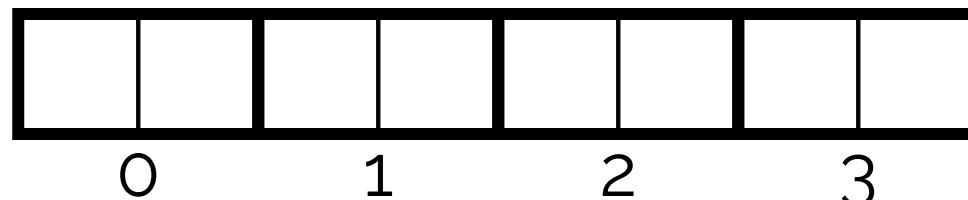
$N = 1$



$N = 8$



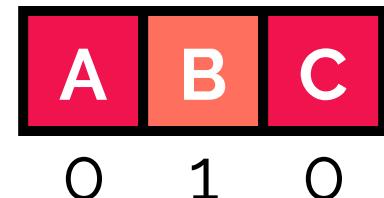
$N = 2$



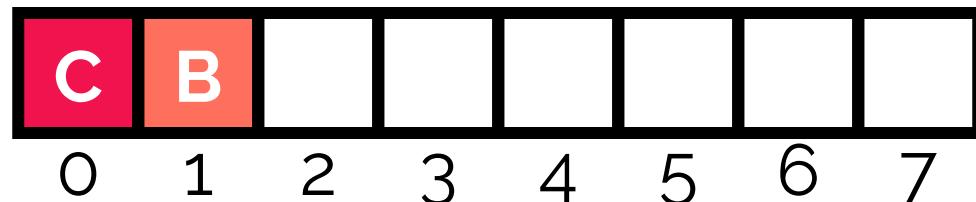
N-way set associative cache

Cache lines:

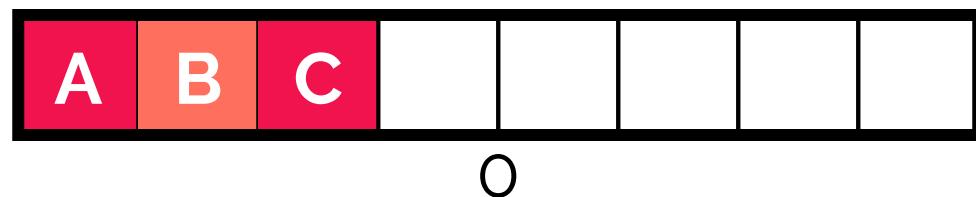
Index bits:



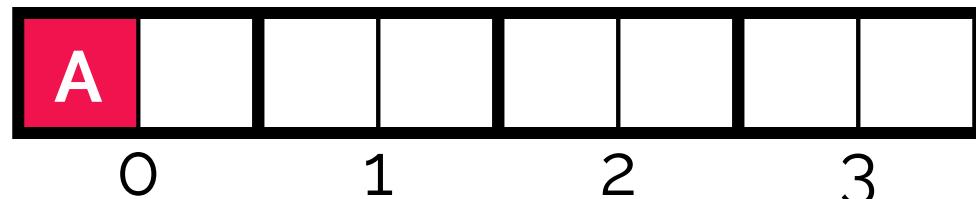
$N = 1$



$N = 8$



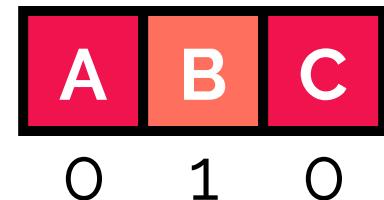
$N = 2$



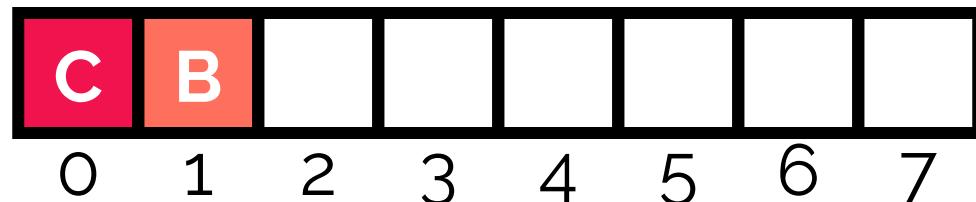
N-way set associative cache

Cache lines:

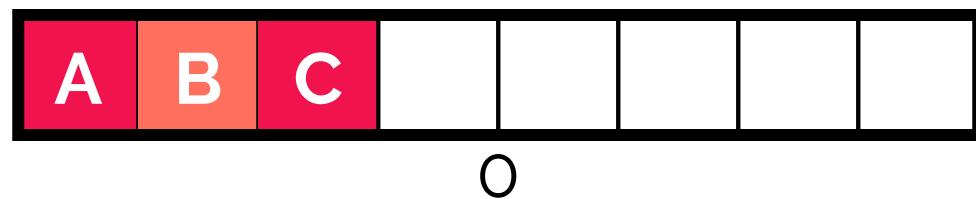
Index bits:



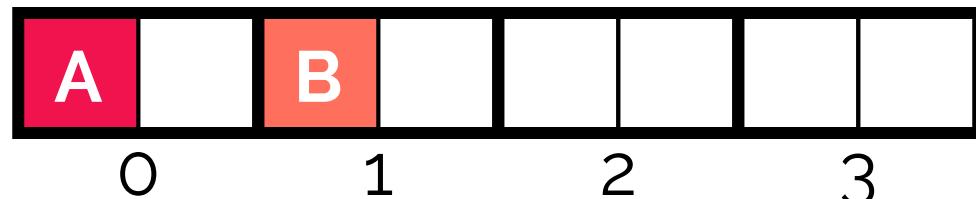
$N = 1$



$N = 8$



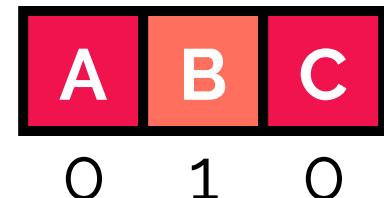
$N = 2$



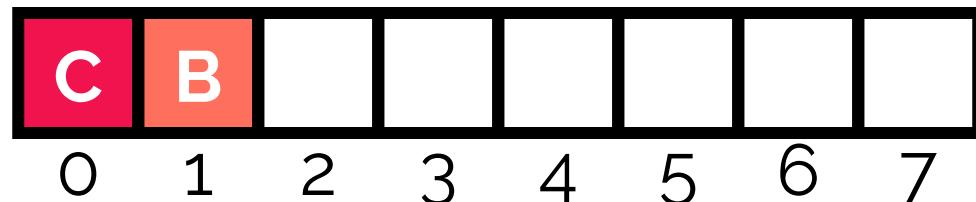
N-way set associative cache

Cache lines:

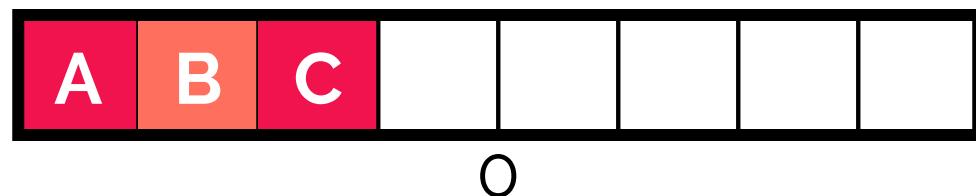
Index bits:



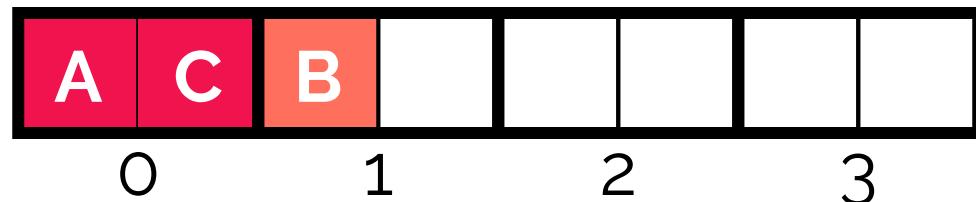
$N = 1$



$N = 8$



$N = 2$



Intel L1 cache

```
$ getconf -a | grep LEVEL1_DCACHE  
LEVEL1_DCACHE_SIZE          32768  
LEVEL1_DCACHE_ASSOC          8  
LEVEL1_DCACHE_LINESIZE      64
```



Intel L1 cache

```
$ getconf -a | grep LEVEL1_DCACHE  
LEVEL1_DCACHE_SIZE          32768  
LEVEL1_DCACHE_ASSOC          8  
LEVEL1_DCACHE_LINESIZE      64
```

- **Cache line size** - 64 B (6 offset bits)



Intel L1 cache

```
$ getconf -a | grep LEVEL1_DCACHE  
LEVEL1_DCACHE_SIZE          32768  
LEVEL1_DCACHE_ASSOC          8  
LEVEL1_DCACHE_LINESIZE      64
```

- **Cache line size** - 64 B (6 offset bits)
- **Associativity (N)** - 8



Intel L1 cache

```
$ getconf -a | grep LEVEL1_DCACHE  
LEVEL1_DCACHE_SIZE          32768  
LEVEL1_DCACHE_ASSOC          8  
LEVEL1_DCACHE_LINESIZE      64
```

- **Cache line size** - 64 B (6 offset bits)
- **Associativity (N)** - 8
- **Size** - 32768 B



Intel L1 cache

```
$ getconf -a | grep LEVEL1_DCACHE  
LEVEL1_DCACHE_SIZE          32768  
LEVEL1_DCACHE_ASSOC          8  
LEVEL1_DCACHE_LINESIZE      64
```

- **Cache line size** - 64 B (6 offset bits)
- **Associativity (N)** - 8
- **Size** - 32768 B
- $32768 / 64 \Rightarrow 512$ cache lines



Intel L1 cache

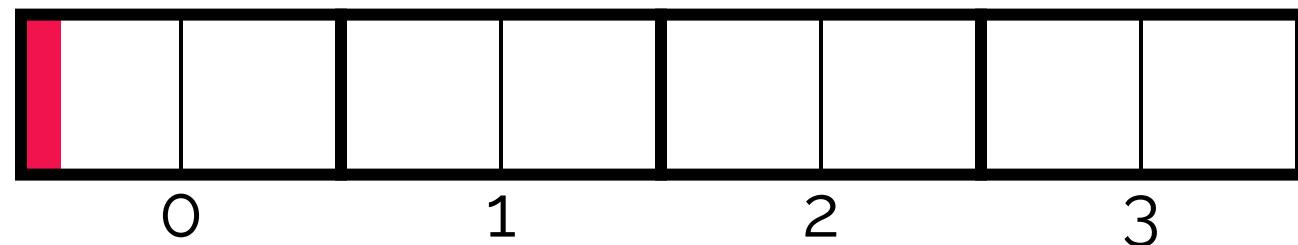
```
$ getconf -a | grep LEVEL1_DCACHE  
LEVEL1_DCACHE_SIZE          32768  
LEVEL1_DCACHE_ASSOC          8  
LEVEL1_DCACHE_LINESIZE      64
```

- **Cache line size** - 64 B (6 offset bits)
- **Associativity (N)** - 8
- **Size** - 32768 B
- $32768 / 64 \Rightarrow 512$ cache lines
- $512 / 8 \Rightarrow 64$ buckets (6 index bits)



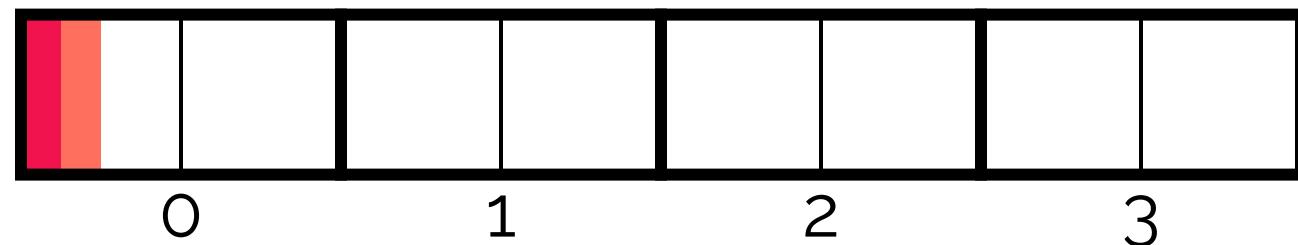
Offset = 4B

Number	Tag	Index	Offset
A	..100000	000000	000000



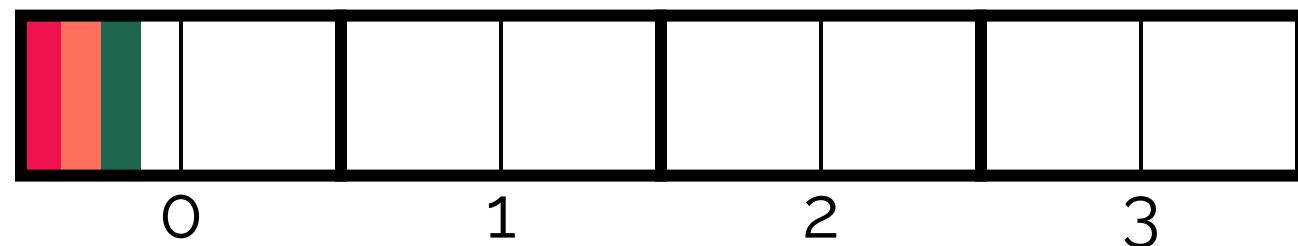
Offset = 4B

Number	Tag	Index	Offset
A	.100000	000000	000000
B	.100000	000000	000100



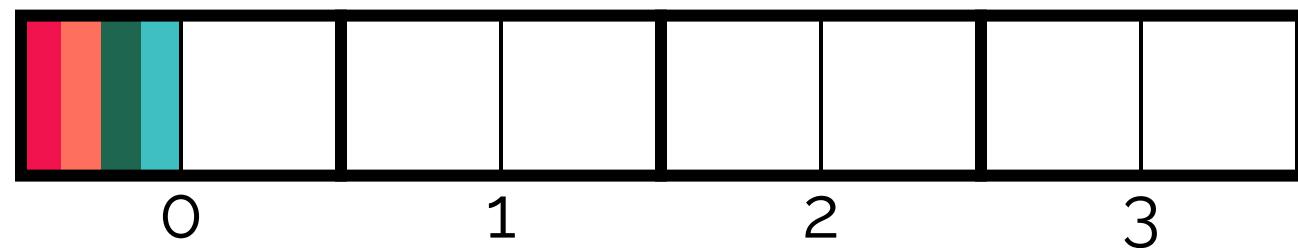
Offset = 4B

Number	Tag	Index	Offset
A	..100000	000000	000000
B	..100000	000000	000100
C	..100000	000000	001000



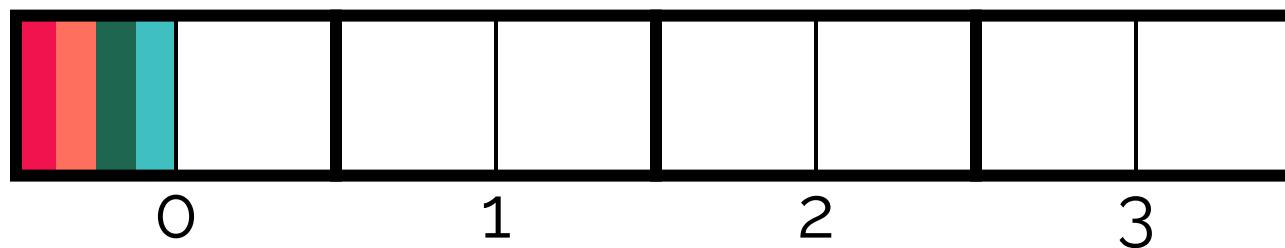
Offset = 4B

Number	Tag	Index	Offset
A	..100000	000000	000000
B	..100000	000000	000100
C	..100000	000000	001000
D	..100000	000000	001100



Offset = 4B

Number	Tag	Index	Offset
A	.100000	000000	000000
B	.100000	000000	000100
C	.100000	000000	001000
D	.100000	000000	001100

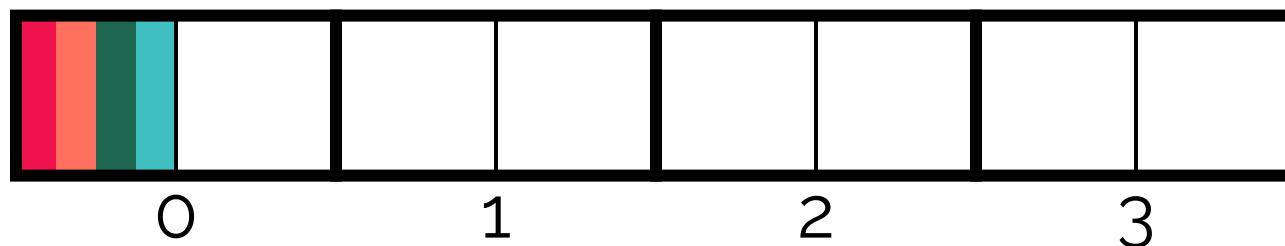


- Same bucket, same cache line for each number



Offset = 4B

Number	Tag	Index	Offset
A	.100000	000000	000000
B	.100000	000000	000100
C	.100000	000000	001000
D	.100000	000000	001100

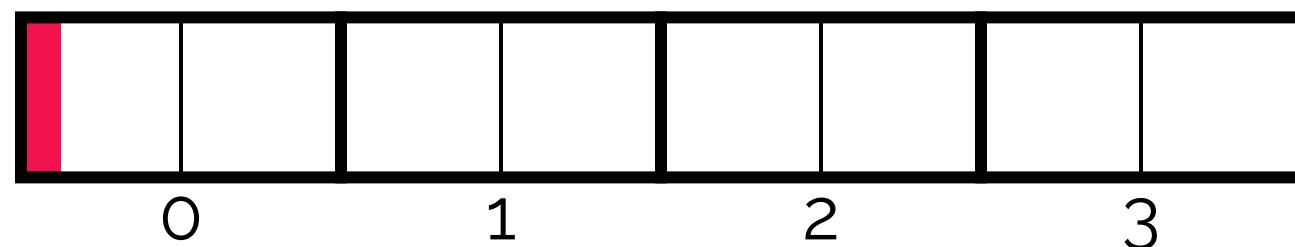


- Same bucket, same cache line for each number
- Most efficient, no space is wasted



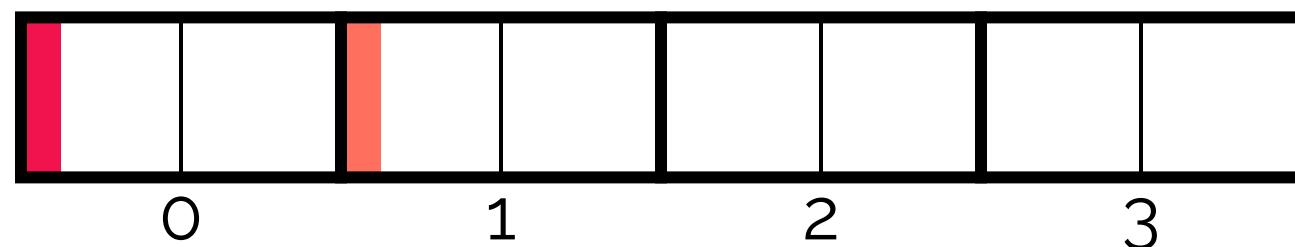
Offset = 64B

Number	Tag	Index	Offset
A	..100000	000000	000000



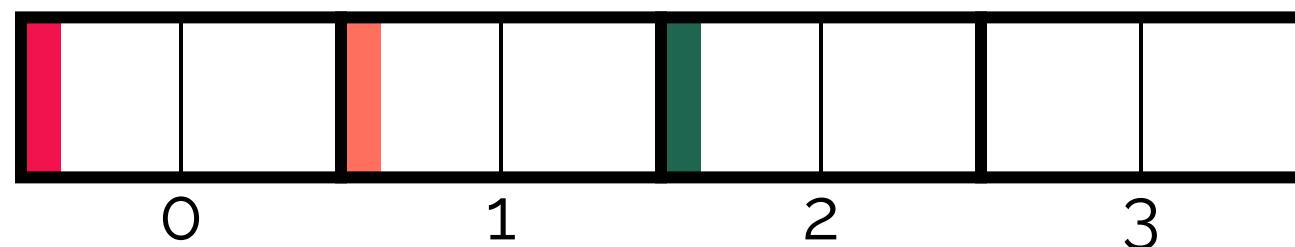
Offset = 64B

Number	Tag	Index	Offset
A	..100000	000000	000000
B	..100000	000001	000000



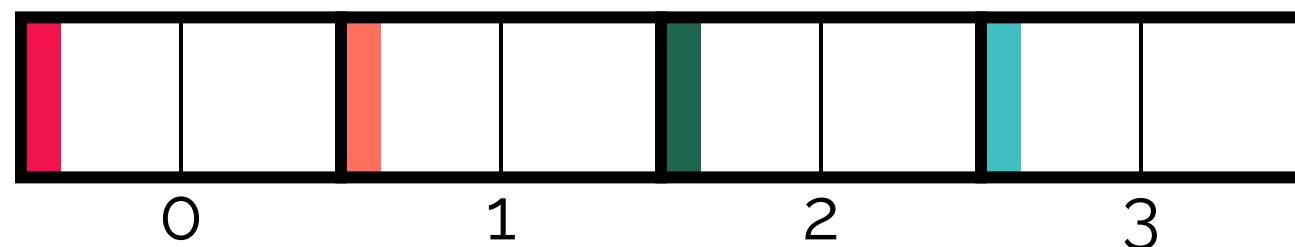
Offset = 64B

Number	Tag	Index	Offset
A	..100000	000000	000000
B	..100000	000001	000000
C	..100000	000010	000000



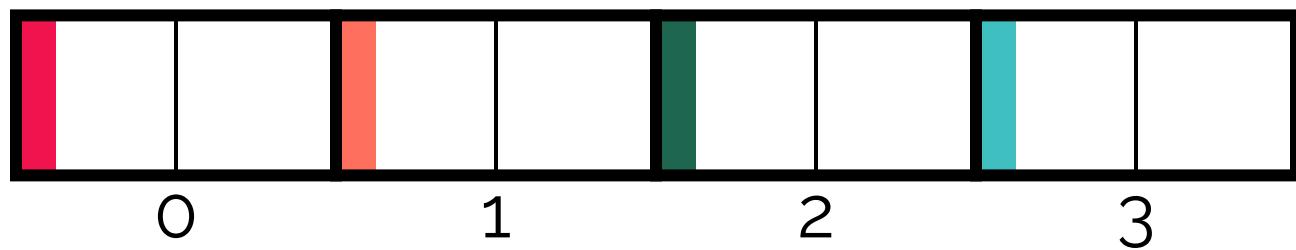
Offset = 64B

Number	Tag	Index	Offset
A	..100000	000000	000000
B	..100000	000001	000000
C	..100000	000010	000000
D	..100000	000011	000000



Offset = 64B

Number	Tag	Index	Offset
A	..100000	000000	000000
B	..100000	000001	000000
C	..100000	000010	000000
D	..100000	000011	000000

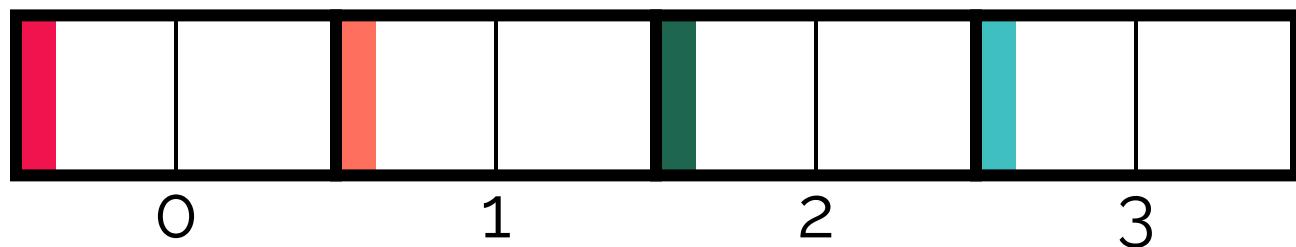


- Different bucket for each number



Offset = 64B

Number	Tag	Index	Offset
A	..100000	000000	000000
B	..100000	000001	000000
C	..100000	000010	000000
D	..100000	000011	000000

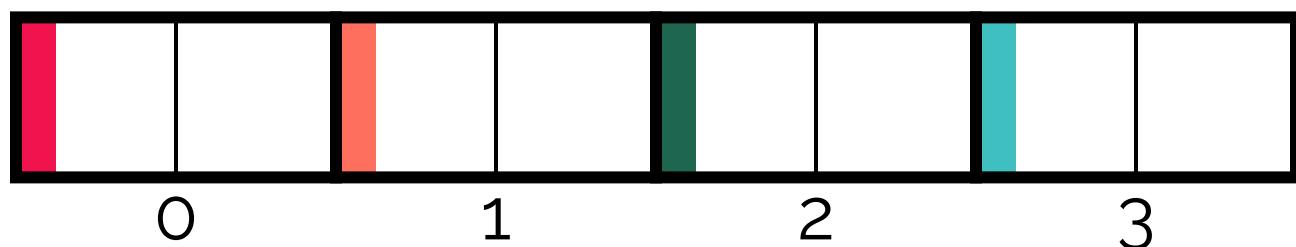


- Different bucket for each number
- Wastes 60B in each cache line



Offset = 64B

Number	Tag	Index	Offset
A	..100000	000000	000000
B	..100000	000001	000000
C	..100000	000010	000000
D	..100000	000011	000000

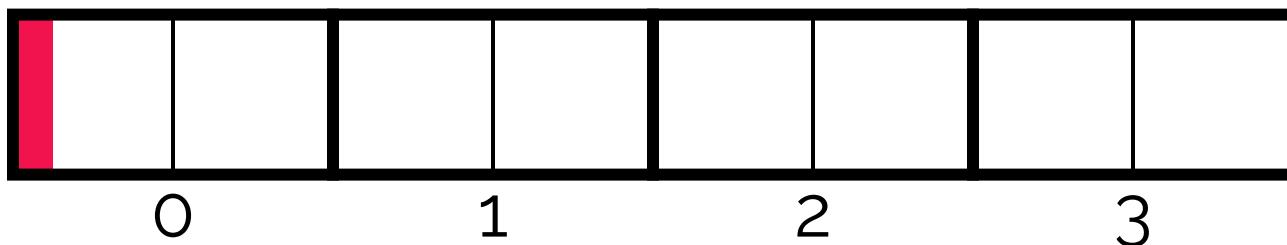


- Different bucket for each number
- Wastes 60B in each cache line
- Equally distributed among buckets



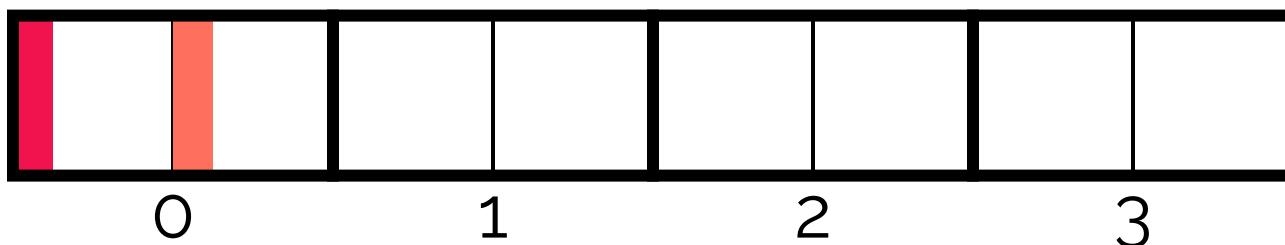
Offset = 4096B

Number	Tag	Index	Offset
A	.100000	000000	000000



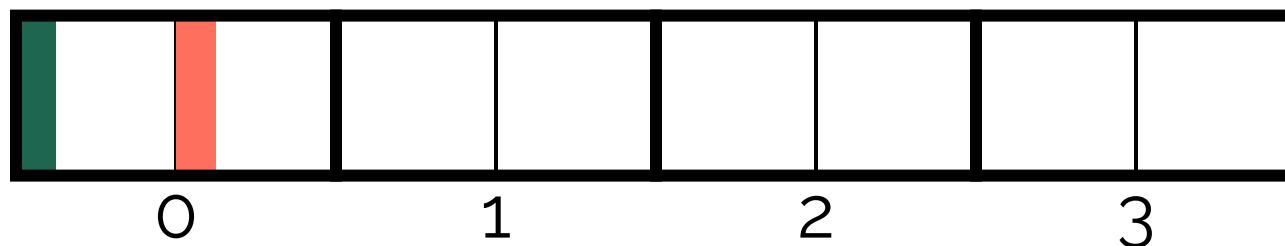
Offset = 4096B

Number	Tag	Index	Offset
A	.100000	000000	000000
B	.100001	000000	000000



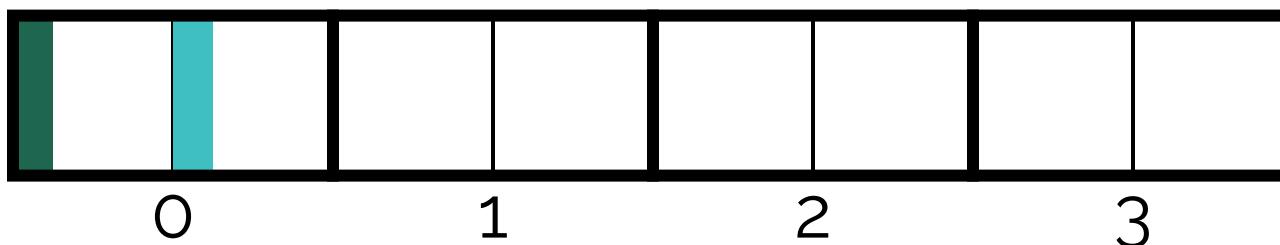
Offset = 4096B

Number	Tag	Index	Offset
A	.100000	000000	000000
B	.100001	000000	000000
C	.100010	000000	000000



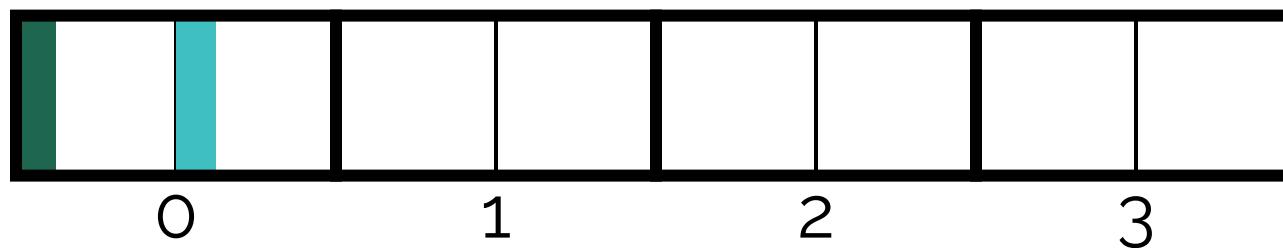
Offset = 4096B

Number	Tag	Index	Offset
A	.100000	000000	000000
B	.100001	000000	000000
C	.100010	000000	000000
D	.100011	000000	000000



Offset = 4096B

Number	Tag	Index	Offset
A	.100000	000000	000000
B	.100001	000000	000000
C	.100010	000000	000000
D	.100011	000000	000000

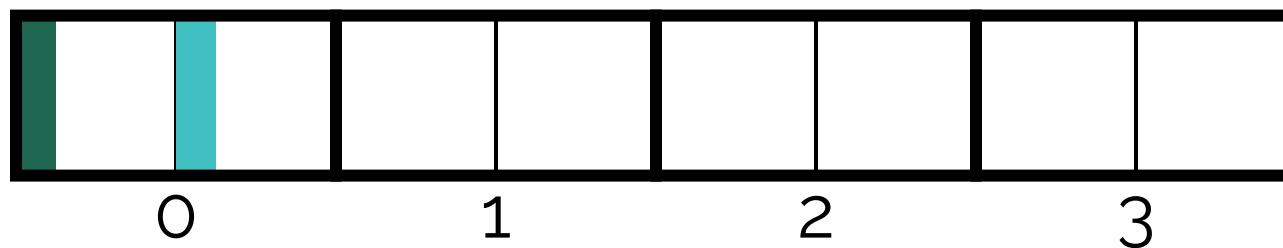


- Same bucket, but different cache lines for each number!



Offset = 4096B

Number	Tag	Index	Offset
A	.100000	000000	000000
B	.100001	000000	000000
C	.100010	000000	000000
D	.100011	000000	000000



- Same bucket, but different cache lines for each number!
- Bucket full => evictions necessary



How to measure?

l1d.replacement

How many times was a cache line loaded into L1?



How to measure?

l1d.replacement

How many times was a cache line loaded into L1?

```
$ perf stat -e l1d.replacement ./example1
4B    offset ->    149 558
4096B offset -> 426 218 383
```



Code (backup)

```
float F = static_cast<float>(std::stof(argv[1]));
std::vector<float> data(4 * 1024 * 1024, 1);

for (int r = 0; r < 100; r++)
{
    for (auto& item: data)
    {
        item *= F;
    }
}
```



Result (backup)

```
> time -p ./example2 0
real 0,12
user 0,12
sys 0,00
> time -p ./example2 0.1
real 0,47
user 0,46
sys 0,00
> time -p ./example2 0.3
real 0,70
user 0,69
sys 0,00
```



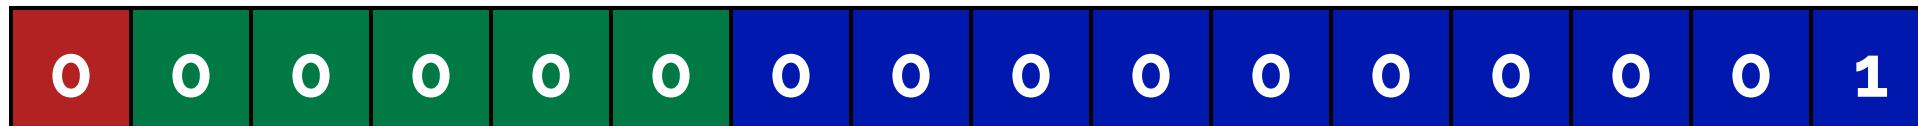
Denormal floating point numbers

0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	1
---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---

$$(-1)^0 * 2^{00000-01110} * 0.0000000001$$



Denormal floating point numbers



Zero exponent

Non-zero significand

$$(-1)^0 * 2^{00000-01110} * 0.0000000001$$



Denormal floating point numbers



Zero exponent

Non-zero significand

$$(-1)^0 * 2^{00000-01110} * 0.0000000001$$

- Numbers close to zero
- Hidden bit = 0, smaller bias



Denormal floating point numbers



Zero exponent

Non-zero significand

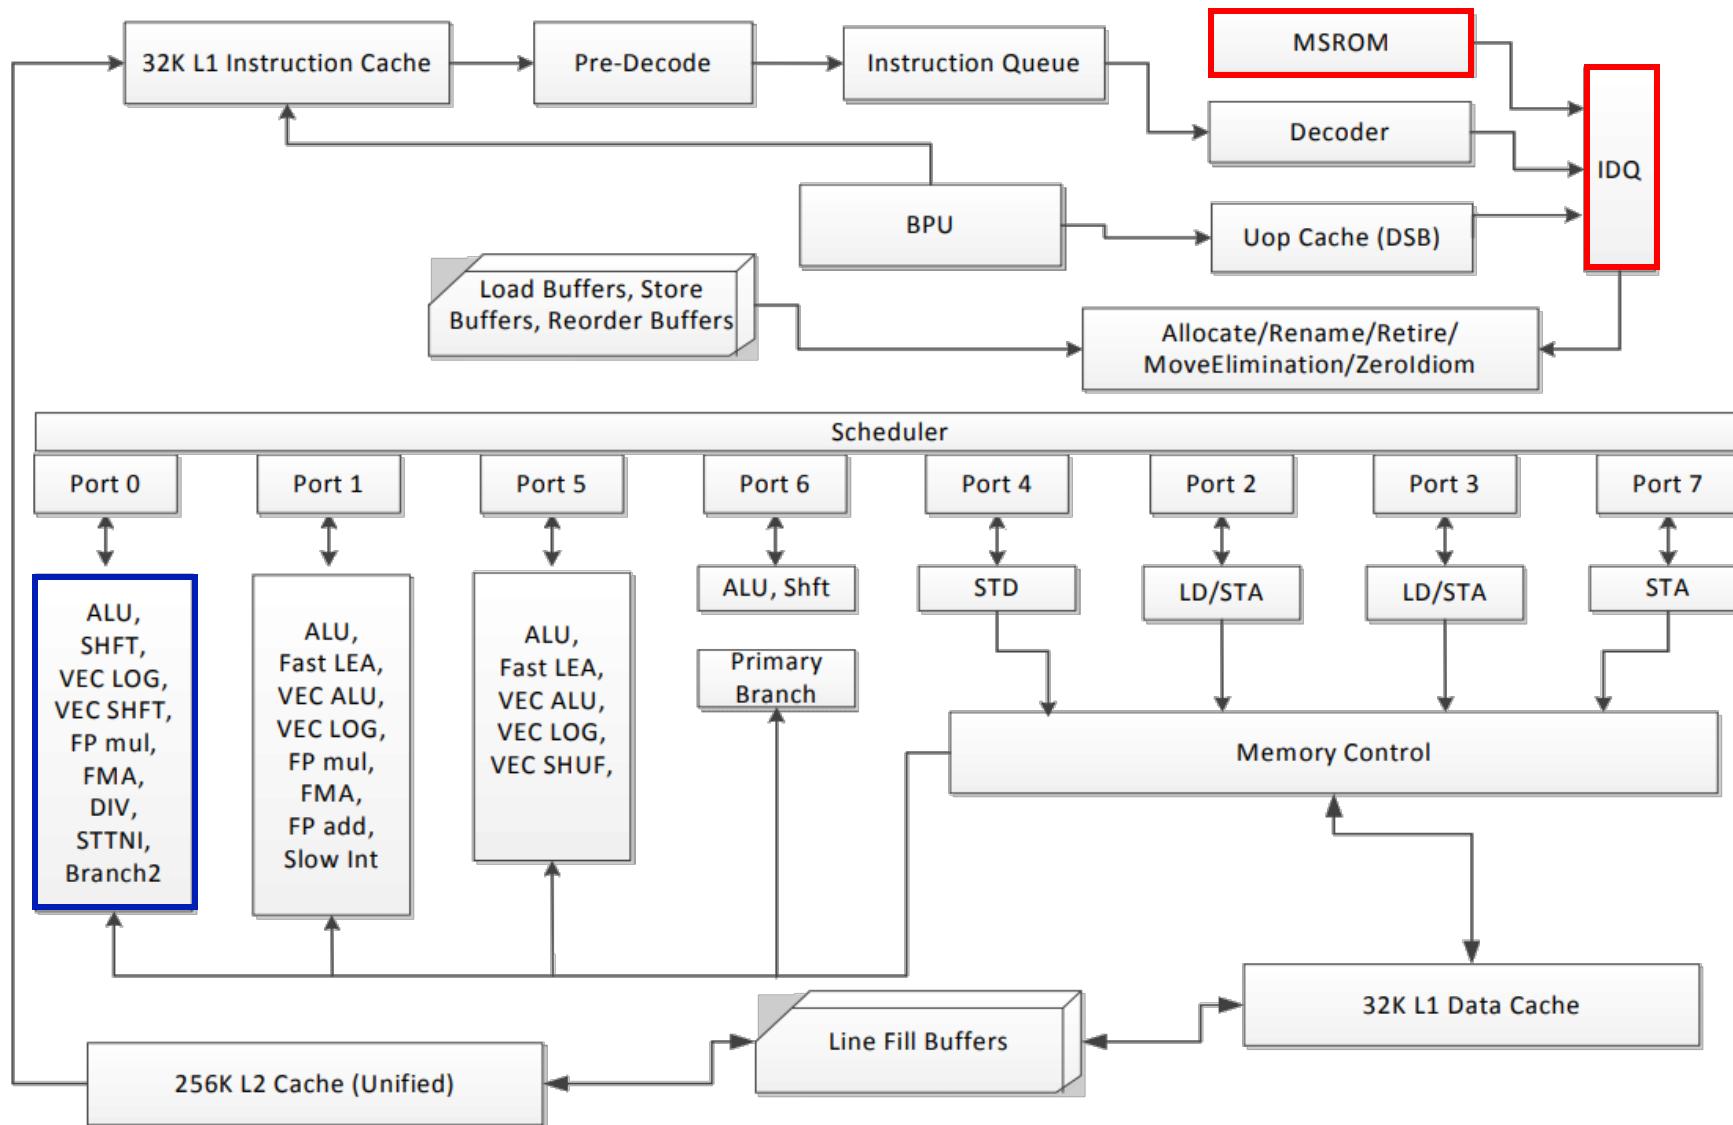
$$(-1)^0 * 2^{00000-01110} * 0.0000000001$$

- Numbers close to zero
- Hidden bit = 0, smaller bias

Operations on denormal numbers are slow!



Floating point handling



How to measure?

fp_assist.any

How many times the CPU switched to the microcode FP handler?



How to measure?

`fp_assist.any`

How many times the CPU switched to the microcode FP handler?

```
$ perf stat -e fp_assist.any ./example2
0      ->          0
0.3    -> 15 728 640
```



How to fix it?

- The nuclear option: `-ffast-math`
 - Sacrifice correctness to gain more FP performance



How to fix it?

- The nuclear option: `-ffast-math`
 - Sacrifice correctness to gain more FP performance
- Set CPU flags:
 - Flush-to-zero - treat denormal outputs as 0
 - Denormals-to-zero - treat denormal inputs as 0



How to fix it?

- The nuclear option: `-ffast-math`
 - Sacrifice correctness to gain more FP performance
- Set CPU flags:
 - Flush-to-zero - treat denormal outputs as 0
 - Denormals-to-zero - treat denormal inputs as 0

```
_mm_setcsr(_mm_getcsr() | 0x8040);
// or
_MM_SET_FLUSH_ZERO_MODE(_MM_FLUSH_ZERO_ON);
_MM_SET_DENORMALS_ZERO_MODE(_MM_DENORMALS_ZERO_ON);
```



There are many other effects

- NUMA
- 4k aliasing
- Misaligned accesses, cache line boundaries
- Instruction data dependencies
- Software prefetching
- Non-temporal stores & cache pollution
- Bandwidth saturation
- DRAM refresh intervals
- AVX/SSE transition penalty
- ...



Thank you!

For more examples visit:
github.com/kobzol/hardware-effects

Jakub Beránek

Slides built with github.com/spirali/elsie



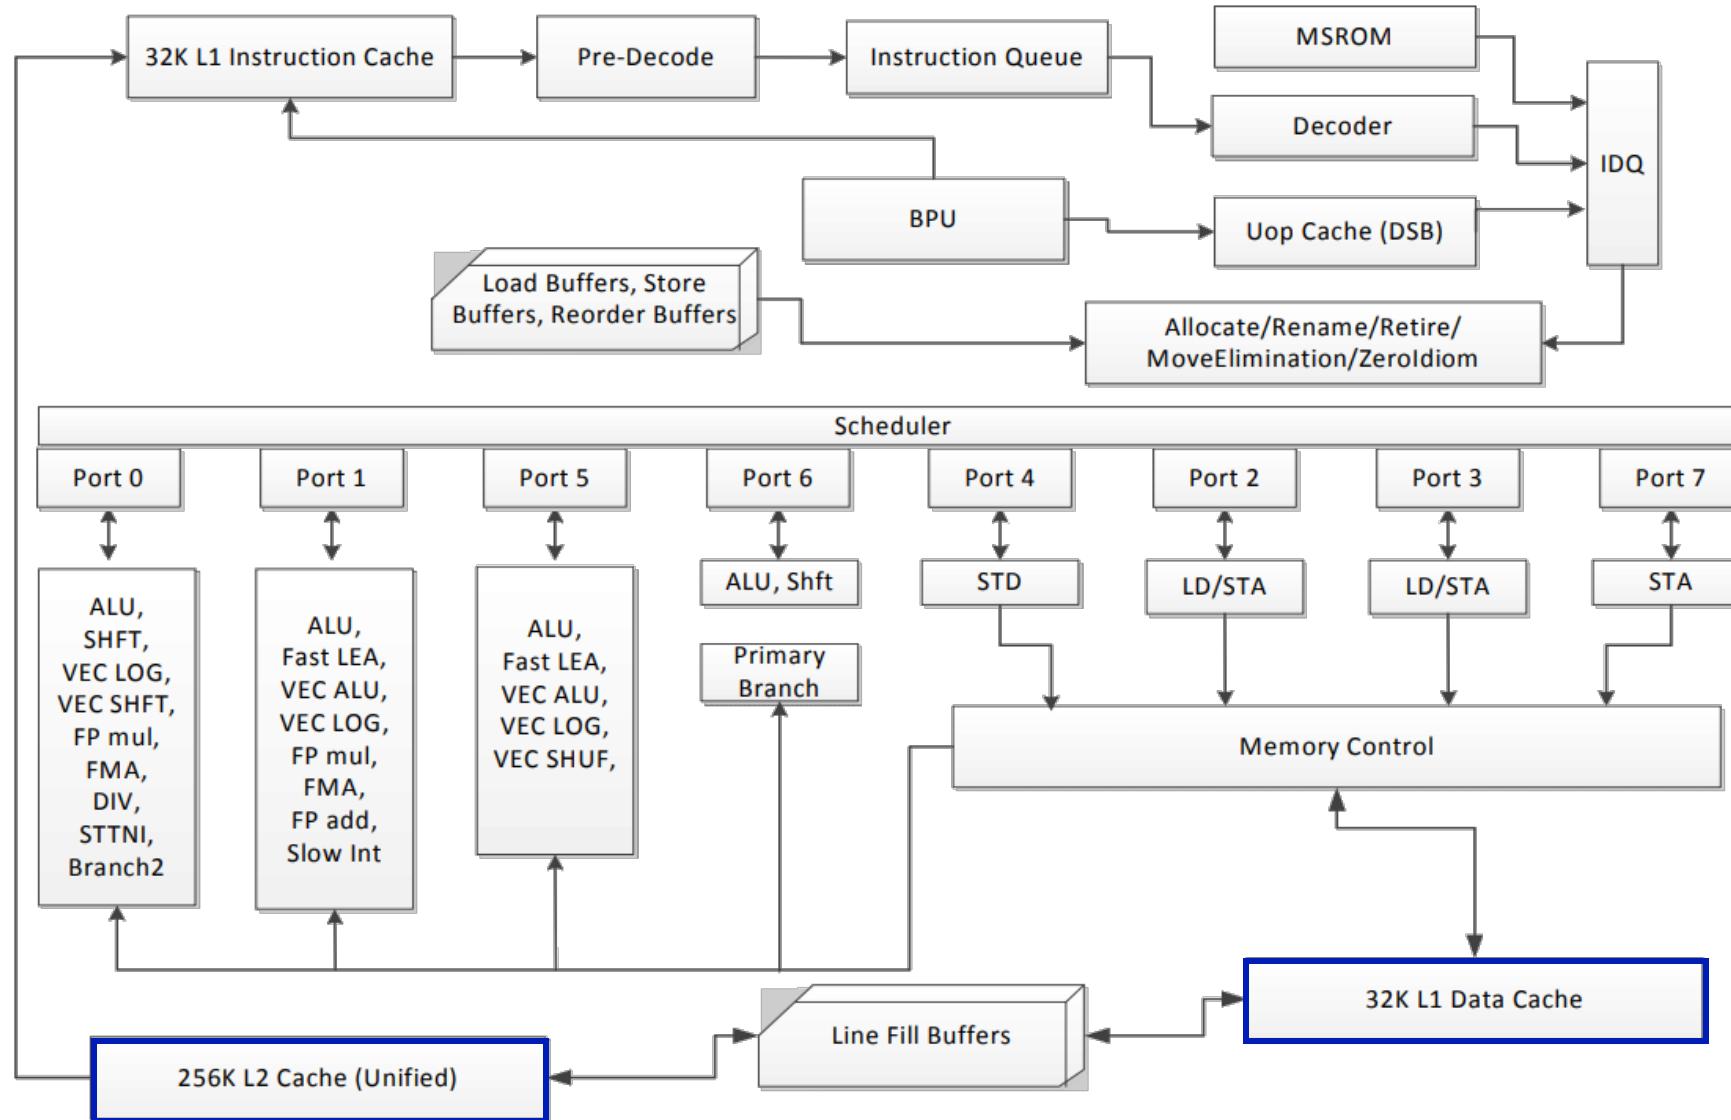
Code (backup)

```
// tid - [0, NO_OF_THREADS)
void thread_fn(int tid, double* data)
{
    size_t repetitions = 1024 * 1024 * 1024UL;
    for (size_t i = 0; i < repetitions; i++)
    {
        data[tid] *= i;
    }
}
```

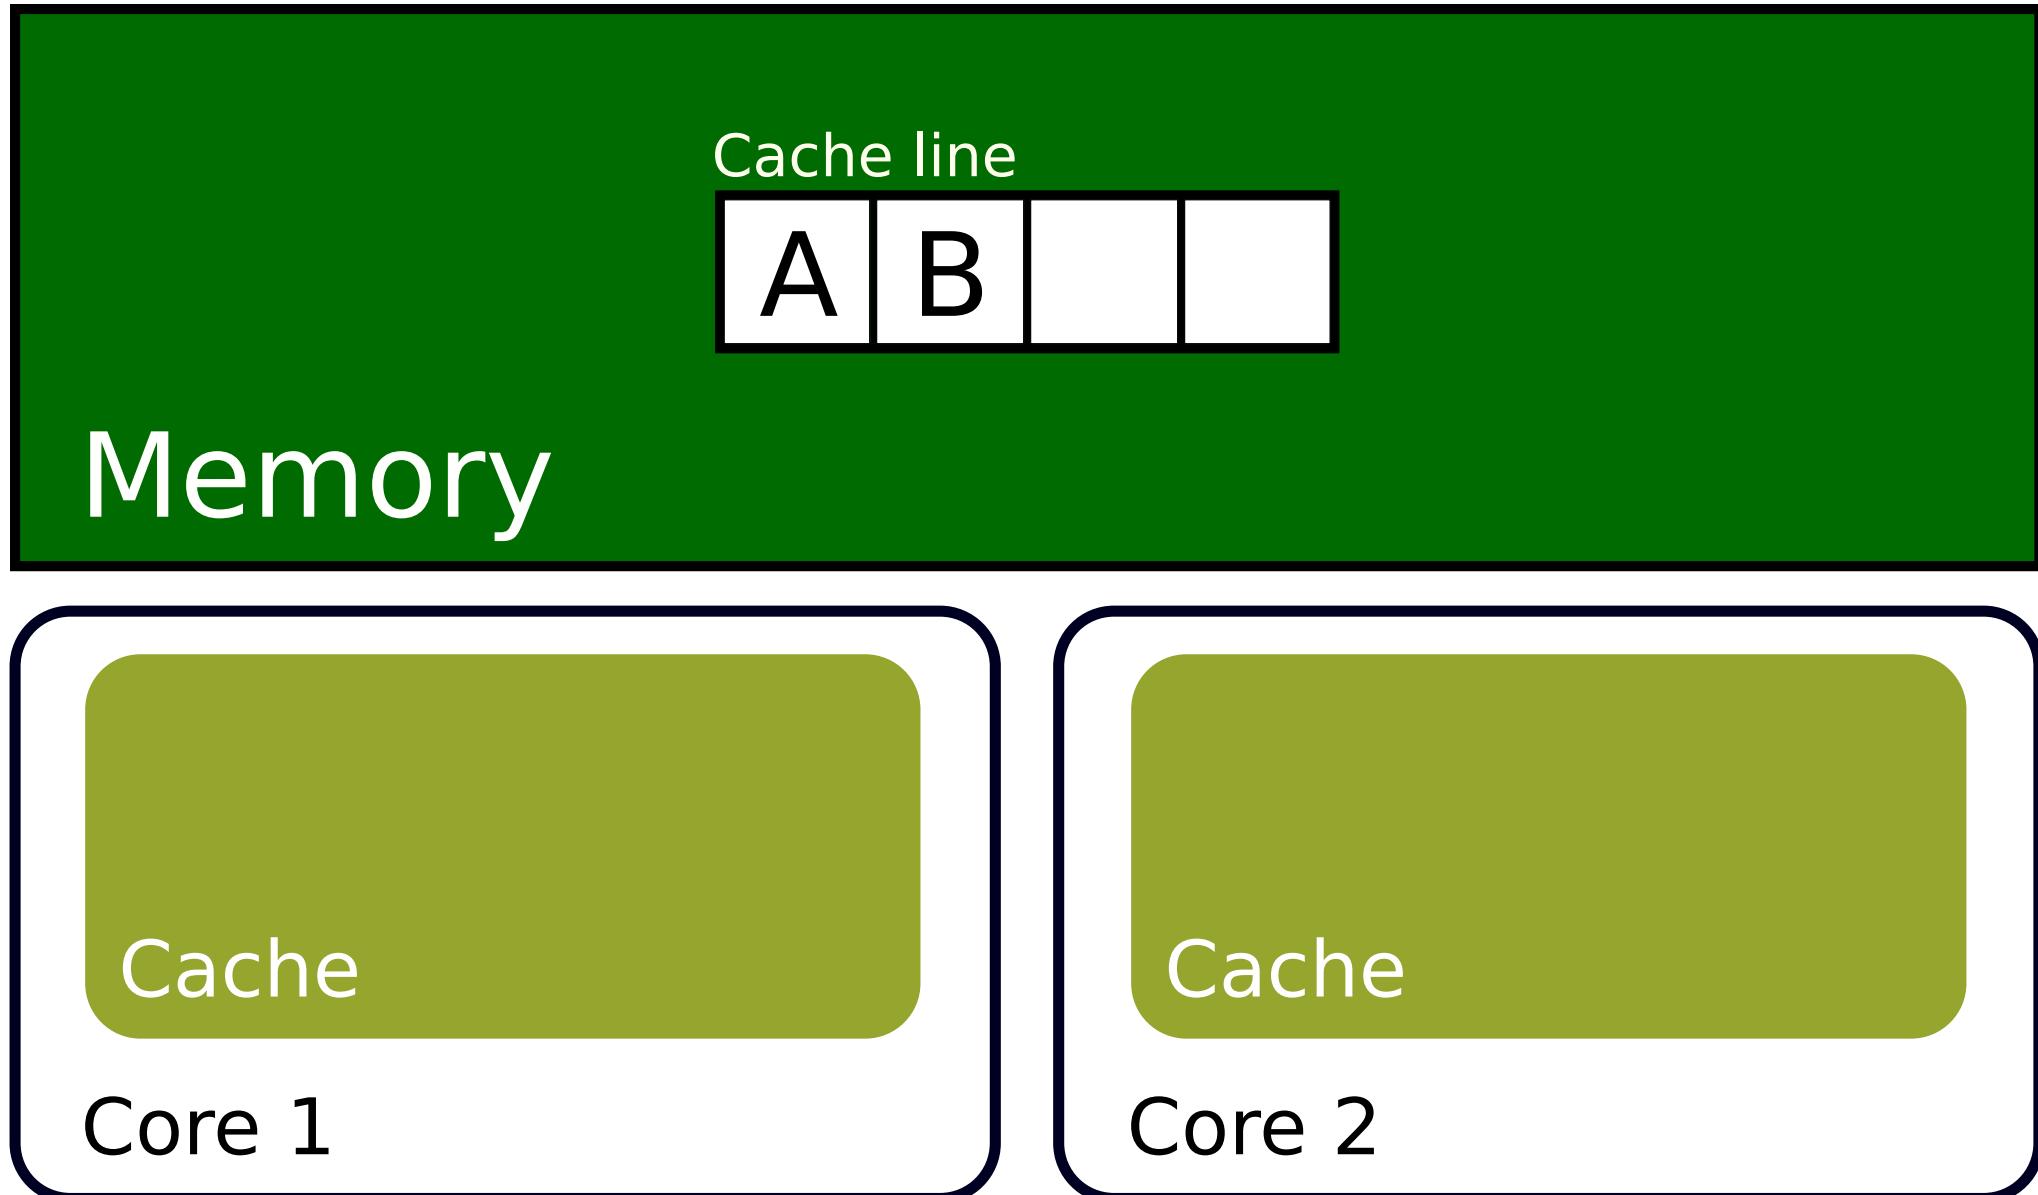
Result (backup)

```
> time -p ./example2 1
real 2,60
user 2,60
sys 0,00
> time -p ./example2 2
real 2,92
user 5,83
sys 0,00
> time -p ./example2 3
real 3,15
user 9,04
sys 0,01
> time -p ./example2 4
real 3,39
user 13,44
sys 0,00
```

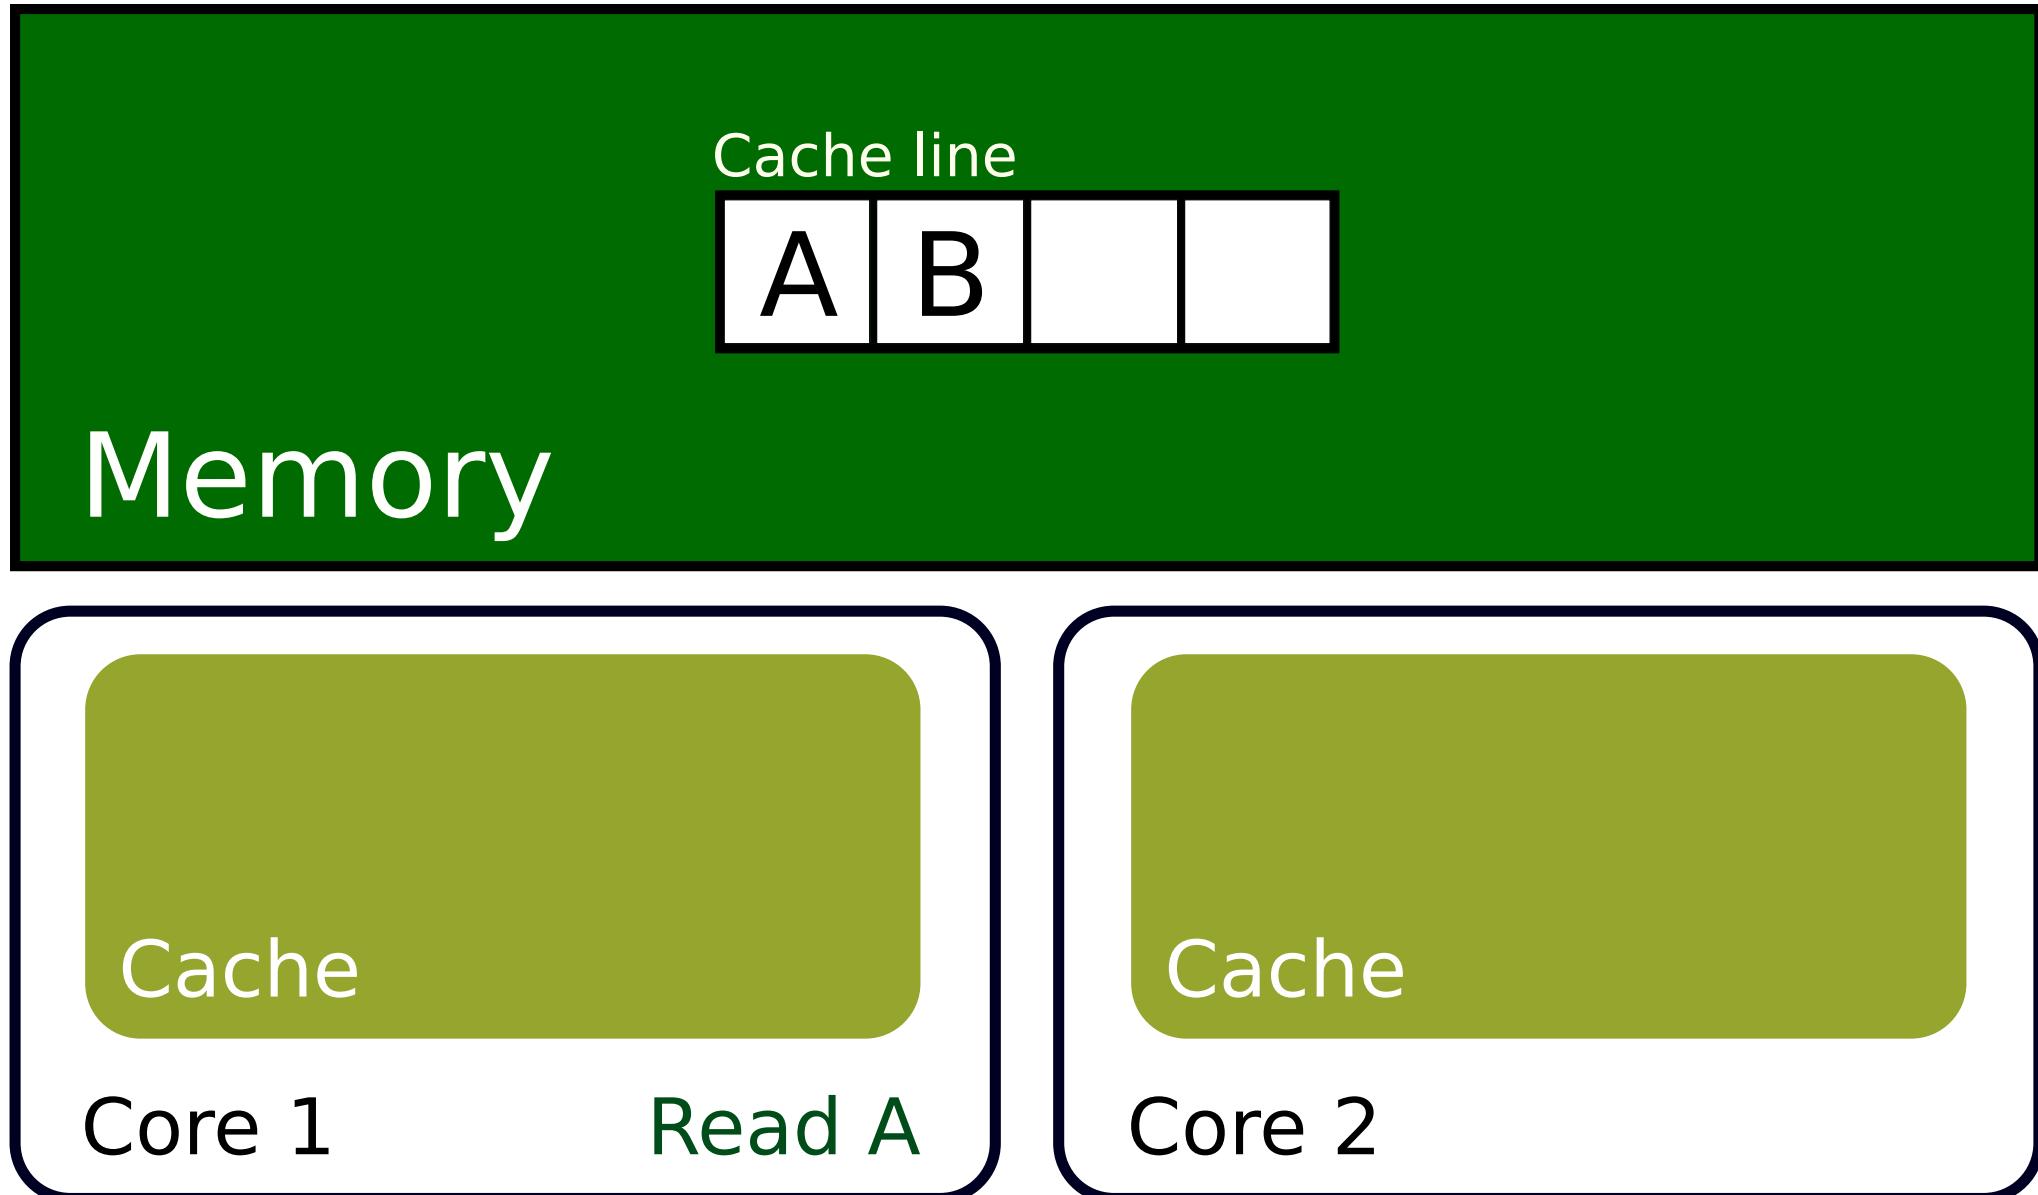
Cache system



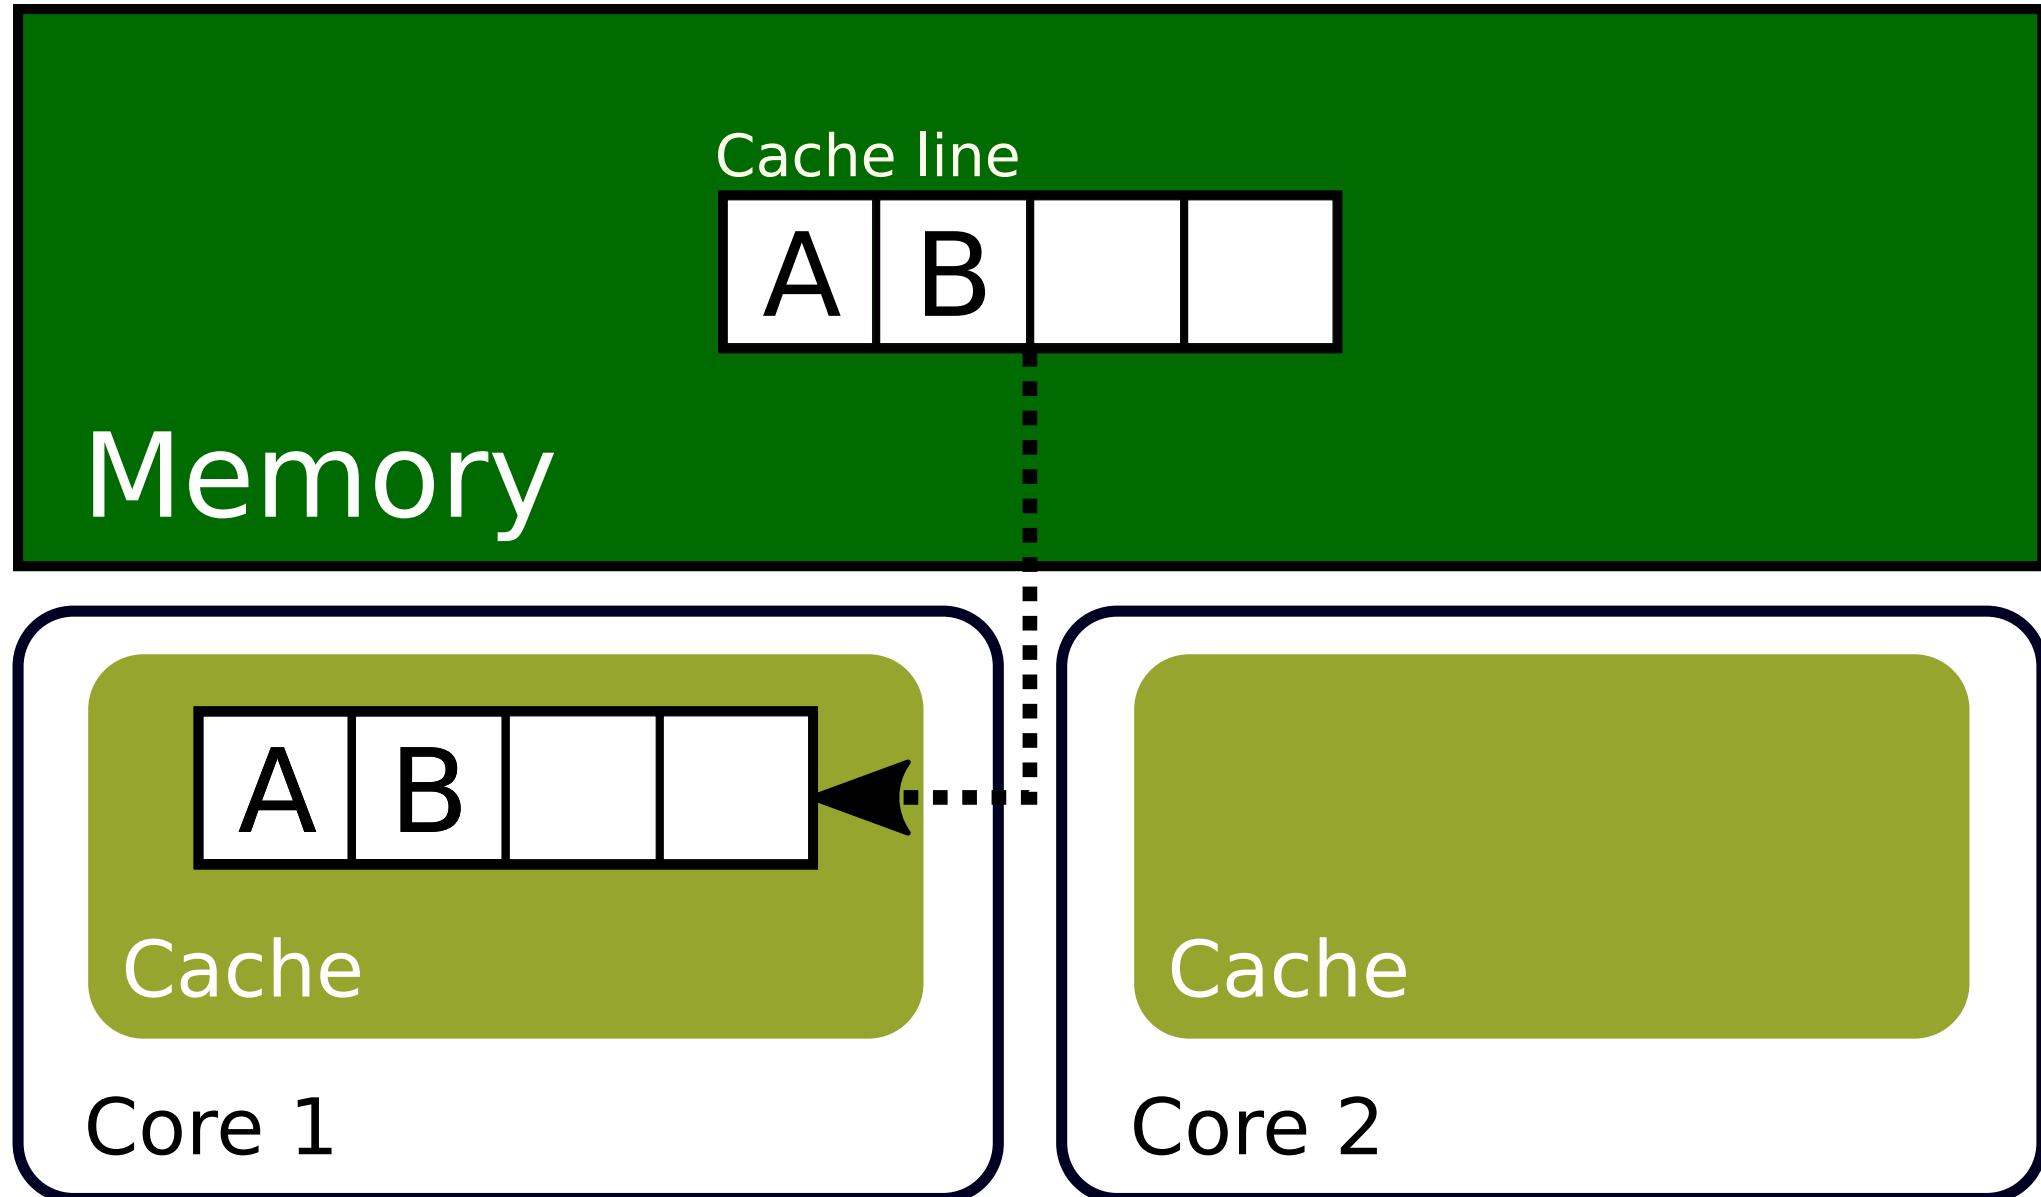
Cache coherency



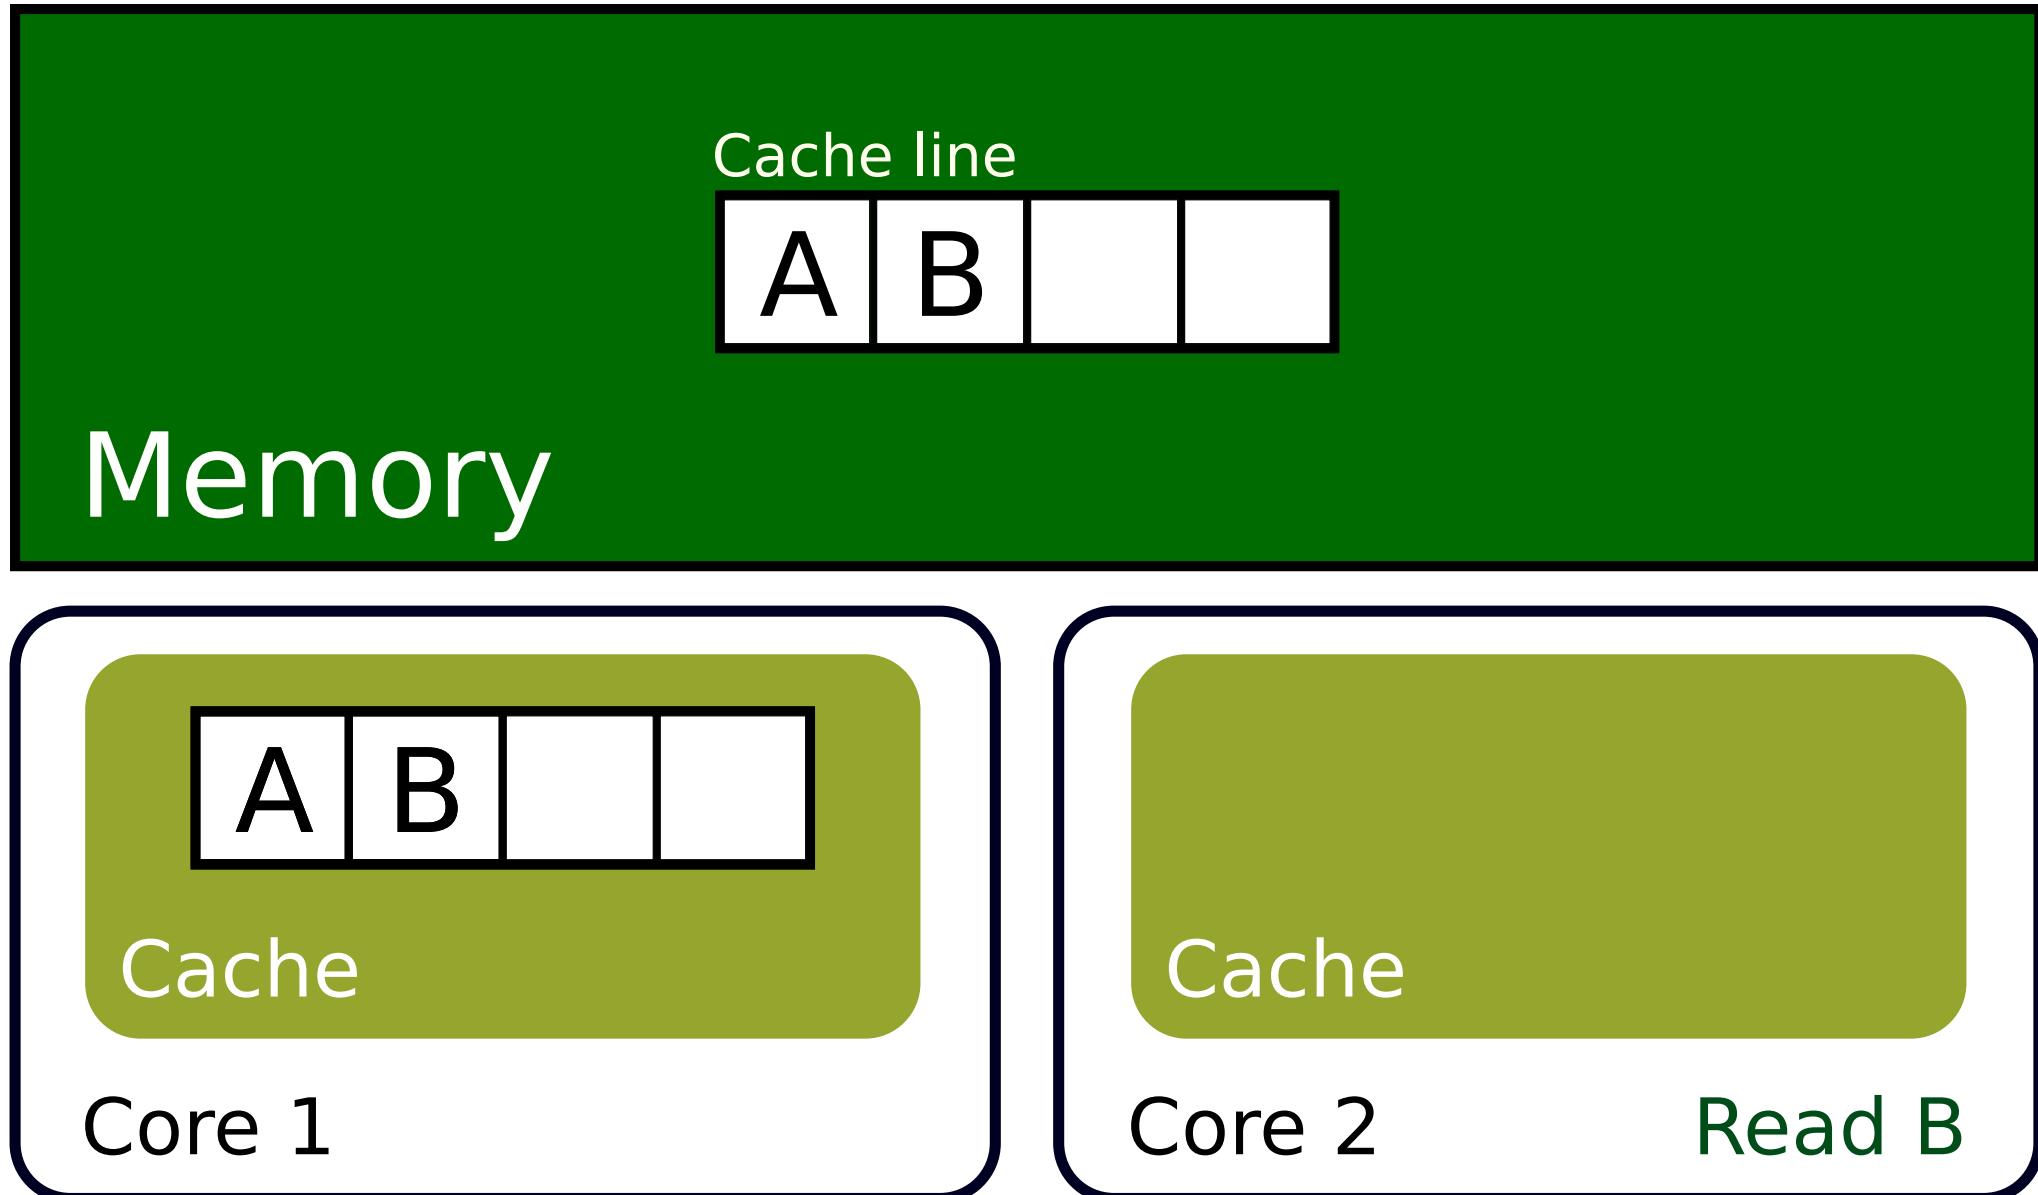
Cache coherency



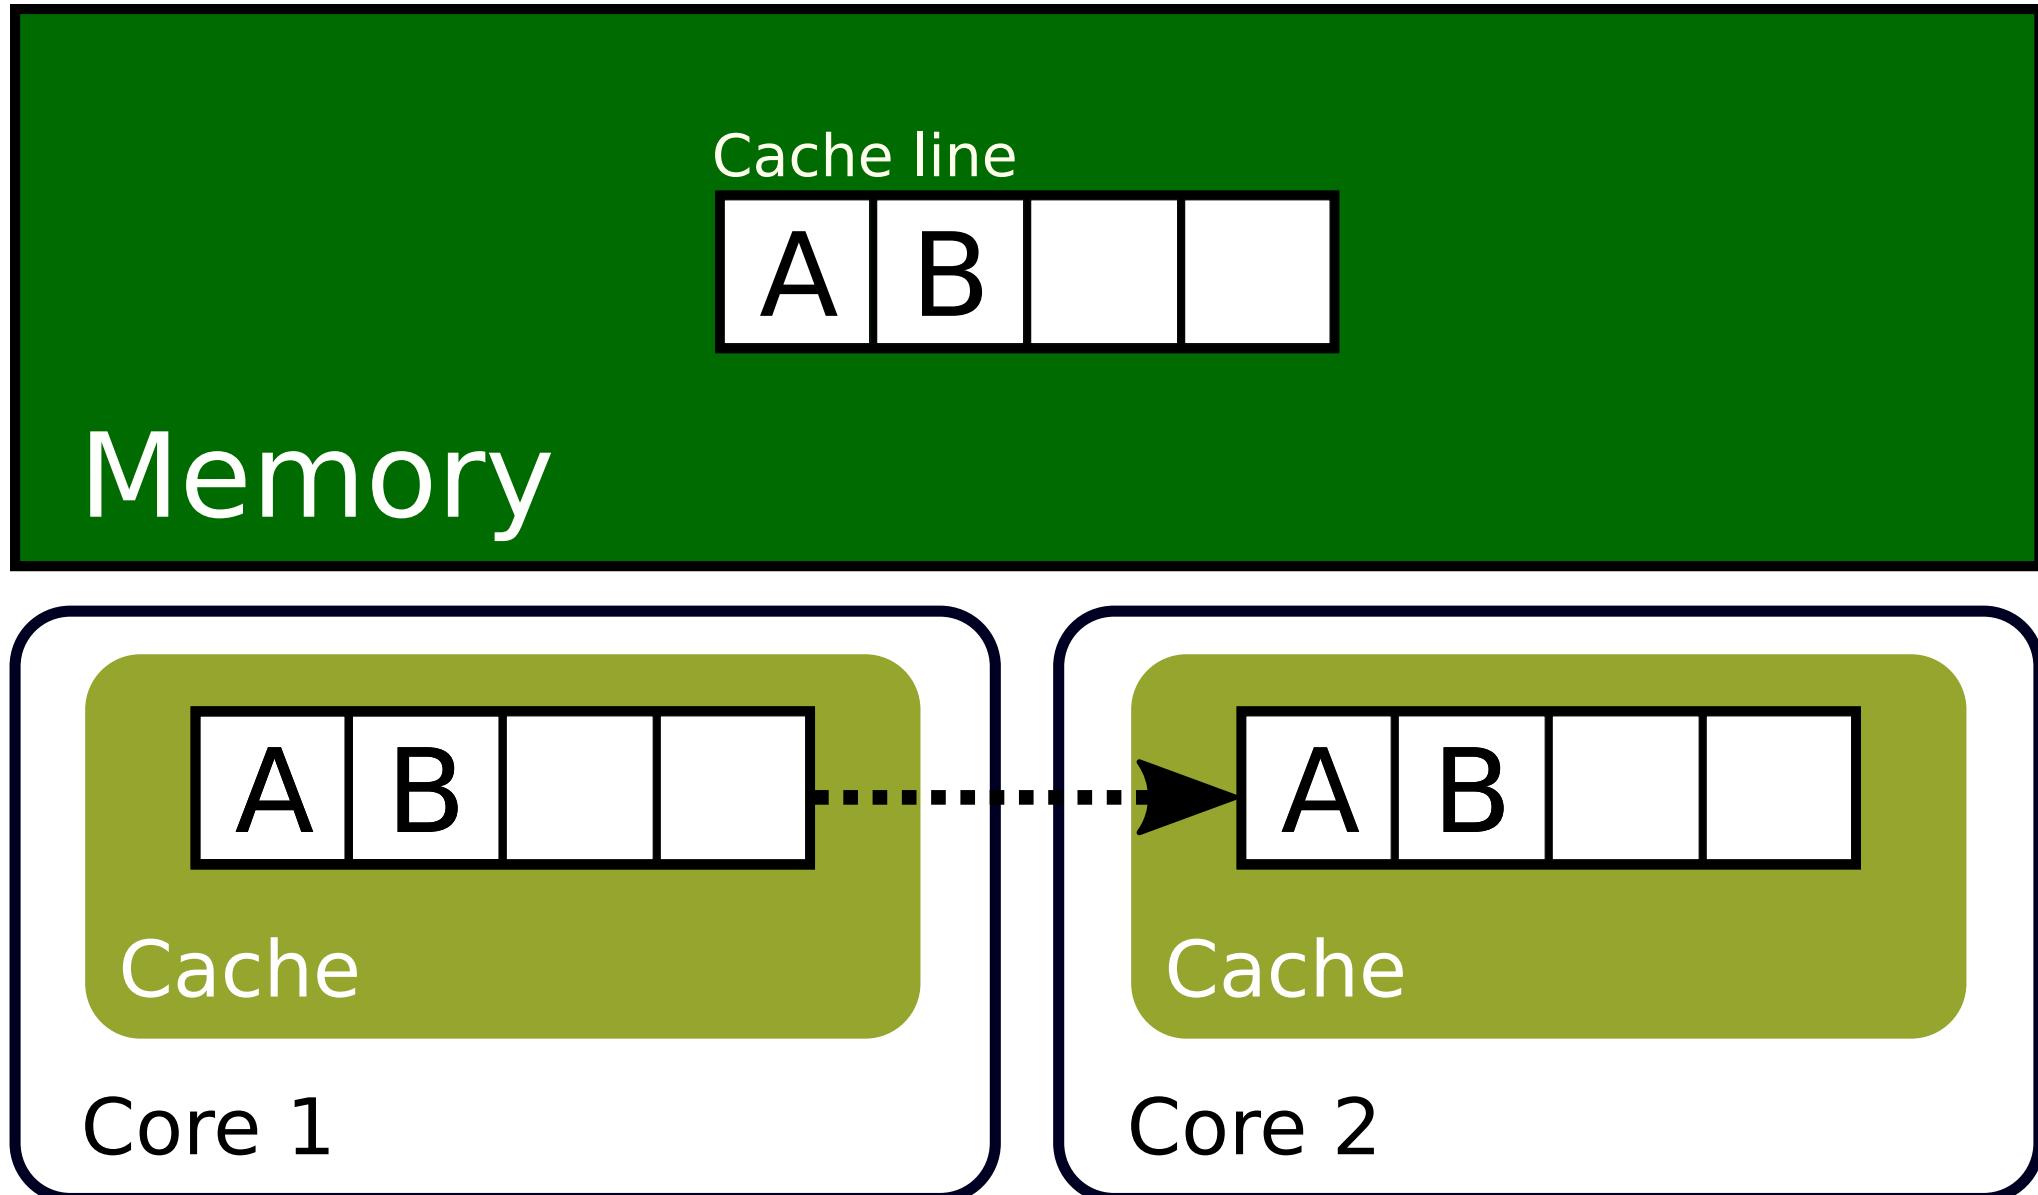
Cache coherency



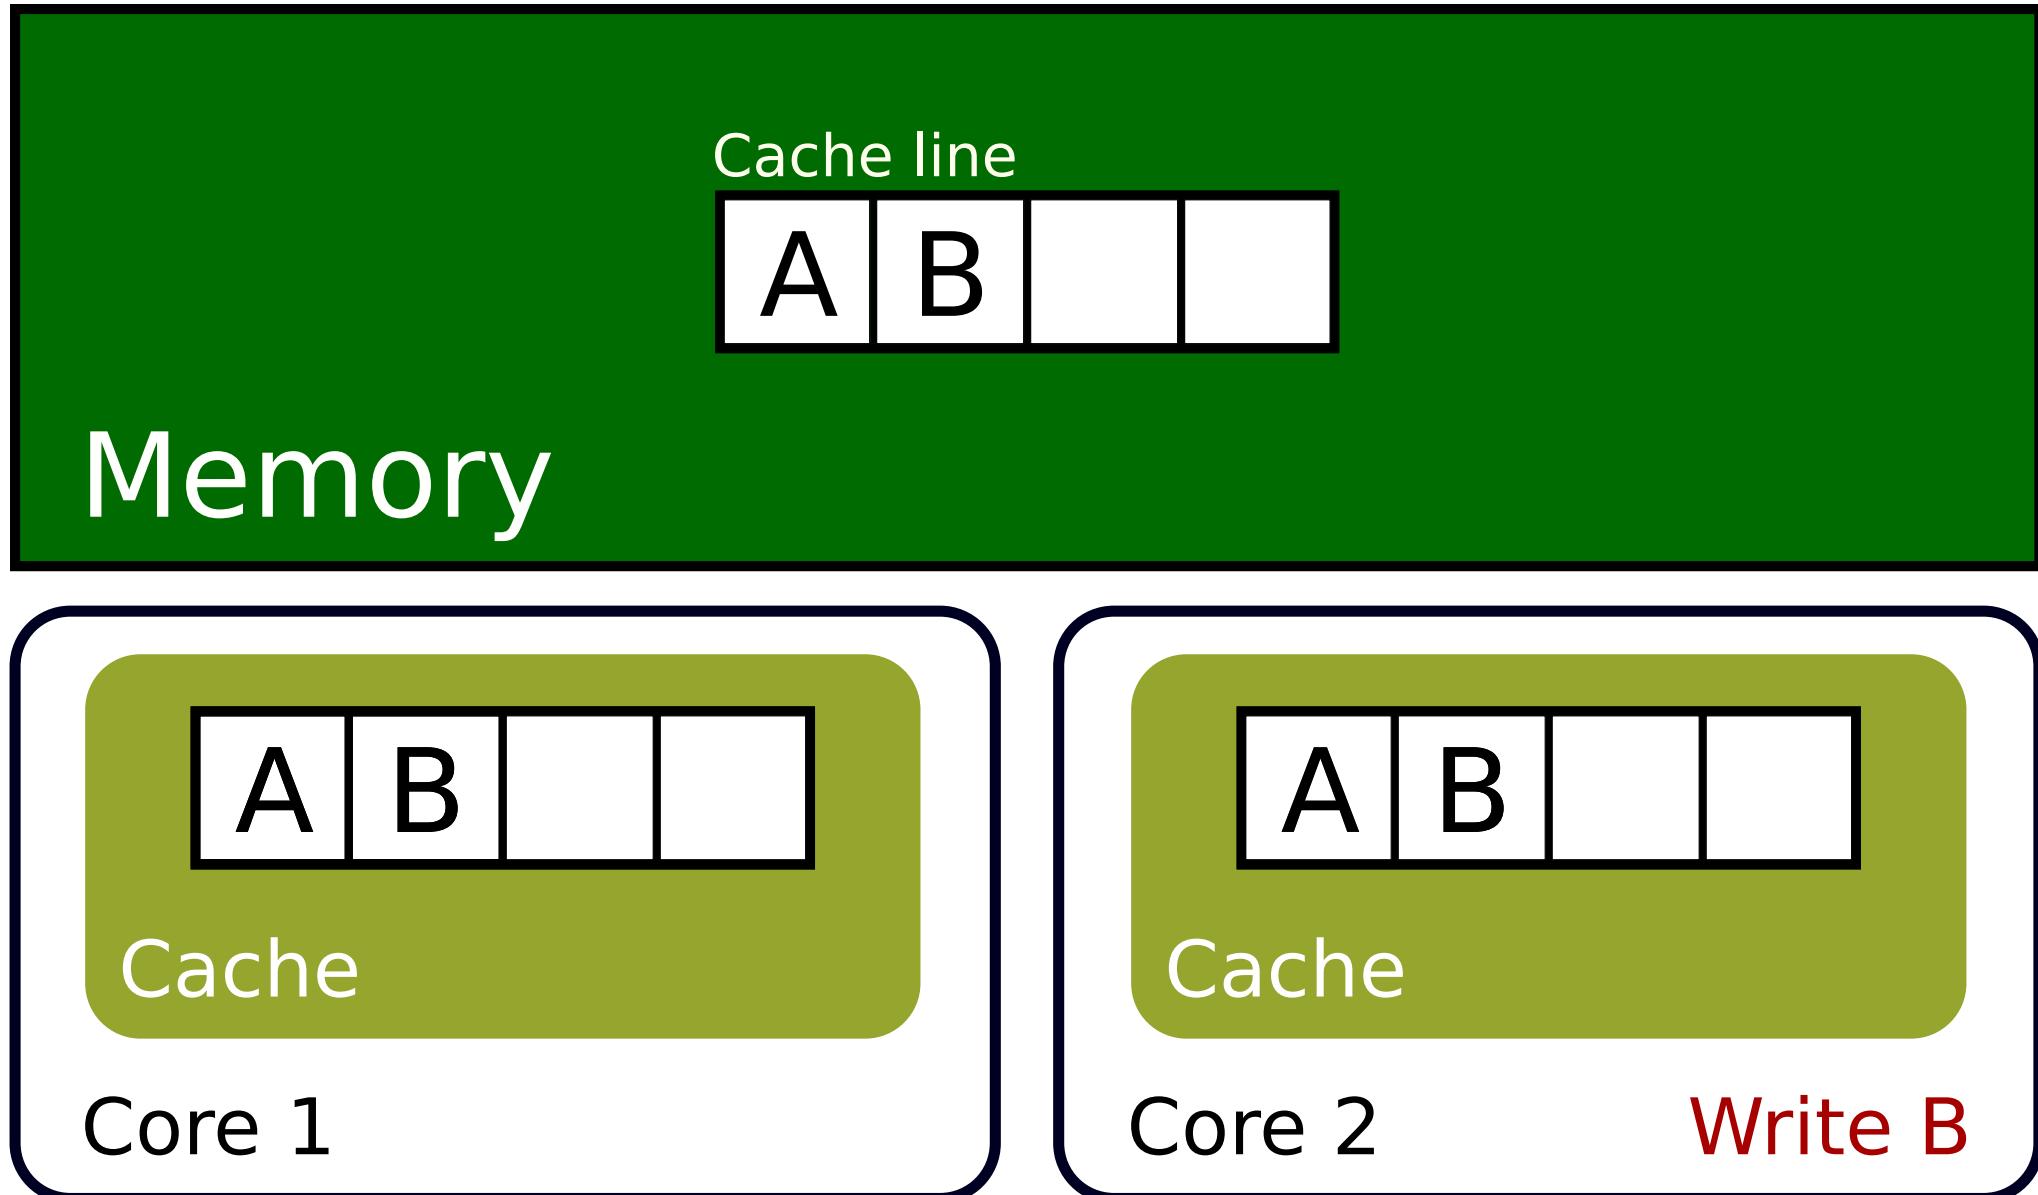
Cache coherency



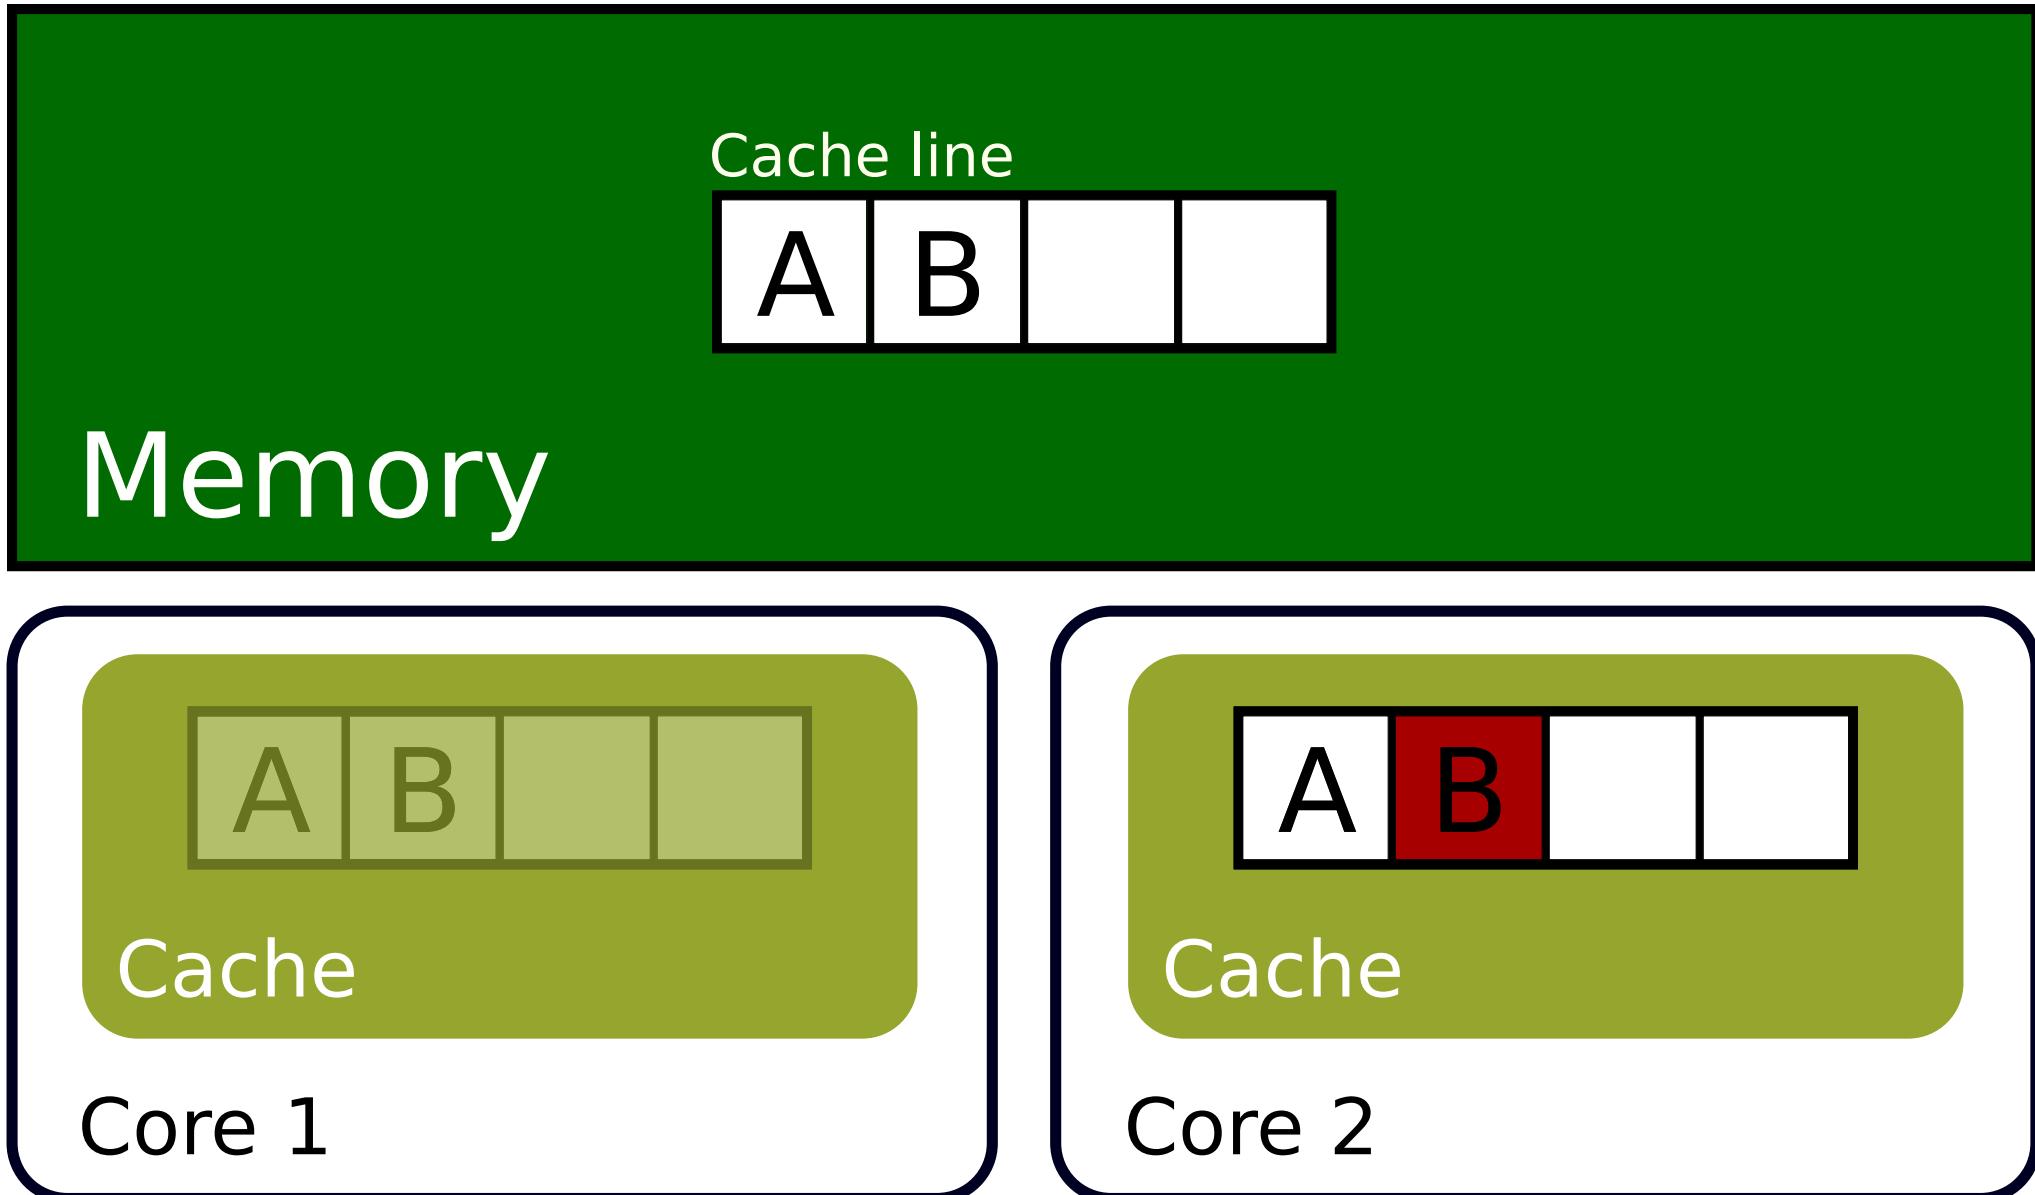
Cache coherency



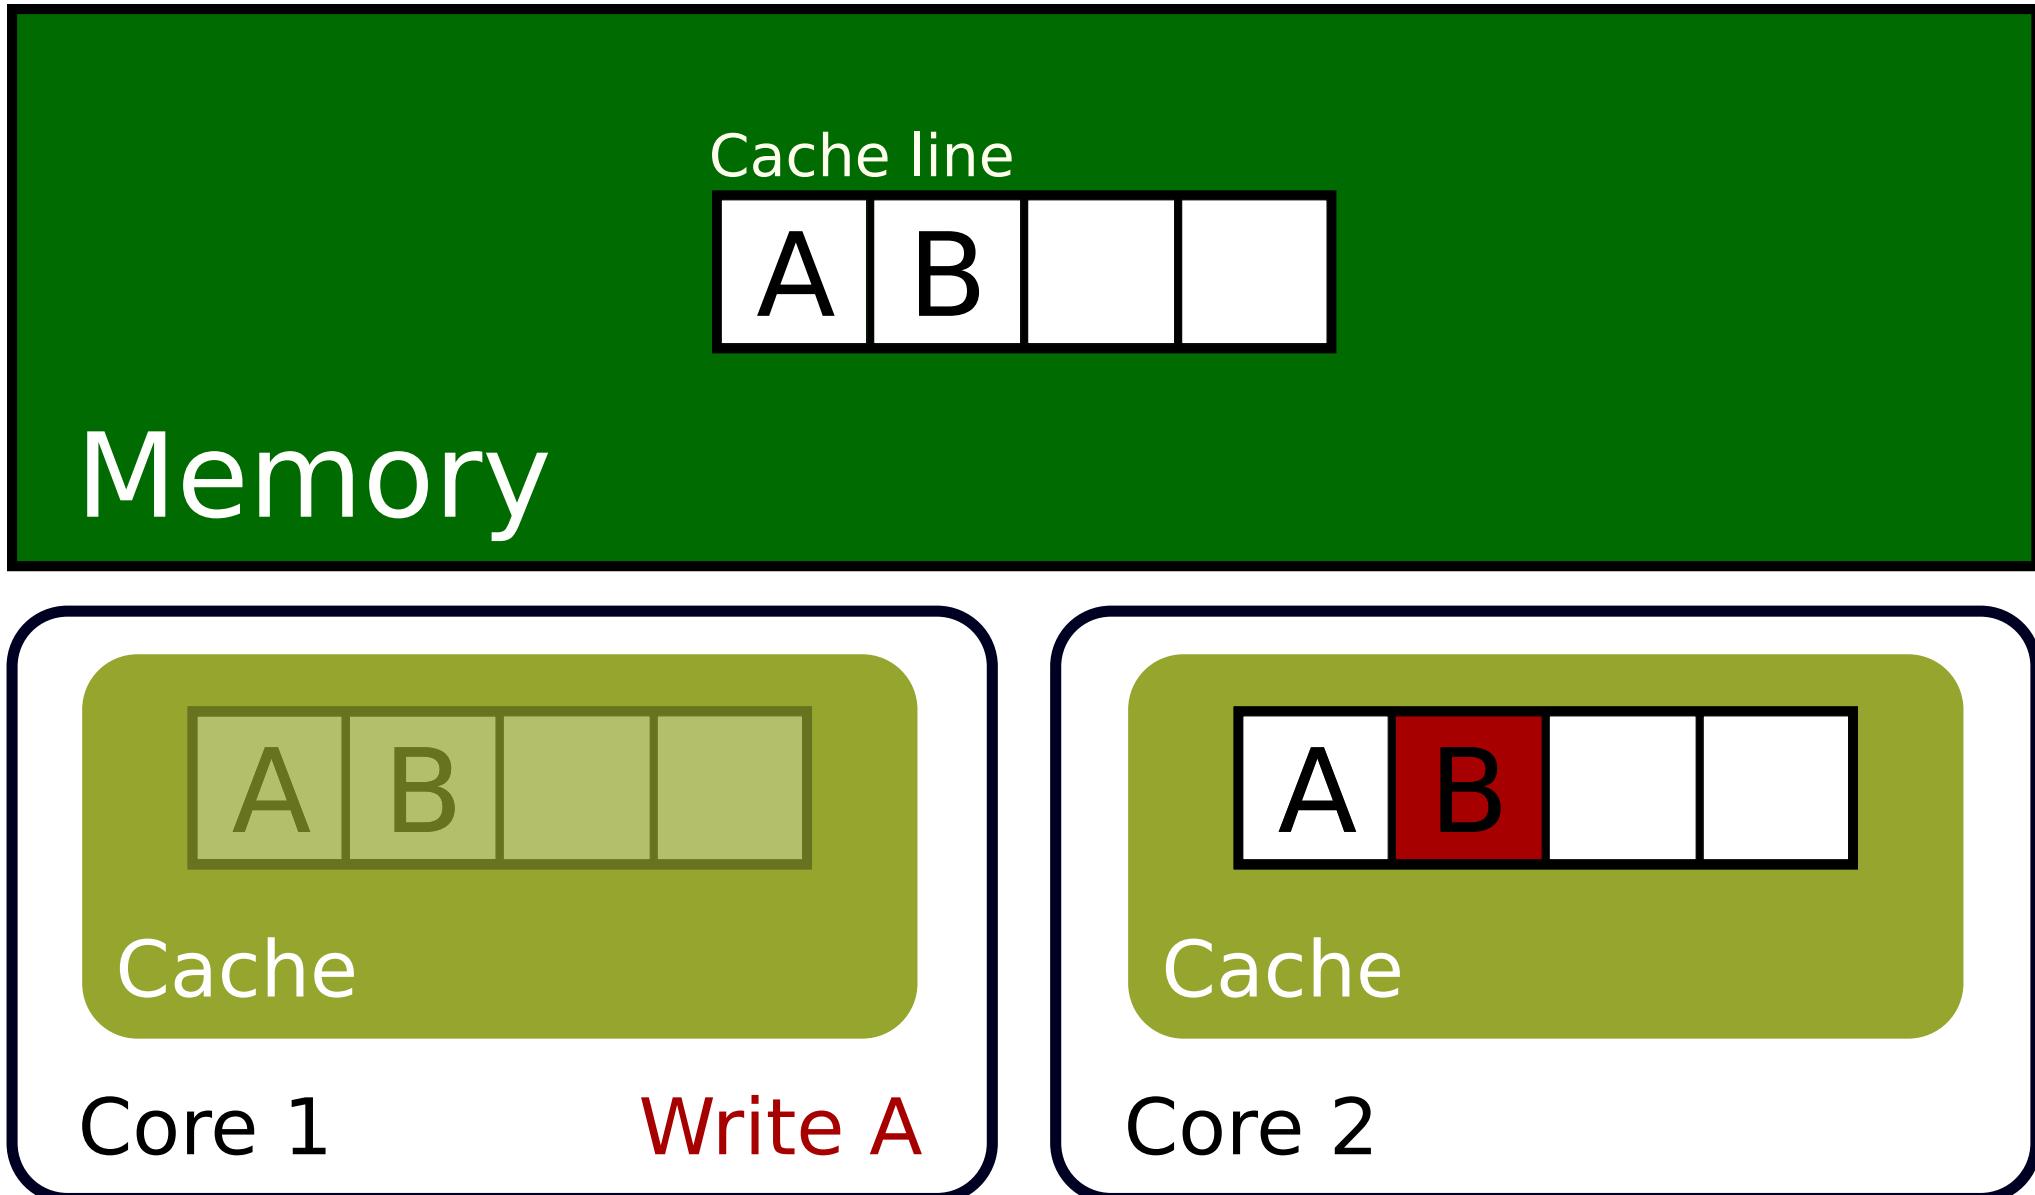
Cache coherency



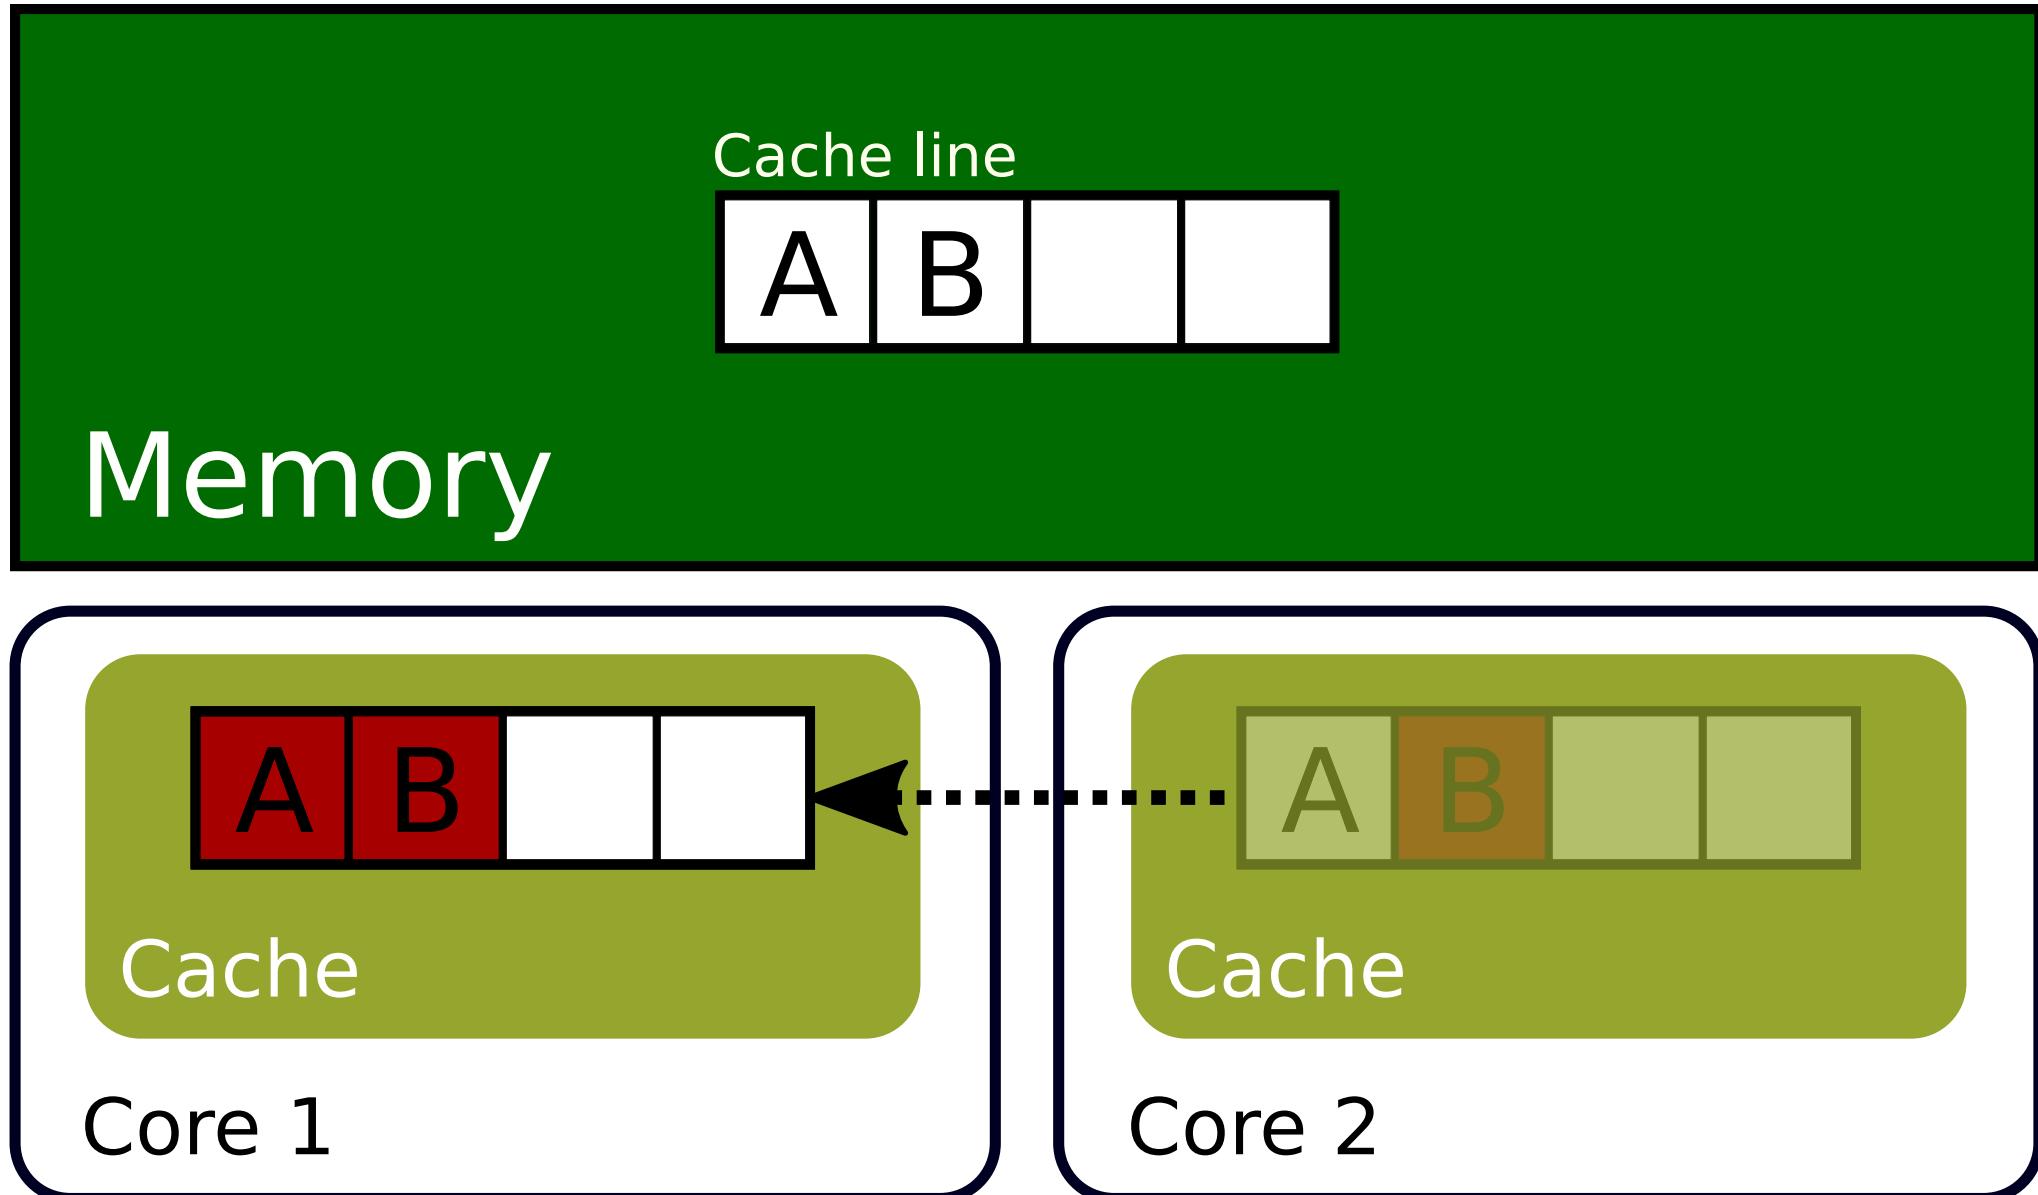
Cache coherency



Cache coherency



Cache coherency



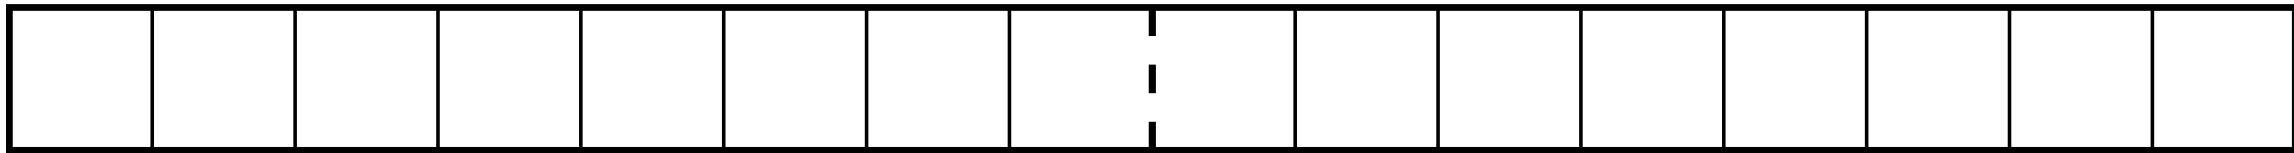
False sharing

double arr[16];

arr[0]

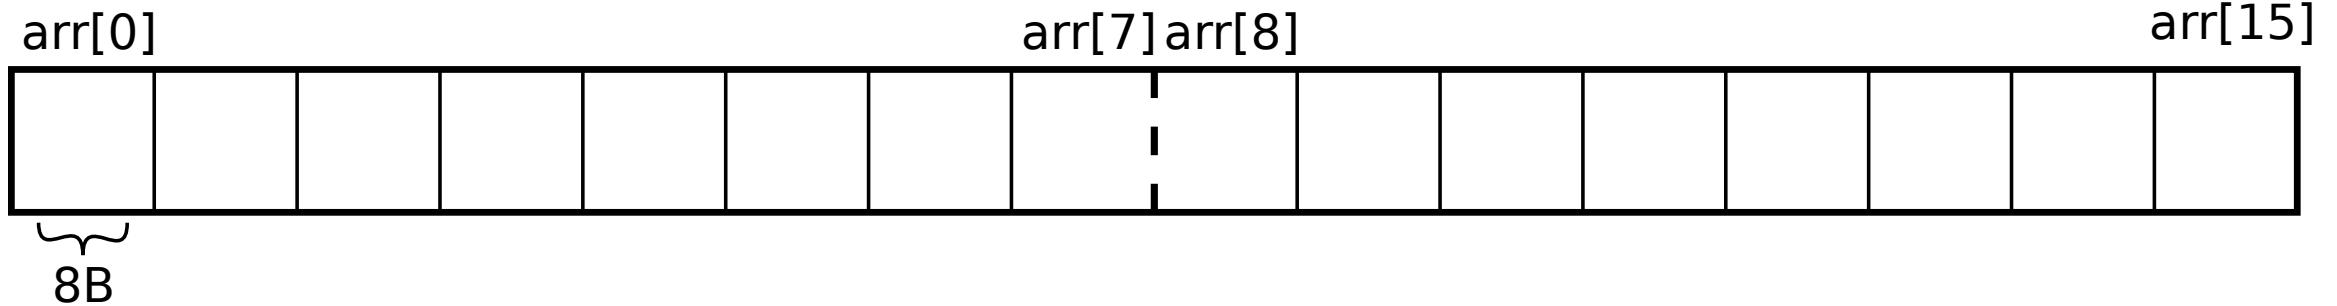
arr[7] arr[8]

arr[15]



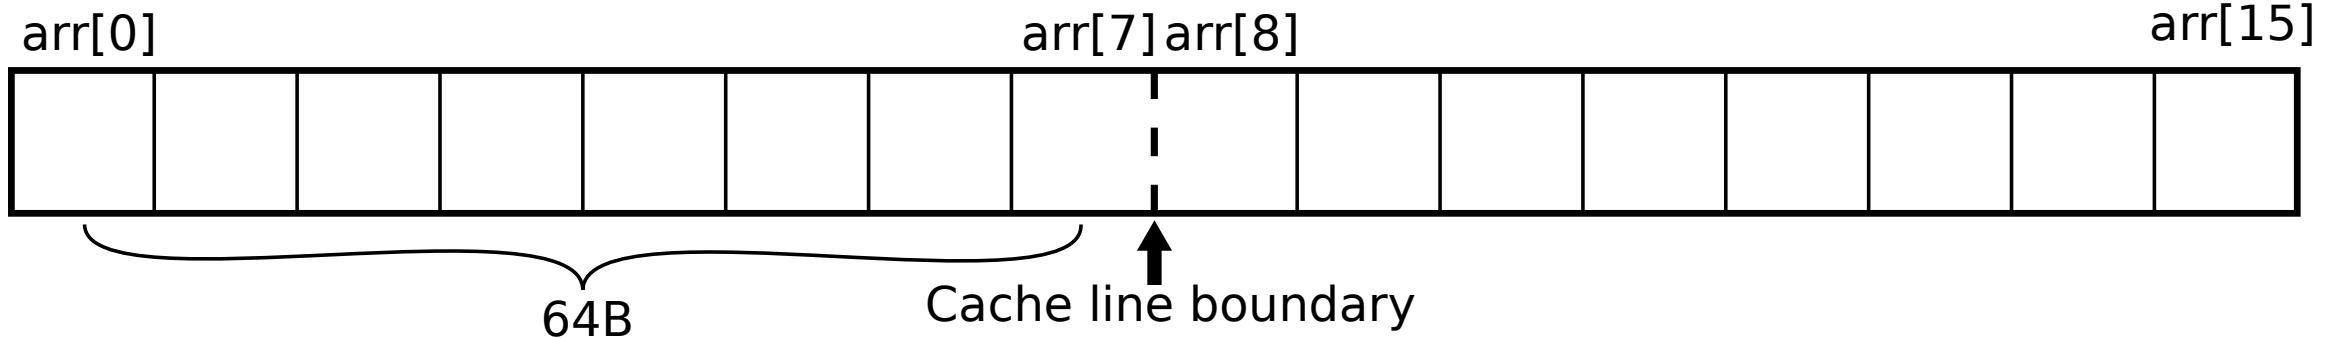
False sharing

double arr[16];



False sharing

double arr[16];



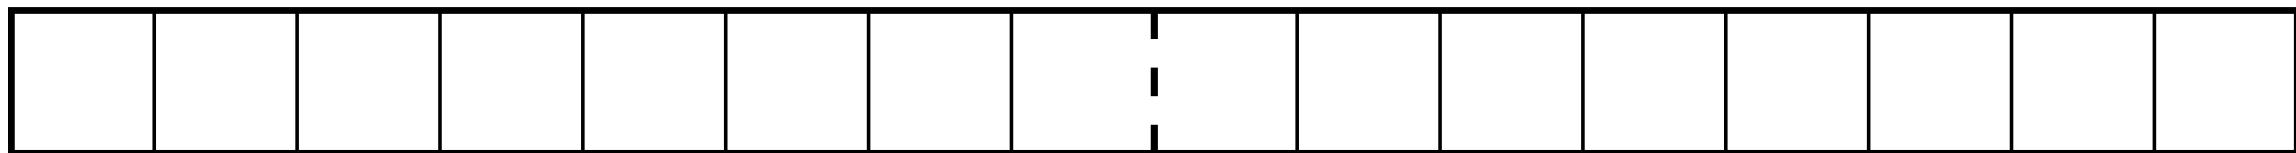
False sharing

double arr[16];

arr[0]

arr[7] arr[8]

arr[15]

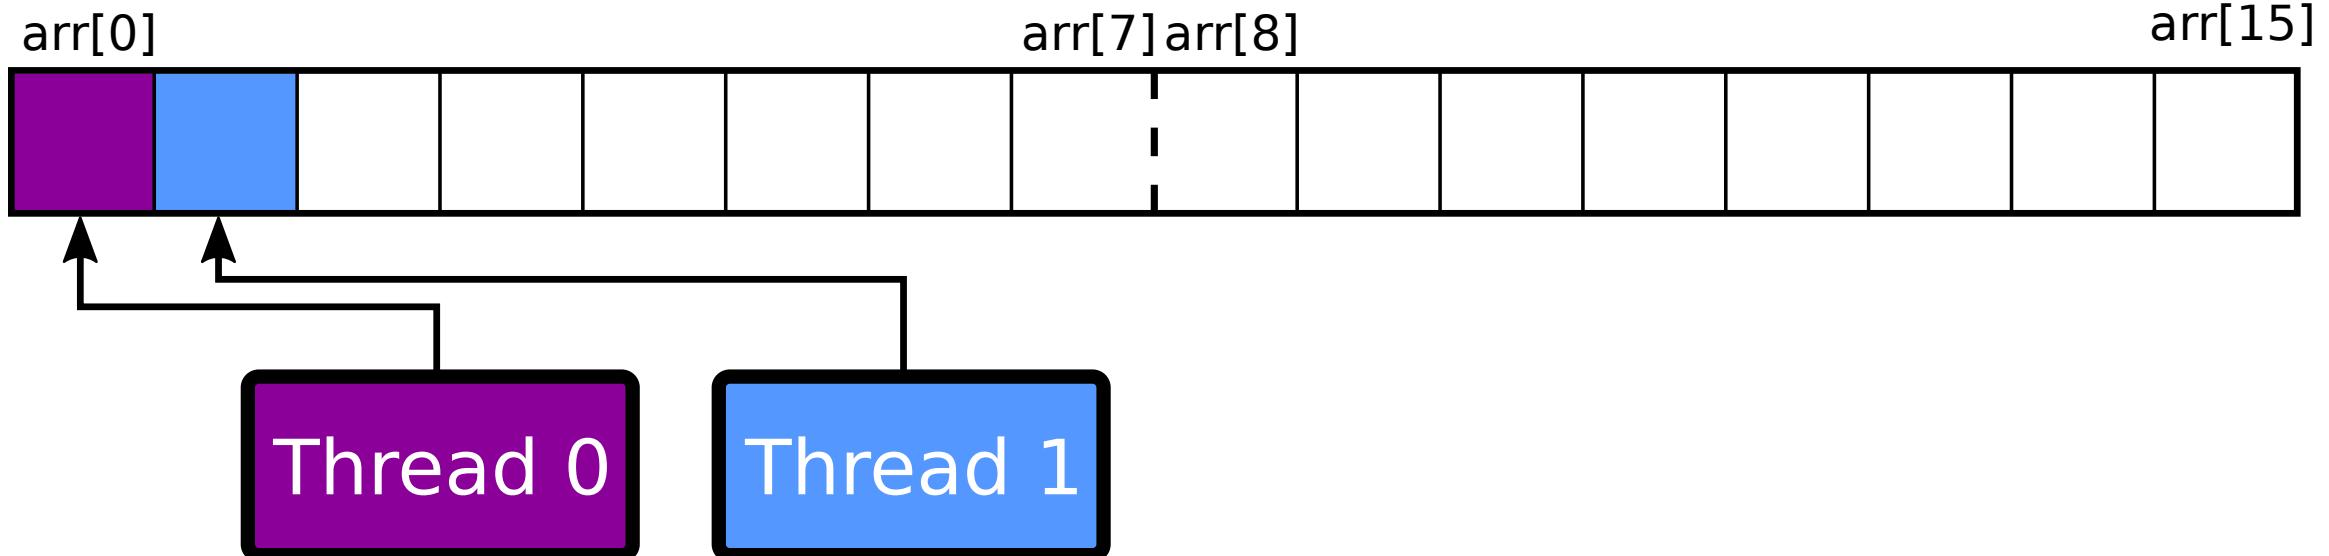


Thread 0

Thread 1

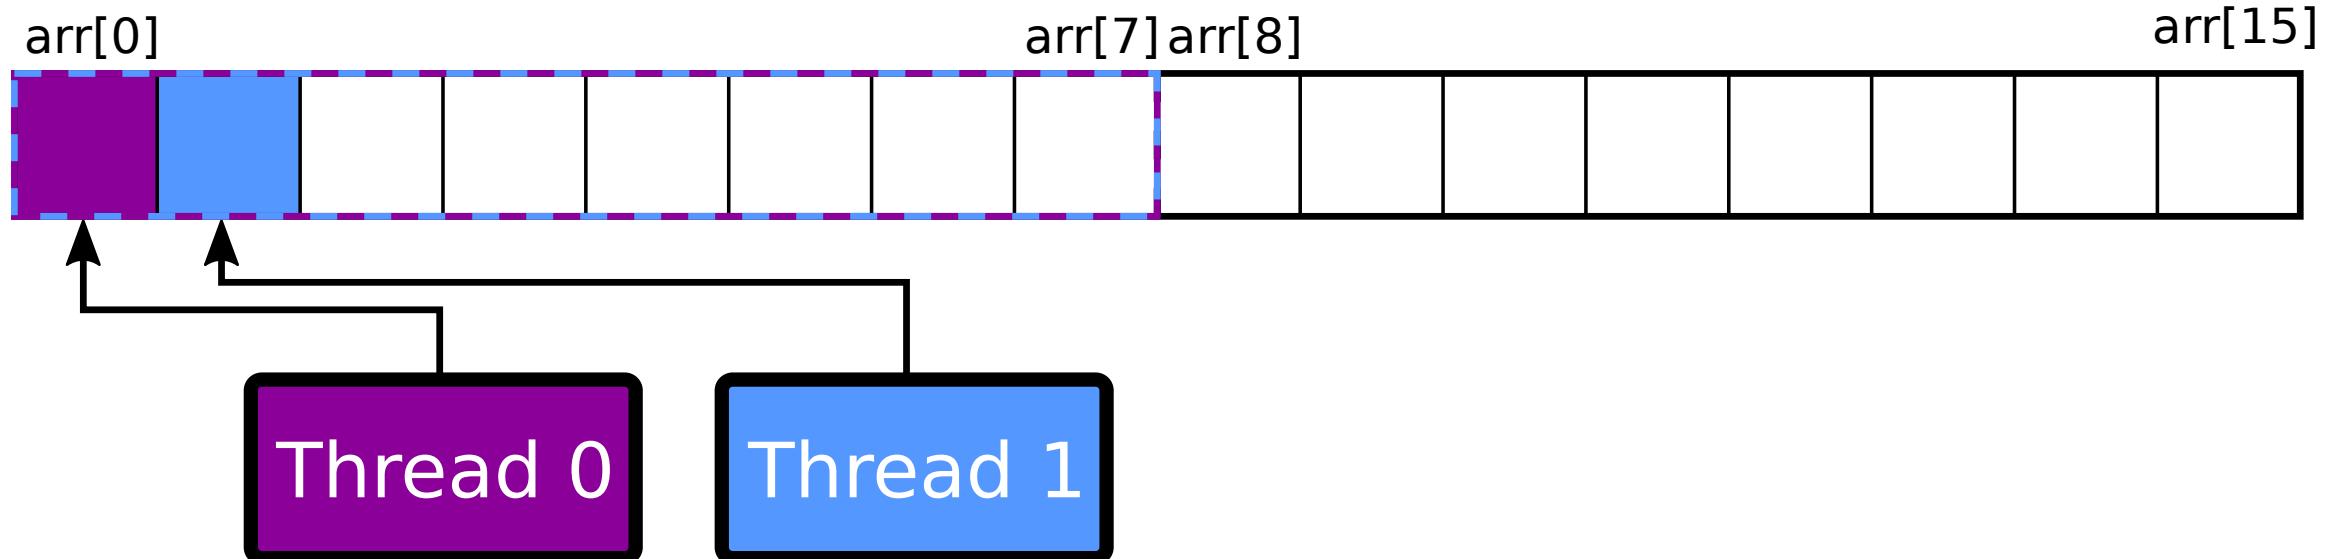
False sharing

double arr[16];



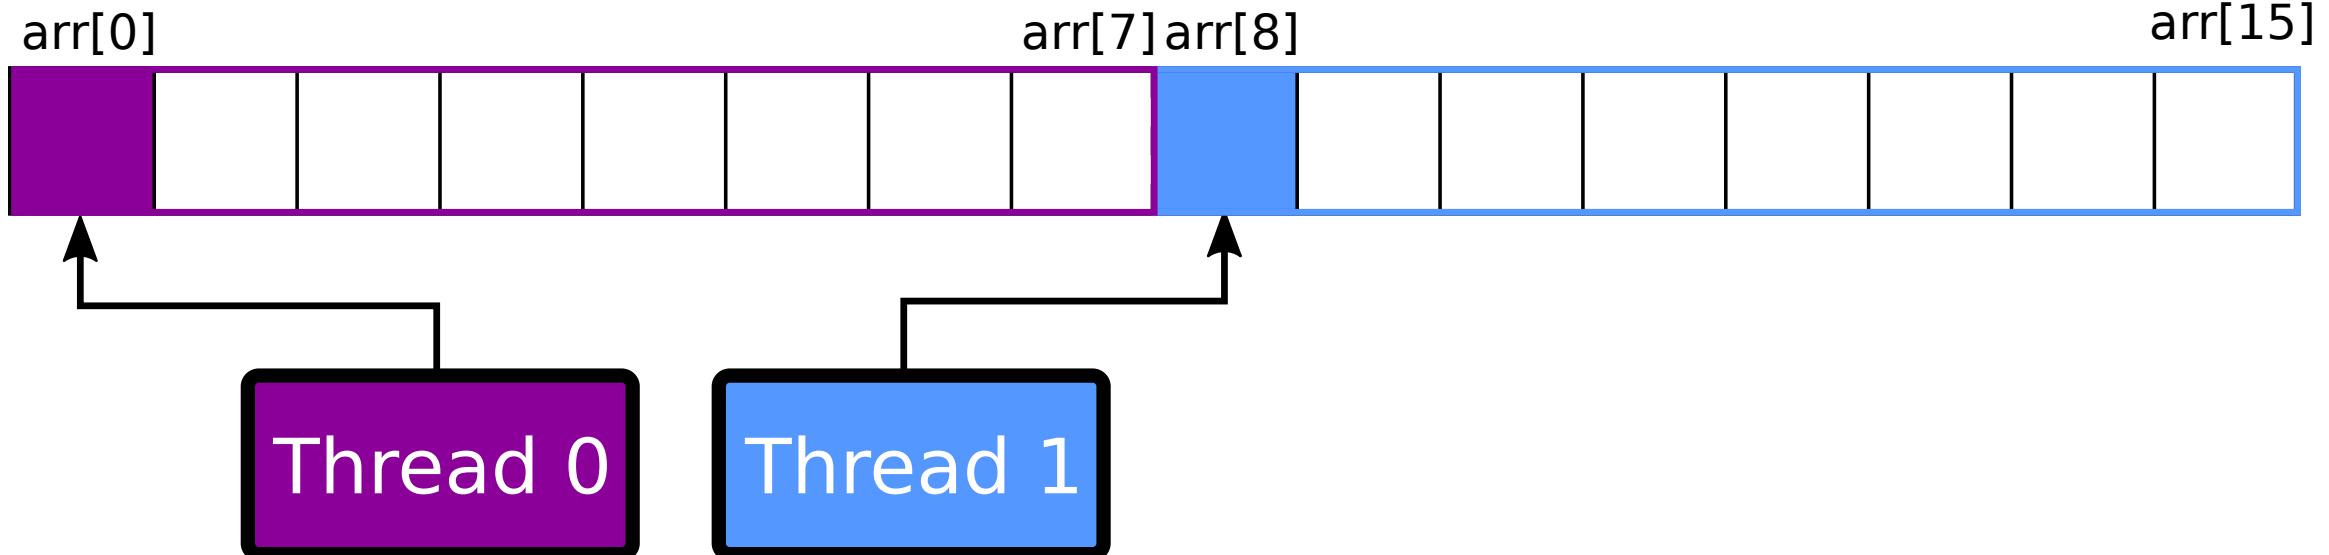
False sharing

double arr[16];



False sharing

double arr[16];



How to measure?

`l2_rqsts.all_rfo`

How many times some core invalidated data in other cores?

How to measure?

`l2_rqsts.all_rfo`

How many times some core invalidated data in other cores?

```
$ perf stat -e l2_rqsts.all_rfo ./example3
1 thread    ->      59 711
2 threads   -> 1 112 258 710
```